The Design and Implementation of the Synchronous Language CÉU

Francisco Sant'Anna, Departamento de Informática e Ciência da Computação, UERJ
Roberto Ierusalimschy, Departamento de Informática, PUC–Rio
Noemi Rodriguez, Departamento de Informática, PUC–Rio
Silvana Rossetto, Departamento de Ciência da Computação, UFRJ
Adriano Branco, Departamento de Informática, PUC–Rio

CÉU is a synchronous language inspired by Esterel with a simpler semantics and more fine-grained control over program execution. CÉU uses an event-triggered notion of time that enables compile-time checks to detect conflicting concurrent statements, resulting in deterministic and concurrency-safe programs. Using Esterel as a base, we present the particularities of our design, such as stack-based internal events, concurrency detection, safe integration with *C*, and first-class timers. CÉU targets soft real-time systems. more abstract and less assumptions about the platform less amenable We also present two implementation back ends: one aiming for resource efficiency and interoperability with *C*, and another as a virtual machine that allows remote reprogramming.

Additional Key Words and Phrases: Concurrency, Determinism, Embedded Systems, Esterel, Synchronous, Reactivity

1. INTRODUCTION

An established alternative to *C* in the field of embedded systems is the family of reactive synchronous languages [Benveniste et al. 2003]. Two major styles of synchronous languages have evolved: in the *control–imperative* style, programs are structured with control flow primitives, such as parallelism, repetition, and preemption; in the *dataflow–declarative* style, programs can be seen as graphs of values, in which a change to a value is propagated through its dependencies without explicit programming. Of the control-based languages, Esterel [Boussinot and De Simone 1991] was the first to appear and succeed, influencing a number of embedded languages, such as *Reactive-C* [Boussinot 1991], *OSM* [Kasten and Römer 2005], *Sync-C* [Von Hanxleden 2009], and *PRET-C* [Andalam et al. 2010].

Despite its success and influence, Esterel has a complex semantics that requires careful static analysis to detect and refuse programs with *causality* and *schizophrenia* problems [Berry 1996; Shiple et al. 1996; Sentovich 1997; Boussinot 1998; Schneider and Wenz 2001; Tardieu and De Simone 2004; Edwards 2005; Yun et al. 2013]. A complex semantics not only challenges the analysis and compilation of programs, but also results in incompatible and non-compliant implementations. Above all, it also affects the programmer's understanding about the code, who, ultimately, has to solve the errors when facing corner cases. Another drawback of the Esterel semantics consists of non-deterministic execution for intra-reaction statements. Non determinism prevents threads from interacting with stateful system calls safely and makes shared-memory concurrency not as straightforward as reading and writing to shared variables.

In this work, we present CÉU, a new programming language that inherits the synchronous and imperative mindset of Esterel but diverges in some fundamental semantic aspects. CÉU proposes a semantics with fine-grained control for intra-reaction execution, and is amenable to concurrency checks that improve safety. The list that follows summarizes the contributions behind the design of CÉU:

— Unique and queue-based external events, which define the notion of time in CÉU.
— Stack-based internal events for intra-reaction communication, which also provides a limited form of coroutines.
— Static concurrency checks to detect suspicious concurrent statements.

— Safe integration with *C* that enforces finalization for external resources.
— First-class timers with dedicated syntax and automatic synchronization.

We also present a lightweight single-threaded implementation of CÉU with two back ends: one aiming for resource efficiency and interoperability with *C*, and another as a virtual machine that allows remote reprogramming. Our implementations target resource-constrained devices, such as *Arduino* and *MICAz* sensor nodes based on 8-bit microcontrollers[1], showing a practical aspect of the proposed semantics.

In previous work [Sant'Anna et al. 2013; Branco et al. 2015], we employed CÉU in the context of wireless sensor networks, developing a number of applications, protocols, and device drivers. We evaluated the expressiveness of CÉU in comparison to event-driven code in *C* and attested a reduction in source code size (around 25%) with a small increase in memory usage (around 5–10% for *text* and *data*) [Sant'Anna et al. 2013]. For the *VM* back end, applications have a bytecode footprint in the order of hundreds of bytes and can be transmitted over the air in a few packets [Branco et al. 2015].

The rest of the paper is organized as follows: Section 2 discusses the design of CÉU, focusing on the fundamental differences to Esterel. Section 3 presents the *C* and *VM* implementation back ends. Section 4 discusses other synchronous languages targeting embedded systems. Section 5 concludes the paper.

2. THE DESIGN OF CÉU

CÉU is a synchronous reactive language inspired by Esterel in which programs advance in a sequence of discrete reactions to external events. Like Esterel, CÉU is designed for control-intensive applications, supporting concurrent lines of execution, known as *trails*, and broadcast communication through events. Internal computations within a reaction (e.g. expressions, assignments, and system calls) are considered to take no time in accordance with the synchronous hypothesis [de Simone et al. 2005]. An `await` is the only statement that halts a running reaction and allows a program to advance in this discrete notion of time. To ensure that reactions run in bounded time and programs always progress, loops are statically required to contain at least one `await` statement in all possible paths [Sant'Anna et al. 2013; Berry 2000]. CÉU shares the same limitations with (core) Esterel and synchronous languages in general [Berry 1993]: computations that run in unbounded time (e.g., cryptography, image processing) do not fit the zero-delay hypothesis, and cannot be directly implemented.

Figure 1 illustrates the syntactic similarities between the languages, showing side-by-side the implementations in Esterel `[a]` and CÉU `[b]` for the following control specification: *"Emit an output O as soon as inputs A and B occur. Reset this behavior each time input R occurs"* [Berry 2000]. The first phrase of the specification, awaiting and emitting the events, is translated almost identically in the two languages (ln. 5–10, in both implementations), as Esterel's '‖' and CÉU's par/and constructs are equivalent. For the second phrase, the reset behavior, the Esterel version uses an `abort-when` statement (ln. 4–11) which, in this case, serves the same purpose as CÉU's par/or (ln. 4–13): the occurrence of event R aborts the awaiting statements in parallel and restarts the enclosing `loop`.

In the subsections that follow, we discuss the main differences between CÉU and Esterel: Unique and queue-based external events (2.1); Stack-based internal events (2.2); Static concurrency checks (2.3); Safe integration with *C* (2.4); and First-class synchronized timers (2.5). We finish the section with a summary of our design (2.6). We present the formal specification of the semantics of CÉU in a separate paper [Sant'Anna 2013].

---

[1]Both *Arduino* and *MICAz* use the 8-bit *ATmega328* microcontroller with 32K of FLASH and 2K of SRAM.

```
1  input  A, B;              1  input  void A, B;
2  output O;                 2  output void O;
3  loop                      3  loop do
4      abort                 4      par/or do
5          [                 5          par/and do
6              await A       6              await A;
7          ||                7          with
8              await B       8              await B;
9          ];                9          end
10         emit O           10          emit O;
11     when R               11      with
12 end                      12          await R;
13                          13      end
14 .                        14  end
```

| [a] Esterel | [b] CÉU |

Fig. 1. A control specification implemented in Esterel and CÉU: *"Emit O after A and B, resetting each R."* A par/and terminates when both trails in parallel terminate. A par/or terminates when any trail terminates, aborting the other.

2.1. Unique and Queue-Based External Events

Esterel defines time as a discrete sequence of logical unit instants or "ticks". At each tick, the program reacts to an arbitrary number of simultaneous input events from the environment. In contrast, CÉU defines time as a discrete sequence of reactions to unique input events. At each input event, which constitutes a logical unit of time, the program reacts exclusively to it. The event-triggered execution of a program in CÉU is as follows [Sant'Anna et al. 2013]:

(1) The program initiates the "boot reaction" in a single trail (but parallel constructs may create new trails).
(2) Active trails execute until they await or terminate, one after the other. This step is named a *reaction chain*, and always runs in bounded time.
(3) The program goes idle and the environment takes control.
(4) On the occurrence of a new external input event, the environment awakes *all* trails awaiting that event. It then goes to step 2.

A program must react to an event completely before handling the next one. Based on the synchronous hypothesis, a program takes a negligible time on step 2 and is always idle on step 3. In practice, if a new external input event occurs while a reaction chain is running, it is enqueued to occur in a subsequent reaction.

Figure 2 compares the discrete notions of time in two variations of Esterel and in CÉU. The box Real World shows event occurrences over a continuous timeline. The other boxes show how the same occurring events fit differently in each logical notion of time.

— [Box-1]: *Esterel with fixed-length ticks* [Li et al. 2005]. We assume a reaction R(boot) at tick-0 which happens before any input. The input A "physically" occurs during the boot reaction but, because time is discrete, its corresponding reaction only executes in the next tick. Note that R(A) takes more time than tick-1 and invades tick-2, causing a *timing violation* [Li et al. 2005]. The events B and C occur during tick-1 and are delayed to happen *simultaneously* at tick-2 with R(B+C). Since no new events occur during tick-2, the CPU stays idle during the whole tick-3. Finally, one instance of event D and two instances of event E occur during the idle tick-3. However, only one occurrence of E can be considered in R(D+E).
— [Box-2]: *Esterel with variable-length ticks* [Roop et al. 2004]. This approach avoids the timing violation for R(A) and also results in smaller idle periods because it adjusts the tick lengths to match the CPU times for the reactions. For instance, the

Fig. 2. The discrete notions of time in Esterel and CÉU.

occurrence of D interrupts the idle `tick-3` to react alone as `R(D)` at `tick-4`. Similarly to the fixed-tick approach, only one of the two simultaneous occurrences of `E` is considered at `tick-5`, now because `R(D)` takes too long.

— [`Box-3`]: CÉU *with unique and queue-based input events*. We also assume a reaction `R(boot)` before any input. Because the occurrence of event `A` is unique during `R(boot)`, the behavior in CÉU is similar to `Box-2` for its first two reactions (`tick-0` and `tick-1`). However, CÉU does not consider the events `B` and `C` as simultaneous, and handles each in subsequent reactions `R(B)` and `R(C)`. We assume the CPU times for `R(B+C)` in Esterel and `R(B)+R(C)` in CÉU to be roughly the same. This way, the first idle periods in `Box-2` and `Box-3` coincide. Finally, CÉU recognizes and reacts to the two instances of `E` independently, which are handled in sequence. Note that, by definition, each reaction corresponds to an event occurrence of the same length.

We decided for the unique and queue-based semantics in CÉU for the reasons that follow:

— **A "tick" is too abstract and imprecise:** Outside the domain of hardware specification, a tick has no natural counterpart in the real world. Also, since ticks require no time regularity [Berry and Sentovich 2001], the two approaches for Esterel in Figure 2 are legitimate, but lead to different behaviors for the same sequence of inputs.

```
1   input  A;   // external          1   input void A;   // external (in uppercase)
2   signal B;   // internal          2   event void b;   // internal (in lowercase)
3   [[                               3   par/and do
4       await A;                     4       await A;
5       emit B;                      5       emit b;
6       call f();                    6       _f();
7   ||                               7   with
8       await B;                     8       await b;
9       call g();                    9       _g();
10  ]]                               10  end
```

[a] Esterel   [b] CÉU

Fig. 3.   Internal signals (events) in Esterel and CÉU: similar syntax, but different semantics.

— **Events are never absolutely simultaneous:** From a rigorous point of view, event occurrences are infinitesimal, having zero probability of being simultaneous. This way, we consider that the notion of simultaneity should not be imposed by the language, but defined explicitly for each use case. In the case of Esterel, simultaneity depends on the imprecise length of discrete ticks. For instance, in `box-1` and `box-2` of Figure 2, Esterel considers the events B and C to be simultaneous, even though A and B "physically" happen much closer to one another. In Section 2.5, after introducing internal events and first-class timers, we show an example to illustrate how to detect simultaneous button clicks in CÉU.

— **Unique input events imply mutual exclusion:** Reactions to multiple events in CÉU never overlap because they are atomic. Automatic mutual exclusion simplifies reasoning about concurrency and is a prerequisite for the concurrency checks to be discussed in Section 2.3.

The synchronous hypothesis for CÉU holds if the reactions run faster than the rate of incoming input events. Otherwise, the application continuously accumulates delays between the real occurrence and actual reaction of a given event. This is also the case for the variable-length-tick approach of Esterel, since the more inputs to handle, the longer the reaction takes, and the more inputs accumulate for subsequent ticks. For the fixed-length-tick approach of Esterel, a breach in the synchronous hypothesis causes timing violations, requiring *worst case reaction time* analysis to infer appropriate tick lengths [Li et al. 2005].

A limitation of event-triggered execution is that all program behavior is purely reactive, given that no code can execute in the absence of inputs. Tick-triggered execution allows for active behavior, since code can execute regularly on every tick. Although CÉU supports active *asynchronous* execution [Sant'Anna et al. 2012], its synchronous core is still purely reactive.

2.2. Stack-Based Internal Events

In Esterel, the behavior of internal and external signals is equivalent. In CÉU, in contrast with queue-based external events, internal events follow a stack-based execution policy similar to subroutine calls in typical programming languages. Figure 3 illustrates the use of internal signals (events) in Esterel [a] and CÉU [b]. In Esterel, when A occurs, the program emits B (ln. 4–5) and both events become active, resulting in the invocation of f() and g() in no particular order (ln. 6,9). In CÉU, when A occurs, the program behaves as follows:

(1) 1st trail awakes, broadcasts b, and pauses (ln. 4–5).
(2) 2nd trail awakes, calls _g(), and terminates (ln. 8–9). *(No other trails awake to b.)*
(3) 1st trail (on top of the stack) resumes, calls _f(), and terminates (ln. 5–6).
(4) Both trails have terminated, so the par/and rejoins, and the program also terminates.

```
1   event int* inc; // subroutine 'inc'
2   par/or do
3       loop do       // definitions are loops
4           var int* p = await inc;
5           *p = *p + 1;
6       end
7   with
8       var int v = 1;
9       <...>
10      emit inc => &v; // call 'inc'
11      _assert(v==2);  // after return
12  end
```

Fig. 4.   Subroutine `inc` is defined in a loop (ln. 3–6), in parallel with the caller (ln. 8–11).

Internal events provide fine-grained execution control and can express a limited form of subroutines, as depicted in Figure 4. The "subroutine" `inc` is defined as a loop (ln. 3–6) that continuously awaits its identifying event (ln. 4), incrementing the value passed by reference (ln. 5). A trail in parallel (ln. 8–11) invokes the subroutine through an `emit inc` (ln. 10) in reaction to some code (ln. 9). Given the stacked execution for internal events, the calling trail pauses, the subroutine awakes (ln. 4), runs its body (yielding v=2), loops, and awaits the next "call" (ln. 4, again). Only after this sequence, the calling trail resumes and passes the assertion test (ln. 10–11).

CÉU also supports nested `emit` invocations for internal events. For instance, the body of the subroutine `inc` in Figure 4 could `emit` another event after awaking (ln. 4), creating a new level in the stack. The runtime stack constitutes fine-grained micro reactions, with one on top of the other, all inside the same reaction to an external event.

On the one hand, this form of subroutine has a significant limitation that it cannot express recursive calls: an `emit` to itself is always ignored, given that a running body cannot be awaiting itself. On the other hand, this very same limitation brings some important safety properties to subroutines: first, they are guaranteed to react in bounded time; second, memory for locals is also bounded, not requiring data stacks. Also, this form of subroutine can use the other primitives of CÉU, such as parallel compositions and the `await` statement. In particular, they await keeping context information such as locals and the program counter, similarly to coroutines [Moura and Ierusalimschy 2009]. In previous work, we build other advanced control mechanisms on top of internal events, such as resumable exceptions and reactive variables [Sant'Anna et al. 2013].

Another distinction regarding event handling in comparison to CÉU is that Esterel supports same-cycle bi-directional communication [Edwards 1999], i.e., two threads can react to one another during the same cycle due to mutual signal dependencies. CÉU takes a different approach, posing a tradeoff that an `await` is only valid for the next reaction, i.e., if an `await` and `emit` occur simultaneously in parallel trails, the `await` does not awake. These *delayed awaits* avoid corner cases of instantaneous termination and re-execution of statements in the same reaction (known as *schizophrenic statements* [Berry 1996]).

TODO: mesmo sem mutual dependency, falar de semantica construtiva e lei da coerencia

The example in Figure 5 illustrates delayed awaits, which prevents infinite execution by design. Both sides of the `par/or` have an `await` statement (ln. 4,7), which characterizes the enclosing `loop` as non instantaneous (ln 2–9). However, if the `emit e` (ln. 6) could awake the `await e` instantaneously (ln. 4), the `par/or` would terminate and restart the `loop` instantaneously, resulting in infinite execution. In atypical scenarios requiring immediate awake, delayed awaits can be circumvented by placing the code

```
1  event void e,f;
2  loop do
3      par/or do
4          await e;
5      with
6          emit e;       // w/o delayed awaits, the emit awakes 1st trail
7          await f;      // and restarts the loop instantaneously
8      end
9  end
```

Fig. 5.   Delayed awaits prevents re-execution of statements by design.

to execute before the `await`. On the one hand, we transfer the burden of dealing with these corner cases to the programmer. On the other hand, we simplify the semantics of the language and eliminate the need for complex analysis to deal with schizophrenic statements.

2.3. Static Concurrency Checks

Embedded applications make extensive use of global memory and shared resources, such as through memory-mapped registers and system calls to device drivers. Hence, an important goal of CÉU is to ensure a reliable behavior for programs with concurrent lines of execution sharing memory and interacting with the environment.

Esterel is only deterministic with respect to external behavior: "the same sequence of inputs always produces the same sequence of outputs" [Berry 2000]. However, the execution order for operations within a reaction is non-deterministic: "if there is no control dependency, as in (`call f1() || call f2()`), the order is unspecified and it would be an error to rely on it" [Berry 2000]. A number of Esterel-inspired synchronous languages enforce an arbitrary execution order for statements in multiple lines of execution to achieve intra-reaction determinism (*Reactive C* [Boussinot 1991], *Protothreads* [Dunkels et al. 2006], *SOL* [Karpinski and Cahill 2007], *SC* [Von Hanxleden 2009], and *PRET-C* [Andalam et al. 2010]). CÉU also takes the deterministic approach and, when multiple trails are active during the same reaction, they are scheduled in lexical order, i.e., in the order they appear in the program source code.

Even so, we consider that enforcing an arbitrary execution order can be misleading in some cases. For instance, consider the two examples in Figure 6, both defining a shared variable (ln. 2), and assigning to it in parallel trails (ln. 5, 8). In the example `[a]`, the two assignments to x can only execute in reactions to different events A and B, which cannot occur simultaneously by definition (Section 2.1). Hence, for the sequence A->B, x becomes 4 (`(1+1)*2`), while for B->A, x becomes 3 (`(1*2)+1`). In the example `[b]`, the two assignments to y are simultaneous because they execute in reaction to the same event A. Since CÉU employs lexical order for intra-reaction statements, the execution is still deterministic, and y always becomes 4 (`(1+1)*2`). However, an (apparently innocuous)

```
1  input void A, B;          1  input void A;
2  var int x = 1;            2  var int y = 1;
3  par/and do               3  par/and do
4      await A;             4      await A;
5      x = x + 1;           5      y = y + 1;
6  with                     6  with
7      await B;             7      await A;
8      x = x * 2;           8      y = y * 2;
9  end                      9  end
```

       [a] Accesses to x are safe   [b] Accesses to y are unsafe

Fig. 6.   Shared-memory concurrency in CÉU: example `[a]` is safe because the trails access x atomically in different reactions; example `[b]` is unsafe because both trails access y in the same reaction.

change in the order of trails modifies the semantics of the program, which we consider unsafe.

To mitigate this threat, CÉU performs concurrency checks at compile time to detect conflicting accesses to shared variables: if a variable is written in a trail segment, then a concurrent trail segment cannot read or write to that variable, nor dereference a pointer of that variable type. Concurrency in CÉU is characterized when two or more trail segments in parallel react to the same input event. A trail segment is a sequence of statements followed by an `await` (or termination). Considering the examples in Figure 6:

— The assignments to x (`[a]`:2,5) **cannot** be concurrent because they are **not** in parallel trails.
— The assignments to x (`[a]`:5,8) **cannot** be concurrent because they **cannot** execute during the same reaction.
— The assignments to y (`[b]`:5,8) **can** be concurrent because they are in parallel trails and **can** execute during the same reaction.

The detection algorithm, which is depicted in Section 3.1, inspects all possible `await` statements that precede a variable access and keeps a list with all corresponding awaking events. Then, it checks all accesses in parallel trails to see if they share an awaking event. If it is the case, the compiler warns about the suspicious accesses.

The uniqueness of input events within reactions makes this analysis possible, otherwise, any two trail segments in parallel could be concurrent, even if they react to different input events. Note that the static checks are optional and do not affect the semantics of the program.

2.4. Safe Integration with *C*

In CÉU, any identifier prefixed with an underscore is passed unchanged to the *C* compiler that generates the final binary. Therefore, access to *C* is straightforward and syntactically trackable. Similarly to Esterel with the `call` primitive, external calls are assumed to be instantaneous [Berry 2000]. This way, programs should only resort to *C* for asynchronous functionality, such as non-blocking I/O, or simple `struct` accessors, but never for control purposes.[2]

2.4.1. Concurrency Checks. As a safety measure, the concurrency checks of Section 2.3 also consider concurrent calls and accesses to external symbols in *C*. As an example, the program in Figure 7.a defines four external symbols inside a `native` block with standard declarations in *C* (ln. 1–6). During the boot reaction, two trails react con-

---

[2]In CÉU, it is possible to restrict the available *C* symbols as a compile-time option.

```
1   native do                                native @const _NUM;
2       ##define NUM 10                       native @pure  _id();
3       void f  (void)  { <...> }             native @safe  _f() with _g();
4       void g  (int v) { <...> }
5       int  id (int v) { <...> }
6   end
7   par/and do
8       _f();
9   with
10      _g(_id(_NUM));
11  end
```

[a] Definitions and uses of symbols   [b] Annotations for the symbols in [a]

Fig. 7. The unsafe program in [a] only compiles with the annotations in [b].

```
par do                                          1   native do
    <...> // animate and redraw "background"    2       #define redraw_non_commutative redraw
    _redraw(background);                         3   end
with                                            4   @safe _redraw_non_commutative with
    <...> // animate and redraw "foreground"    5       _redraw_non_commutative;
        _redraw(foreground);                    6   par do
end                                             7       <...> // animate and redraw "background"
                                                8           _redraw_non_commutative(background);
                                                9   with
                                                10      <...> // animate and redraw "fore"
                                                11          _redraw_non_commutative(foreground);
.                                               12  end
```

[a] _redraw **cannot** be concurrent   [b] _redraw_non_commutative **can** be concurrent

Fig. 8.   Making the non-commutative redrawing calls from [a] to compile in [b].

currently inside the par/and (ln. 7–11): the first trail calls symbol _f (ln. 8), while the second calls _g and _id, and also reads _NUM (ln. 10). Since CÉU does not inspect any code in C, it complains about suspicious concurrent accesses between _f and all symbols in the second trail.

2.4.2. Annotations. The annotations in Figure 7.b provide hints to the compiler about the semantics of the C symbols in program [a], which now compiles successfully:

— _NUM is a constant symbol, meaning that it is safe to use it concurrently with any other symbol in the program.
— _id is a pure function, also meaning that it is safe to call it concurrently with any other symbol in the program.
— Both _f and _g are impure, but have non-conflicting commutative effects, and can be safely called concurrently.

From our experience, however, we find that programs often need non-commutative concurrent calls. This is the case for logging (e.g., calls to _printf in trails in parallel) and for redrawing objects in the screen. Figure 8.a shows an abstract code to animate and redraw the objects background and foreground in trails in parallel. In typical graphical APIs, consecutive calls to _redraw overwrites conflicting pixels, which makes the calls non commutative and prevents the code to compile. However, in this case we want to rely on lexical order to always redraw the background object before the foreground object. Therefore, in Figure 8.b, we redefine _redraw to an "intimidating name" _redraw_non_commutative (ln. 2, to be explicit about its effect), and annotate it as safe (ln. 4–5) to make the code compile successfully.

2.4.3. Finalization. Esterel's abort and CÉU's par/or statements provide orthogonal abortion of lines of execution, which is a distinctive feature of synchronous languages in comparison to asynchronous languages [Berry 1993]. However, aborting lines of execution that deal with external resources may lead to inconsistencies. For this reason, Esterel and CÉU provide a finalize construct to unconditionally execute a series of statements even if the enclosing block is aborted and does not terminate normally.

In the example in Figure 9 in Esterel [a] and CÉU [b], the calls to lock and unlock represent accesses to an external resource. After we lock the resource (ln. [a]:4 and [b]:3), we perform some operations in subsequent reactions to input A (ln. [a]:5–8 and [b]:7–10), and then we unlock the resource (ln. [a]:10 and [b]:5). Note that if the aborting input B (ln. [a]:12 and [b]:12) occurs after the lock but before the reactions to A, we still want to call unlock to safely release the resource. In Esterel and CÉU, the finalize clause (ln. [a]:10 and [b]:5) executes automatically if the enclosing block (ln. [a]:3–1 and [b]:3–10) is externally aborted (ln. [a]:12 and [b]:12).

```
1   input A, B;                           1   input void A, B;
2   abort                                 2   par/or do
3       finalize                          3       _lock();
4           call lock();                  4       finalize with
5           await A;                      5           _unlock();  // defer execution
6           <...>;      // do something   6       end
7           await A;                      7       await A;
8           <...>;      // do something   8       <...>;              // do something
9       with                              9       await A;
10          call unlock();               10       <...>;              // do something
11      end                              11   with
12  when B                               12       await B;
13  .                                    13   end
```

[a] Esterel   [b] CÉU

Fig. 9.   Finalization in Esterel and CÉU: after the call to lock, both languages guarantee to call unlock if the enclosing block aborts when B occurs.

CÉU also enforces the use of `finalize` for system calls that deal with pointers representing resources:

— If CÉU **passes** a pointer to a system call, the pointer represents a **local** resource that requires finalization.
— If CÉU **receives** a pointer from a system call return, the pointer represents an **external** resource that requires finalization.

CÉU tracks the interaction of system calls with pointers and requires finalization clauses to accompany them. In the example in Figure 10.a, the local variable msg (ln. 2) is an internal resource passed as a pointer to _send_request (ln. 5), which is an asynchronous call that transmits the buffer in the background. If the block aborts (ln. 11) before receiving an acknowledge from the environment (ln. 9), the local msg goes out of scope and the external transmission now holds a *dangling pointer*. The finalization ensures that the transmission also aborts (ln. 7). In the example in Figure 10.b, the call to _fopen (ln. 4) returns an external file resource as a pointer. If the block aborts (ln. 12) during the await A (ln. 9), the file remains open as a *memory leak*. The finalization ensures that the file closes properly (ln. 6). In both cases, the code does not compile without the `finalize` construct.

Note that the illustrative example in Figure 9 does not manipulate pointers (i.e., the resource is a *singleton*). That case is an example of a bad and unsafe API to expose to CÉU because the compiler will not enforce the use of finalization.

```
1   par/or do                            1   par/or do
2       var _buffer_t msg;               2       var _FILE* f;
3       <...> // prepare msg             3       finalize
4       finalize                         4           f = _fopen(...);
5           _send_request(&msg);         5       with
6       with                             6           _fclose(f);
7           _send_cancel(&msg);          7       end
8       end                              8       _fwrite(..., f);
9       await SEND_ACK;                  9       await A;
10  with                                10       _fwrite(..., f);
11      <...>                            11   with
12  end                                 12       <...>
13  .                                   13   end
```

[a] Local resource finalization   [b] External resource finalization

Fig. 10.   CÉU enforces the use of finalization to prevent *dangling pointers* for local resources and *memory leaks* for external resources.

```
1   var int v;                          1   par/or do
2   await 10ms;                         2       await 10ms;
3   v = 1;                              3       <...>       // any non—awaiting sequence
4   await 1ms;                          4       await  1ms;
5   v = 2;                              5       v = 1;
6                                       6   with
7                                       7       await 12ms;
8                                       8       v = 2;
9   .                                   9   end
```

[a]   [b]

Fig. 11.   First-class timers in CÉU.

2.5. First-Class Timers

Activities that involve reactions to *wall-clock time*[3] appear in typical patterns of embedded development [Bourke and Sowmya 2009], such as timeout watchdogs and sensor samplings. However, the interaction between system clocks and programs is not absolutely precise, a fact that is usually ignored in the development process. We define the difference between a requested timeout and the actual expiring time as the *residual delta time (delta)*. Without explicit manipulation, the recurrent use of timed activities in a row (or in a loop) may accumulate a considerable amount of deltas that can lead to incorrect behavior in programs.

The `await` statement of CÉU supports wall-clock time and handles deltas automatically, resulting in more robust applications. Due to platform overheads and occasional slow reactions, an `await 10ms` suspends a line of execution for *at least* 10 milliseconds. In the example in Figure 11.a, suppose that after the first `await` request, the underlying system gets busy and takes 15ms to notify CÉU. The scheduler will notice that the `await 10ms` (ln. 2) has not only already expired, but is delayed with `delta=5ms`. Then, the awaiting trail awakes, sets `v=1` (ln. 3), and invokes `await 1ms` (ln. 4). As the current delta is still higher than the requested timeout (i.e. $5ms > 1ms$), the trail is rescheduled for execution, now with `delta=4ms`.

Delta compensation also implies that timers can be added and compared when reasoning about time as a physical quantity. For instance, in the example in Figure 11.b, although the scheduler cannot guarantee that the first trail terminates exactly in 11ms (ln. 2,4), it can at least ensure that the program always terminates with `v=1`. Given that any non-awaiting sequence is considered to take no time in the synchronous model, the first trail (ln. 2–5) is guaranteed to awake first and terminate the `par/or` before the second trail executes (ln. 7–8), because $10 + 1 < 12$. A similar program in a language without first-class support for timers would depend on the execution timings for the code marked as <...>, making the reasoning about the execution behavior more difficult.

In Section 2.1, we argue that event occurrences are infinitesimal and can never be absolutely simultaneous. However, the "sensation of simultaneity" is not infinitesimal, but actually increases with the inaccuracy of the observer (e.g., a human being). Therefore, we consider that simultaneity should be defined case by case, and should not be imposed by the language. First-class timers simplify the implementation of application-defined simultaneity. Figure 12 emulates a `middle_click` event (ln. 3) in terms of "simultaneous" occurrences of LEFT_CLICK and RIGHT_CLICK (ln. 1–2). If both events occur, we emit the internal event `middle_click` (ln. 6–7). However, if one of them occurs and the *200ms* timer expires (ln. 9–10), we abort the whole behavior with the `par/or` (ln. 5) and try again with the enclosing `loop` (ln. 4). In this specification, "simultaneous" means "within 200 milliseconds", which is a huge amount of time for a

---

[3]By wall-clock time we mean the passage of time from the real world, measured in hours, minutes, etc.

```
1   input void LEFT_CLICK;                             13   #define AWAIT_AND(e1, e2)    \
2   input void RIGHT_CLICK;                            14       par/and do           \
3   event void middle_click;                           15           await e1;        \
4   loop do                                            16       with                 \
5       par/or do                                      17           await e2;        \
6           AWAIT_AND(LEFT_CLICK, RIGHT_CLICK);        18       end                  \
7           emit middle_click;                         19   #define AWAIT_OR(e1, e2)     \
8       with                                           20       par/or do            \
9           AWAIT_OR(LEFT_CLICK, RIGHT_CLICK);         21           await e1;        \
10          await 200 ms;                              22       with                 \
11      end                                            23           await e2;        \
12  end                                                24       end
```

Fig. 12. Application defines that a middle_click event occurs whenever both LEFT_CLICK and RIGHT_CLICK occur within 200 milliseconds. The macros AWAIT_AND (ln. 13–18) and AWAIT_OR (ln. 19–24) are simple expansions to a par/and and par/or for better readability.

language-defined tick. For instance, a similar implementation of this specification in Esterel would not rely on the tick notion of simultaneity either.

2.6. Summary

CÉU aims to offer a simpler semantics than Esterel with more determinism and fine-grained control over program execution. The following list summarizes the contributions of our design in this direction:

— *Event-triggered notion of time bound to the semantics of the language.* Event-driven programming is popular in many domains, such as server and GUI development. We believe that programmers are more familiar with dealing with events in isolation, which simplifies the reasoning about concurrency. In addition, the uniqueness of external events is a prerequisite for the concurrency checks of CÉU.
— *Deterministic intra-reaction execution and communication.* Determinism in CÉU is "all-inclusive" and does not depend on additional levels of static analysis. It encompasses the whole language, including memory accesses, system calls, and stack-based internal events. Programmers can always figure out which statement executes next, making runtime analysis and debugging easier.
— *Static concurrency checks.* Although execution is deterministic, the CÉU compiler still advises about suspicious statements that can react concurrently to the same event.
— *Safe integration with* C. When dealing with concurrent system calls, programmers can provide annotations to reduce false positives in the static checks, or to force non commutative concurrent behavior. CÉU also requires finalization clauses to handle pointers representing resources.
— *First-class timers with dedicated syntax and automatic synchronization.* Given the omnipresence of timers in embedded systems, a dedicated syntax can simplify the development and readability of programs. Furthermore, automatic synchronization releases the programmer from the burden of adjusting timers in sequence and in parallel.

first-class timers non-intrusive
Our synchronous and deterministic approach also leads to some limitations as follows:

— *Execution is purely reactive as result of event-triggered reactions.* Since only event occurrences can start reactions, programs cannot execute proactively in the absence of events. In addition, await statements cannot awake in the same reaction they are reached.

—*Reactions must execute in bounded time due to the synchronous hypothesis.* As a synchronous language, CÉU requires CPU times for reactions to be negligible in comparison to the rate of incoming events.

—*Execution is sequential because of intra-reaction determinism.* The deterministic semantics of CÉU does not make implicit parallelization easy (to be discussed in Section 3.5).

Nonetheless, we advocate keeping a tractable synchronous reactive core with support for shared memory concurrency and deterministic execution. To deal with the limitations above, we recommend memory-isolated parallelizable asynchronous primitives as separate extensions to the synchronous core [Berry et al. 1993] (which are not in the scope of this paper).

3. IMPLEMENTATION

The compilation process of programs in CÉU is composed of three main phases: the *parsing phase* converts the source code in CÉU to an abstract syntax tree (*AST*); the *concurrency checks phase* detects inconsistencies in programs, such as unbounded loops and suspicious concurrent statements; the *code generation phase* converts the *AST* to standard *C* code and augments it with platform-dependent functionality (e.g., system calls) and the runtime of CÉU, compiling everything with *gcc* to generate the final binary.

In previous work [Sant'Anna et al. 2013], we evaluate the implementation of CÉU in comparison to a code base[4] in *nesC* (a *C* variant) [Gay et al. 2003]. Figure 13 compares source size (*number of tokens*), binary size (*ROM*), and memory usage (*RAM*) for a number of standardized network protocols and a radio driver. The small overhead in resource usage shows that the gains in productivity and safety with CÉU make it a viable alternative to *C* in the context of constrained embedded systems. Unfortunately, most real applications in Esterel seem to be closed source. In future work, we plan to cooperate with groups that have more experience and access to real code bases in Esterel for a comprehensive comparison with CÉU.

In the subsections that follow, we discuss implementation details specific to CÉU: concurrency checks for determinism (Section 3.1), static memory allocation for data and trails (Sections 3.2 and 3.3), static scheduling and trail finalization (Section 3.4), single-threaded dispatching (Section 3.5), interaction with the environment (Section 3.6), and the *VM* back end (Section 3.7),

3.1. Concurrency Checks

The compile-time concurrency checks phase detects inconsistencies in CÉU programs. Here, we focus on the algorithm that detects suspicious concurrent statements, such as accesses to shared variables, as discussed in Section 2.3.

For each node representing a statement in the program AST, we keep the set of input events I (*incoming*) that can start the execution of the node, and also the set of input events O (*outgoing*) that can terminate the node. As an example, for the single-statement program `await A`, we have $I = \{boot\}$ and $O = \{A\}$.

A node inherits the set I from its immediate parent and calculates O according to its type, as follows:

---

[4]TinyOS repository: http://github.com/tinyos/tinyos-release/

Fig. 13.   Resource usage for CÉU and *nesC* in the domain of sensor networks.

```
1   input void A, B;
2   var int y;
3   par/or do
4      await A;
5      y = 1;
6   with
7      await B;
8      y = 2;
9   end
10  await A;
11  y = 3;
```

```
1   Stmts I={boot} O={A}
2      Dcl_y I={boot} O={boot}
3      ParOr I={boot} O={A,B}
4         Stmts I={boot} O={A}
5            Await_A I={boot} O={A}
6            Set_y I={A} O={A}
7         Stmts I={boot} O={B}
8            Await_B I={boot} O={B}
9            Set_y I={B} O={B}
10     Await_A I={A,B} O={A}
11     Set_y I={A} O={A}
```

[a] A program in CÉU...   [b] ...with corresponding sets I and O.

Fig. 14. A program with a corresponding AST describing the sets I and O. The program is safe because accesses to y in parallel have no intersections for I.

— Nodes that represent expressions, assignments, C calls, and declarations simply reproduce $O = I$, as they do not await;
— An `await E` statement, where `E` is an external input event, has $O = \{E\}$ (see also internal events below).
— A `break` statement has $O = \{\}$ as it escapes the innermost `loop` and never proceeds to the statement immediately following it (see also `loop` below);
— A *sequence node (;)* modifies each of its children to have $I_n = O_{n-1}$, except for $n = 1$ (which inherits I from the parent node). The set O for the whole node is copied from its last child, i.e., $O = O_n$.
— A `loop` node includes the output of its body on its own I ($I = I \cup O_{body}$), as the loop is also reached from its own body. The union of all O from nested `break` statements forms the set O for a `loop`.
— An `if` node has $O = O_{true} \cup O_{false}$, where `true` and `false` are the two `if` branches.
— A parallel composition may terminate from any of its branches, hence $O = O_1 \cup ... \cup O_n$.
— For internal events, an `await` awakes from any input that leads to any matching `emit` in a trail in parallel:
  — An `await e` has $O = I_{e1} \cup ... \cup I_{eN}$, where $e1...eN$ are `emit e` statements in trails in parallel.
  — An `emit e` terminates in the same reaction, having $O = I$.

With all sets calculated, we take all pairs of nodes that perform side effects and are in parallel branches, and compare their sets I for intersections. For each pair, if the intersection is not empty, we mark both nodes as suspicious.

The code in Figure 14.a has its corresponding AST and sets I and O in Figure 14.b. The assignments to y in parallel (ln. [a]:5,8) have an empty intersection of I (ln. [b]:6,9), hence, they do not conflict. Note that, although the accesses to y in sequence (ln. [a]:5,11) do have an intersection (ln. [b]:6,11), they are not in parallel branches and are also safe.

3.2. Static Memory Layout

CÉU promotes a fine-grained use of trails: it is common to use trails that await a single event and terminate. For this reason, CÉU does not allocate per-trail stacks; instead, all data resides in fixed memory slots—this is true for the program variables as well as for temporary values and runtime flags. Memory for trails in parallel must coexist, while statements in sequence can reuse it. Translating this idea to C is straightforward [Kasten and Römer 2005]: memory for blocks in sequence are packed into a `union`, while blocks in parallel are packed into a `struct`. CÉU reserves a single static block of memory to hold all memory slots, whose size is the maximum the program uses at

```
1  input int A, B, C;
2  do
3      var int a = await A;
4  end
5  do
6      var int b = await B;
7  end
8  par/and do
9      await B;
10 with
11     await C;
12 end
```

```
1  union {              // sequence
2      int a;           //    do_1
3      int b;           //    do_2
4      struct {         //    par/and
5          int _and_1: 1;
6          int _and_2: 1;
7      };
8  } MEM ;
9
10
11
12 .
```

[a] A program in CÉU...   [b] ...with corresponding memory layout

Fig. 15. A program with blocks in sequence and in parallel, with corresponding memory layout generated by the compiler.

any given time. A position in the memory may hold different data (with variable sizes) during runtime. As an example, Figure 15 shows a program with its corresponding memory layout. The do-end blocks and par/and in sequence (ln. [a]:2–4,5–7,8–12) are packed in a union (ln. [b]:2,3,4–7), given that their variables cannot be in scope at the same time, e.g., a and b can safely share the same memory slot. The example also illustrates the presence of runtime flags (ln. [b]:4–7) related to the par/and termination (ln. [a]:8–12), which also reside in reusable slots in the static memory.

3.3. Static and Lightweight Trail Allocation

Each line of execution in CÉU needs to carry associated data, such as which event it is currently awaiting and which code to execute when it awakes. The compiler statically infers the maximum number of trails a program can have at the same time and creates a static vector to hold the runtime information about them. Like normal variables, trails that cannot be active at the same time share slots in the static memory vector.

At any given moment, a trail can be awaiting in one of the following states: INACTIVE, STACKED, FINALIZE, or in any of the events defined in the program:

```
enum {
    INACTIVE = 0,
    STACKED,
    FINALIZE,
    EVT_A,       // input void A;
    EVT_e,       // event int e;
    <...>        // other events
}
```

All terminated or not-yet-started trails stay in the INACTIVE state and are ignored by the scheduler. A STACKED trail holds an associated numeric stack level and can only execute when scheduler runtime drops to that level. A FINALIZE trail represents a pending finalization block which is scheduled only when its corresponding block goes out of scope. A trail waiting for an event stays in that event, also holding the minimum sequence reaction number (*seqno*) in which it can awake (to respect *delayed awaits*). In concrete terms, a trail is represented by the following struct:

```
struct trail_t {
    state_t evt;              // awaiting event
    label_t lbl;              // awaking execution label
    union {
        unsigned char seqno;  // if evt=EVT_*
        stack_t       stk;    // if evt=STACKED
    };
};
```

The field evt holds the state of the trail (or the event it is awaiting); the field lbl holds the entry point in the code to execute when the trail segment is scheduled; the

third field depends on the `evt` field and may hold the `seqno` for an event, or the stack level `stk` for a `STACKED` trail.

The size of `state_t` depends on the number of events in an application; for an application with less than 253 events (plus the 3 states), one byte is enough. The size of `label_t` depends primarily on the number of `await` statements in the application—each `await` splits the code into two segments and requires a unique entry point in the code for its continuation. Additionally, split & join points for parallel compositions, `emit` continuations, and finalization blocks also require labels. The fields `seqno` requires only 2 bits because the scheduler adjusts them while traversing all trails. The size of `stack_t` depends on the maximum depth of nested emissions but is bounded by the maximum number of trails: in the worst case, a trail emits an event that awakes another trail, which emits an event that awakes another trail, and so on; the last trail cannot awake any other trail, because they are all blocked in the `STACKED` state.

In the context of embedded systems, the size of `trail_t` is typically only 3 bytes (1 byte for each field), imposing a negligible memory overhead even for trails that only await a single event and terminate. For instance, the *CTP* collection protocol ported to CÉU reaches eight simultaneous lines of execution but has a memory overhead of only 2% in comparison to the original single-threaded version in *C* [Sant'Anna et al. 2013].

3.4. Static Scheduling and Trail Finalization

In the final generated code in *C*, each trail segment label representing an entry point becomes a *switch case* with the associated code to execute. Figure 16 illustrates the generation process. For the program in [a], the compiler extracts the entry points and associated trails, e.g., the label `Awake_e` will execute on `TRAIL-0` (ln. [a]:7 and [b]:3). For each yielding statement (e.g., `emit`, `await`, `par/and`, etc.), the compiler splits the trail into two segments with associated entry points. The entry points translate to an `enum` in the generated code (ln. [b]:1–10). The state of trails translates to a vector of type `trail_t` with the maximum number of simultaneous trails (ln. [b]:12–15). On initialization, `TRAIL-0` is set to execute the `Main` entry point (ln. [b]:13), while all others are set to `INACTIVE` (ln. [b]:14).

The scheduler executes in two passes: in the *broadcast* pass, it sets all trails that are waiting for the current event to the state `STACKED` in the current numeric stack level; in the *dispatch* pass, it executes each trail that is `STACKED` to run in the current level, setting it immediately to `INACTIVE`.

During the dispatch pass, if a trail executes and emits an internal event, the scheduler increments the stack level and re-executes the two passes. After all trails are properly dispatched, the scheduler decrements the stack level and resumes the previous execution. For the boot reaction, the scheduler starts from the *dispatch* pass, given that the `Main` label is the only one that can be active at the stack level 0 (ln. [b]: 13).

The code in [c] dispatches a trail segment according to the current label to execute. For the first reaction, it executes the `Main` label in `TRAIL-0` (ln. 3–14). When the `Main` label reaches the `par/and` (ln. [a]:4), it first stacks `TRAIL-1` (ln. [c]:4–7) and then executes the `await e` (ln. [a]:6) in `TRAIL-0` (ln. [c]:9–14), respecting lexical execution order. The dispatcher sets the running `TRAIL-0` to await `EVT_e` on label `Awake_e`, and then halts with a `break`. Then, it switches to `TRAIL-1` and executes label `And_sub_2` (ln. [c]:6,16–21), which sets `TRAIL-1` to await `EVT_A` and also halts.

Regarding abortion and finalization, when a `par/or` terminates, the scheduler makes a *broadcast* pass for the `FINALIZE` event, but limited to the range of trails covered by the terminating `par/or`. Trails that do not match the `FINALIZE` are set to `INACTIVE`, as they have to be aborted. Given that trails in parallel are allocated in subsequent slots in the static vector `TRLS` (ln. [b]:12–15), this pass only aborts the desirable trails. The

```
input void A;                         1   enum {                           1   void dispatch (trail_t* t) {
event void e;                         2     Main = 1,   // ln 3          2     switch (t->lbl) {
// TRAIL 0 — lbl Main                 3     Awake_e,    // ln 7          3       case Main:
par/and do                            4     And_chk,    // ln 8,15       4         // activate TRAIL 1
  // TRAIL 0 — lbl Main               5     And_sub_2,  // ln 10         5         TRLS[1].evt = STACKED;
  await e;                            6     Awake_A_1,  // ln 12         6         TRLS[1].lbl = And_sub_2;
  // TRAIL 0 — lbl Awake_e            7     Emit_cont,  // ln 14         7         TRLS[1].stk = cur_stack;
  // TRAIL 0 — lbl And_chk            8     And_out,    // ln 17         8
with                                  9     Awake_A_2   // ln 19         9         // code in the 1st trail
  // TRAIL 1 — lbl And_sub_2         10   };                            10         // await e;
  await A;                           11                                 11         TRLS[0].evt = EVT_e;
  // TRAIL 1 — lbl Awake_A_1         12   trail_t TRLS[2] = {           12         TRLS[0].lbl = Awake_e;
  emit e;                            13     { STACKED,  Main, 0 };      13         TRLS[0].seq = cur_seqno;
  // TRAIL 1 — lbl Emit_cont         14     { INACTIVE, 0,    0 };      14         break;
  // TRAIL 1 — lbl And_chk           15   };                            15
end                                  16                                 16       case And_sub_2:
// TRAIL 0 — lbl And_out             17                                 17         // await A;
await A;                             18                                 18         TRLS[1].evt = EVT_A;
// TRAIL 0 — lbl Awake_A_2           19                                 19         TRLS[1].lbl = Awake_A_1;
                                     20                                 20         TRLS[1].seq = cur_seqno;
                                     21                                 21         break;
                                     22                                 22
                                     23                                 23       <...>  // other labels
                                     24                                 24     }
.                                    25   .                             25   }

         [a]                                  [b]                                  [c]
```

Fig. 16. [a] Static allocation of trails: the comments identify the trail indexes inferred by the compiler; [b] Entry-point labels: each trail segment has an associated numeric identifier generated by the compiler. [c] Dispatch function: uses a `switch` to associate each segment identifier with the corresponding code to execute.

subsequent *dispatch* pass executes the finalization code properly. Escaping a `loop` that contains parallel compositions also triggers the same abortion process.

3.5. Single-Threaded Dispatching

The implementation of CÉU dispatches active trails sequentially in a single thread, taking no advantage of multi-core CPUs. This decision comes not only from the fact that CÉU targets constrained single-CPU embedded systems, but also because CÉU imposes deterministic execution for intra-reaction statements.

Note that, as discussed in Section 2.3, the concurrency checks of CÉU infer precisely trails that are concurrent and yet do not share resources. Hence, these non-conflicting trails could potentially execute with real parallelism in multiple cores. However, our experiments with multi-core execution are actually slower than single-core execution in the same system. Considering that we use CÉU primarily in control-dominated applications, this result is not surprising and also appears in related work [Yuan et al. 2011; Haribi 2012]. One reason is the overhead from continuous fork-and-rejoin in small reactions. Another reason is contention from excessive locality of data in stackless trails sharing contiguous static memory.

If we consider data-intensive applications, multi-core implementations can offer considerable speedups. However, data-intensive computations do not typically require a disciplined step-wise execution and can actually execute in isolated asynchronous calls. Esterel provides a `task` primitive for this purpose [Berry 2000], while CÉU provides an equivalent `async/thread` primitive. Asynchronous execution is out of the scope of this paper.

Single-threaded dispatching may not be suitable for hard real-time activities. In previous work [Sant'Anna et al. 2013], we measure how synchronous lengthy computations in *C* (e.g., hashing and compression) can block the scheduler and affect higher-priority activities such as a radio driver. In such cases, the system requires careful

testing to avoid undersized hardware deployment. For instance, we currently do not perform any worst-case reaction times analysis [Boldt et al. 2008; Li et al. 2005].

3.6. Interaction with the Environment

As a reactive language, the execution of programs in CÉU is guided entirely by the occurrence of external input events. The binding for a specific platform (environment) calls hook functions in the API of the CÉU runtime whenever an external event occurs. These calls must never interleave or parallelize execution in order to preserve the sequential/discrete notion of time in CÉU.

Figure 17 shows our binding for *TinyOS* [Hill et al. 2000], which maps system callbacks to input events in CÉU. The file `ceu_app.h` (ln. 3) contains all definitions for the compiled CÉU program, which are further queried through `#ifdef`'s. The file `ceu_app.c` (ln. 4) contains the runtime of CÉU with the scheduler and dispatcher pointing to the labels defined in the program. The callback `Boot.booted` (ln. 6–11) is called by TinyOS on startup, so we initialize CÉU inside it (ln. 7). If the CÉU program uses timers, we also start a periodic timer (ln. 8–10) that triggers callback `Timer.fired` (ln. 13–17) every 10 milliseconds to advance the wall-clock time of CÉU (ln. 15)[5]. The remaining lines map pre-defined TinyOS events that can be used in CÉU programs, such as the light sensor (ln. 19–23) and the radio transceiver (ln. 25–36). The scheduler of TinyOS is already synchronous by default and always executes `event` handlers atomically, hence, the API calls to CÉU are properly serialized.

3.7. The Terra Virtual Machine

Terra is a system for programming wireless sensor network applications which uses CÉU as its scripting language [Branco et al. 2015]. Figure 18 shows the three basic elements of Terra: CÉU as the scripting language, a set of customized pre-built components, and the embedded virtual-machine engine which can disseminate and install bytecode images dynamically. This approach aims to combine the flexibility of remotely uploading code with the expressiveness and safety guarantees of CÉU.

The main difference between the standard *C* back end and the Terra *VM* is the *code generation phase*, which here outputs assembly instructions for the *VM* (instead of statements in *C*). To reduce the memory footprint of applications, the *VM* includes special instructions for complex and recurrent operations from the runtime of CÉU, such as for handling events and trails.

In Terra, CÉU scripts cannot execute arbitrary *C* code, instead, they rely on pre-built components that can be customized for different application domains. In the domain of sensor networks, Terra already provides components organized in four areas: radio communication, group management, data aggregation, and local operations (e.g., access to sensors and actuators). When creating an instance of the *VM*, the programmer can choose whether or not to include each component, setting different abstraction boundaries for scripts. The generated *VM* has to be preloaded into the embedded devices before they are physically distributed.

The communication between scripts in CÉU and the components in the *VM* is mostly through events: scripts `emit` requests through `output` events and `await` answers through `input` events. Terra also provides system calls for initialization and configuration of components (e.g., *getters* and *setters*). Figure 19.a shows a CÉU interface with the available functionality for a customized *VM* (with temperature and radio components). Figure 19.b shows the associated bindings for output events (ln. 1–8), input events (ln. 10–14), and system calls (ln. 16–22). Note that all applications for the customized

---

[5]We also offer a mechanism to start the underlying timer on demand to avoid the "battery unfriendly" 10ms polling.

```
1   implementation
2   {
3       #include "ceu_app.h"
4       #include "ceu_app.c"
5
6       event void Boot.booted () {
7           ceu_init();
8   #ifdef CEU_WCLOCKS
9           call Timer.startPeriodic(10);
10  #endif
11      }
12
13  #ifdef CEU_WCLOCKS
14      event void Timer.fired () {
15          ceu_wclock(10000);
16      }
17  #endif
18
19  #ifdef EVT_PHOTO_READDONE
20      event void Photo.readDone (int val) {
21          ceu_go(EVT_PHOTO_READDONE, &val);
22      }
23  #endif
24
25  #ifdef EVT_RADIO_SENDDONE
26      event void RadioSend.sendDone (message_t* msg) {
27          ceu_go(EVT_RADIO_SENDDONE, &msg);
28      }
29  #endif
30
31  #ifdef EVT_RADIO_RECEIVE
32      event message_t* RadioReceive.receive (message_t* msg) {
33          ceu_go(EVT_RADIO_RECEIVE, &msg);
34          return msg;
35      }
36  #endif
37
38      <...>   // other events
39  }
```

Fig. 17. The *TinyOS* binding for CÉU. This platform-dependent template includes the *C* files generated from the original application in CÉU (`ceu_app.h` and `ceu_app.c`) for the *code generation phase*.

Fig. 18. Terra programming system basic elements.

VM must comply with the same interface. In contrast, the template-based *C* back end (illustrated in Figure 17) allows applications to choose all possible combinations of functionalities from the underlying platform at compile time.

4. RELATED WORK

CÉU has a strong influence from Esterel but differ in the most fundamental aspect of the notion of time. In CÉU, instead of clock ticks, atomic external event occurrences that define time units. The event-driven approach of CÉU is widespread [Ousterhout 1996] and popular in many software communities, such as web frameworks (e.g., *jQuery* [Chaffer 2009] and *Node.js* [Tilkov and Vinoski 2010]), GUI toolkits (e.g., *Tcl/Tk* [Ousterhout 1991] and *Java Swing* [Eckstein et al. 1998]), and Games [Nystrom 2014]. We consider that this approach is more familiar to programmers and simplifies the reasoning about concurrency. As far as we know, CÉU is the first to encrust this event-triggered notion of time in the core of the language, which is a prerequisite for the concurrency checks that enable safe shared-memory concurrency. CÉU also the distinguishes external events from internal events. Internal events support stack-

```
// Output events                                      1   // Output events
output void REQUEST_TEMPERATURE;                      2   void VM.out(int evt_id, void* args) {
output int  REQUEST_SEND;    // sends int value       3       switch (id){
                                                      4           case O_REQUEST_TEMPERATURE:
// Input events                                       5               call TINYOS_TEMP.read();
input int  TEMPERATURE_DONE; // recvs int value       6           <...>; // O_REQUEST_SEND
input void SEND_DONE;                                 7       }
                                                      8   }
// System calls                                       9
function int getRadioID (void);                      10   // Input events
                                                     11   event TINYOS_TEMP.done (int val) {
                                                     12       VM.enqueue(I_TEMPERATURE_DONE, &val);
                                                     13   }
                                                     14   <...> // TINYOS_SEND.done
                                                     15
                                                     16   // System calls
                                                     17   void VM.function(int id, void* params) {
                                                     18       switch (id) {
                                                     19           case F_GET_RADIO_ID:
                                                     20               VM.push(TINYOS_NODE_ID);
                                                     21       }
.                                                    22   }

              [a]                                                   [b]
```

Fig. 19.   [a] CÉU interface with customized *VM*. [b] The routine VM.out redirects all output events to the corresponding OS calls (ln. 1–8). Each *TinyOS* event callback calls VM.enqueue for the corresponding input event (ln 10–14). System calls use VM.push for immediate return values (ln. 16–22).

based micro reactions within external reactions, providing more fine-grained control for intra-reaction execution.

Like CÉU, many synchronous languages rely on deterministic scheduling to preserve intra-reaction determinism (*Reactive C* [Boussinot 1991], *Protothreads* [Dunkels et al. 2006], *SOL* [Karpinski and Cahill 2007], *SC* [Von Hanxleden 2009], and *PRET-C* [Andalam et al. 2010]). CÉU also performs concurrency checks to detect trails that, when reordered, change the observable behavior of the program, i.e., trails that actually rely on deterministic scheduling.

Esterel+Delay [Bourke and Sowmya 2009] is a non-intrusive extension to Esterel to program in terms of physical time with delay statements. The global transformation relies on a *platform statement* the specifies the available timers in the system (e.g., sample-driven or interrupt-driven). It then expands delay <expr> statements into existing Esterel statements in a way that the desired semantics can be realised in the platform. CÉU makes physical time a special input event that feeds the runtime with an associated time to elapse which is decremented from all awaiting trails. The compensation scheme of CÉU guarantees that trails awake in the correct order and that errors are not propagated on subsequent awaits. Interrupt-driven timers are also supported with a hook callback that the runtime calls whenever the program awaits an earlier timer.

Regarding resource management, Esterel supports a finalization mechanism to unconditionally execute a series of statements on abortion. In addition, CÉU also tracks pointers representing resources that cross *C* boundaries and forces the programmer to provide associated finalizers.

ReactiveML [Mandel and Pouzet 2005] and *URBI* [Baillie 2005] extend the synchronous model with dynamic lines of execution. The implementations use coroutines or CPS transformations and rely on heap allocation and/or garbage collection, diverging from our goals regarding resource efficiency and static bounds for memory and execution time. We discuss dynamic abstractions in CÉU in previous work [Sant'Anna et al. 2015].

Esterel has different compilation back ends that synthesizes to software and also to hardware circuits [Dayaratne et al. 2005; Edwards 2003]. Among the software-based approaches, *SAXO–RT* [Closse et al. 2002] is the closest to our implementation with respect to trail allocation and scheduling: the compiler slices programs into "control points" (analogous to our "entry points") and rearranges them into a directed acyclic graph respecting the constructive semantics of Esterel. Then, it flattens the graph into sequential code in *C* suitable for static scheduling.

TODO: compilao mais sofisticada, tem que lidar com data dependency

A number of virtual machines have been proposed for embedded systems. *Darjeeling* [Brouwers et al. 2008] and *TakaTuka* [Aslam et al. 2010] are complete *Java VMs* targeting constrained embedded systems with support for multithreading and garbage collection. Java has antagonistic design choices in comparison to CÉU: it does not impose static bounds on memory usage and execution time, and provides preemptive multithreading which requires synchronization primitives for accessing shared memory. Plummer et al. [Plummer et al. 2006] propose a Esterel-based *VM* with similar design choices to our work. To reduce code size, the *VM* has a specialized instruction set to deal with events and concurrency constructs that are particular to Esterel. However, the proposed *VM* is only a proof of concept, with no support for arithmetic operations, external system calls, or remote reprogramming.

5. CONCLUSION

We present the design and implementation of CÉU, a synchronous reactive language inspired by Esterel with simpler semantics and more fine-grained control for intra-reaction execution.

CÉU is a concurrency-safe language, employing static checks to ensure that the high degree of concurrency in embedded systems does not pose safety threats to applications. As a summary, the following safety properties hold for all programs that successfully compile in CÉU: time and memory-bounded reactions to the environment (except for system calls), no race conditions in shared memory, reliable abortion for activities handling resources, and automatic synchronization for timers. These properties are usually desirable in embedded applications and are guaranteed as preconditions in CÉU by design.

CÉU is a resource-efficient language suitable for constrained embedded systems. The reference implementation compiles to portable event-driven code in *C*, with no special requirements for OS threads or per-trail data stacks. The *VM* implementation uses the same front end and imposes no extra restrictions, being equally suitable for constrained systems.

CÉU is a practical language with expressive control constructs, such as lexically scoped parallel compositions, convenient first-class timers, and a unique stack-based mechanism for internal signalling. Programs interoperate seamlessly with *C*, and can take advantage of existing libraries, lowering the entry barrier for adoption. CÉU has an open source implementation and bindings for *TinyOS*, *Arduino*, and the *SDL* graphical library.[6]

For the past three years, we have been teaching CÉU to undergraduate and graduate students in courses on *distributed systems* and *reactive programming*. Our experience shows that students take advantage of the sequential-imperative style of CÉU and can implement non-trivial concurrent applications in a few weeks. More recently, a company specialized in embedded systems (not related to our research group) released a product based on CÉU.

---

[6]Website of CÉU: http://www.ceu-lang.org/

REFERENCES

Sidharta Andalam, Partha Roop, and Alain Girault. 2010. Predictable multithreading of embedded applications using PRET-C. In *Proceeding of MEMOCODE'10*. IEEE, 159–168.

Faisal Aslam, Luminous Fennell, Christian Schindelhauer, Peter Thiemann, Gidon Ernst, Elmar Haussmann, Stefan Rührup, and Zastash A. Uzmi. 2010. Optimized Java binary and virtual machine for tiny motes. In *Proceedings of DCOSS'10*. Springer, 15–30.

Jean-Christophe Baillie. 2005. Urbi: Towards a universal robotic low-level programming language. In *International Conference on Intelligent Robots and Systems*. IEEE, 820–825.

Albert Benveniste, Paul Caspi, Stephen A. Edwards, Nicolas Halbwachs, Paul Le Guernic, and Robert De Simone. 2003. The synchronous languages twelve years later. In *Proceedings of the IEEE*, Vol. 91. 64–83.

Gérard Berry. 1993. Preemption in Concurrent Systems.. In *FSTTCS (LNCS)*, Vol. 761. Springer, 72–93.

Gérard Berry. 1996. The Constructive Semantics of Pure Esterel. (1996).

Gérard Berry. 2000. *The Esterel-V5 Language Primer*. CMA and Inria, Sophia-Antipolis, France. Version 5.10, Release 2.0.

Gérard Berry, S Ramesh, and RK Shyamasundar. 1993. Communicating reactive processes. In *Proceedings of the 20th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*. ACM, 85–98.

Gérard Berry and Ellen Sentovich. 2001. Multiclock Esterel. In *Correct Hardware Design and Verification Methods*. Springer, 110–125.

Marian Boldt, Claus Traulsen, and Reinhard von Hanxleden. 2008. Worst case reaction time analysis of concurrent reactive programs. *Electronic Notes in Theoretical Computer Science* 203, 4 (2008), 65–79.

T. Bourke and A. Sowmya. 2009. Delays in Esterel. *SYNCHRON09* (2009), 55.

Frédéric Boussinot. 1991. Reactive C: An extension of C to program reactive systems. *Software: Practice and Experience* 21, 4 (1991), 401–428.

Frédéric Boussinot. 1998. SugarCubes implementation of causality. (1998).

Frédéric Boussinot and Robert De Simone. 1991. The Esterel language. *Proc. IEEE* 79, 9 (Sep 1991), 1293–1304.

Adriano Branco, Francisco Sant'anna, Roberto Ierusalimschy, Noemi Rodriguez, and Silvana Rossetto. 2015. Terra: Flexibility and Safety in Wireless Sensor Networks. *ACM Trans. Sen. Netw.* 11, 4, Article 59 (Sept. 2015), 27 pages. DOI:http://dx.doi.org/10.1145/2811267

Niels Brouwers, Peter Corke, and Koen Langendoen. 2008. Darjeeling, a Java compatible virtual machine for microcontrollers. In *Proceedings of the ACM/IFIP/USENIX Middleware'08 Conference Companion*. ACM, 18–23.

Jonathan Chaffer. 2009. *Learning JQuery 1.3: Better Interaction and Web Development with Simple JavaScript Techniques*. Packt Publishing Ltd.

Etienne Closse, Michel Poize, Jacques Pulou, Patrick Venier, and Daniel Weil. 2002. Saxo–RT: Interpreting Esterel semantic on a sequential execution structure. *Electronic Notes in Theoretical Computer Science* 65, 5 (2002), 80–94.

M. W. Sajeewa Dayaratne, Partha S. Roop, and Zoran Salcic. 2005. Direct Execution of Esterel Using Reactive Microprocessors. In *Proceedings of SLAP'05*.

Robert de Simone, Jean-Pierre Talpin, and Dumitru Potop-Butucaru. 2005. The Synchronous Hypothesis and Synchronous Languages. In *Embedded Systems Handbook*, R. Zurawski (Ed.).

Adam Dunkels, Oliver Schmidt, Thiemo Voigt, and Muneeb Ali. 2006. Protothreads: simplifying event-driven programming of memory-constrained embedded systems. In *Proceedings of SenSys'06*. ACM, 29–42.

Robert Eckstein, Marc Loy, and Dave Wood. 1998. *Java swing*. O'Reilly & Associates, Inc.

Stephen A. Edwards. 1999. Compiling Esterel into sequential code. In *7th International Workshop on Hardware/Software Codesign*. ACM, 147–151.

Stephen A. Edwards. 2003. Tutorial: Compiling concurrent languages for sequential processors. *ACM Transactions on Design Automation of Electronic Systems* 8, 2 (2003), 141–187.

Stephen A. Edwards. 2005. Using and Compiling Esterel. MEMOCODE'05 Tutorial. (July 2005).

David Gay, Philip Levis, Robert von Behren, Matt Welsh, Eric Brewer, and David Culler. 2003. The nesC Language: A Holistic Approach to Networked Embedded Systems. In *PLDI'03*. 1–11.

Wahbi Haribi. 2012. Compiling Esterel for Multi-Core Execution. *Synchrone Sprachen* (2012), 45.

Jason Hill, Robert Szewczyk, Alec Woo, Seth Hollar, David Culler, and Kristofer Pister. 2000. System architecture directions for networked sensors. *SIGPLAN Notices* 35 (November 2000), 93–104. Issue 11.

Marcin Karpinski and Vinny Cahill. 2007. High-Level Application Development is Realistic for Wireless Sensor Networks. In *Proceedings of SECON'07*. 610–619.

Oliver Kasten and Kay Römer. 2005. Beyond Event Handlers: Programming Wireless Sensors with Attributed State Machines. In *Proceedings of IPSN '05*. 45–52.

Xin Li, Jan Lukoschus, Marian Boldt, Michael Harder, and Reinhard Von Hanxleden. 2005. An Esterel processor with full preemption support and its worst case reaction time analysis. In *Proceedings of the 2005 international conference on Compilers, architectures and synthesis for embedded systems*. ACM, 225–236.

Louis Mandel and Marc Pouzet. 2005. ReactiveML: a reactive extension to ML. In *Proceedings of PPDP'05*. ACM, 82–93.

Ana Lúcia De Moura and Roberto Ierusalimschy. 2009. Revisiting coroutines. *ACM TOPLAS* 31, 2 (Feb. 2009), 6:1–6:31.

Robert Nystrom. 2014. *Game Programming Patterns*. Genever Benning.

John K. Ousterhout. 1991. An X11 Toolkit Based on the Tcl Language.. In *USENIX Winter*. 105–116.

John K. Ousterhout. 1996. Why Threads Are A Bad Idea (for most purposes). (January 1996).

Becky Plummer, Mukul Khajanchi, and Stephen A. Edwards. 2006. An Esterel virtual machine for embedded systems. In *International Workshop on Synchronous Languages, Applications, and Programming (SLAP'06)*. Citeseer, Vienna, Austria.

Partha S Roop, Zoran Salcic, and MW Dayaratne. 2004. Towards direct execution of Esterel programs on reactive processors. In *Proceedings of the 4th ACM international conference on Embedded software*. ACM, 240–248.

Francisco Sant'Anna. 2013. *Safe System-level Concurrency on Resource-Constrained Nodes with Céu*. Ph.D. Dissertation. PUC–Rio.

Francisco Sant'Anna, Noemi Rodriguez, and Roberto Ierusalimschy. 2012. *Céu: Embedded, Safe, and Reactive Programming*. Technical Report 12/12. PUC-Rio.

Francisco Sant'Anna, Noemi Rodriguez, and Roberto Ierusalimschy. 2013. Advanced Control Reactivity for Embedded Systems. Workshop on Reactivity, Events and Modularity (REM'13). (2013).

Francisco Sant'Anna, Noemi Rodriguez, and Roberto Ierusalimschy. 2015. Structured Synchronous Reactive Programming with Céu. In *Proceedings of Modularity'15*.

Francisco Sant'Anna, Noemi Rodriguez, Roberto Ierusalimschy, Olaf Landsiedel, and Philippas Tsigas. 2013. Safe System-level Concurrency on Resource-Constrained Nodes. In *Proceedings of SenSys'13*. ACM.

Klaus Schneider and Michael Wenz. 2001. A new method for compiling schizophrenic synchronous programs. In *Proceedings of the 2001 international conference on Compilers, architecture, and synthesis for embedded systems*. ACM, 49–58.

Ellen M Sentovich. 1997. Quick conservative causality analysis. In *System Synthesis, 1997. Proceedings., Tenth International Symposium on*. IEEE, 2–8.

Thomas R Shiple, Gerard Berry, and Hemé Touati. 1996. Constructive analysis of cyclic circuits. In *European Design and Test Conference, 1996. ED&TC 96. Proceedings*. IEEE, 328–333.

Olivier Tardieu and Robert De Simone. 2004. Curing schizophrenia by program rewriting in Esterel. In *Formal Methods and Models for Co-Design, 2004. MEMOCODE'04. Proceedings. Second ACM and IEEE International Conference on*. IEEE, 39–48.

Stefan Tilkov and Steve Vinoski. 2010. Node.js: Using JavaScript to build high-performance network programs. *IEEE Internet Computing* 6 (2010), 80–83.

Reinhard Von Hanxleden. 2009. SyncCharts in C: a proposal for light-weight, deterministic concurrency. In *Proceedings of the seventh ACM international conference on Embedded software*. ACM, 225–234.

Simon Yuan, Li Hsien Yoong, and Partha S Roop. 2011. Compiling Esterel for multi-core execution. In *Digital System Design (DSD), 2011 14th Euromicro Conference on*. IEEE, 727–735.

Jeong-Han Yun, Chul-Joo Kim, Seonggun Kim, Kwang-Moo Choe, and Taisook Han. 2013. Detection of harmful schizophrenic statements in Esterel. *ACM Transactions on Embedded Computing Systems (TECS)* 12, 3 (2013), 80.