

The Design, Semantics, and Implementation of the Synchronous Language CÉU

Francisco Sant'Anna, Departamento de Informática, PUC-Rio

Roberto Ierusalimsky, Departamento de Informática, PUC-Rio

Noemi Rodriguez, Departamento de Informática, PUC-Rio

Silvana Rossetto, Departamento de Ciência da Computação, UFRJ

Adriano Branco, Departamento de Informática, PUC-Rio

CÉU is a reactive language inspired by Esterel that targets constrained embedded platforms and ensures safe concurrency by handling threats at compile time. Based on the synchronous programming model, our design allows for a simple reasoning about concurrency that enables compile-time analysis and results in deterministic and memory-safe programs. We discuss the design of CÉU and propose a formal semantics for its particular control mechanisms, such as parallel compositions, finalization, and internal events. We also present two implementation back ends: one aiming for resource efficiency and interoperability with *C*, and another based on a virtual machine that allows remote reprogramming.

Additional Key Words and Phrases: Concurrency, Determinism, Embedded Systems, Esterel, Synchronous, Reactivity

1. INTRODUCTION

An established alternative to *C* in the field of embedded systems is the family of reactive synchronous languages [Benveniste et al. 2003]. Two major styles of synchronous languages have evolved: in the *control-imperative* style, programs are structured with control flow primitives, such as parallelism, repetition, and preemption; in the *dataflow-declarative* style, programs can be seen as graphs of values, in which a change to a value is propagated through its dependencies without explicit programming. Considering the control-based languages, Esterel [Boussinot and de Simone 1991] was the first to appear and succeed, influencing a number of embedded languages, such as *Reactive-C* [Boussinot 1991], *OSM* [Kasten and Römer 2005], *Sync-C* [Von Hanxleden 2009], and *PRET-C* [Andalam et al. 2010].

Despite its success and influence, Esterel has an overly complex semantics that requires ingenious static analysis to detect and refuse programs with *causality* and *schizophrenia* problems, with extensive coverage in the literature [Berry 1996; Edwards 2005; Schneider et al. 2004; Boussinot 1998; Schneider et al. 2006; Schneider and Wenz 2001; Sentovich 1997; Shiple et al. 1996; Tardieu and De Simone 2004]. The complex semantics not only challenges the analysis and compilation of programs, but also results in incompatible and non-compliant implementations. Furthermore, it also affects the capacity of the programmer's understanding, which, after all, has to solve the errors when facing corner cases. Another drawback of the Esterel semantics consists of loose and non-deterministic execution for intra-reaction statements, which prevents threads to share memory and interact with side-effect and stateful system calls appropriately. For instance, although Esterel supports orthogonal abortion of threads [Berry 1993], it doesn't offer a finalization mechanism to release resources in use.

In this work, we present CÉU, a new programming language that inherits the synchronous and imperative mindset of Esterel but diverges in fundamental semantic aspects. Overall, CÉU has a simple semantics with fine-grained execution control, and a straightforward implementation targeting resource-constrained systems. The list that follows summarizes the semantic peculiarities of our design:

- Stack-based execution for internal events, which provide a limited form of coroutines.

- A static temporal analysis and deterministic execution semantics that allows programs to safely share memory.
- A finalization mechanism for safe abortion of lines of execution holding external resources.
- First-class synchronized timers.

We discuss the design of CÉU and present a formal semantics for a small synchronous kernel that represents a subset of the language covering these new functionalities. We also present a lightweight implementation of CÉU with two back ends: one aiming for resource efficiency and interoperability with *C*, and another based on a virtual machine that allows remote reprogramming. Our implementations target resource-constrained devices, such as *Arduino* and *MICAz* sensor nodes based on 8-bit microcontrollers¹, showing that the peculiarities in the semantics of CÉU does not pose practical obstacles.

In previous work [Sant’Anna et al. 2013; Branco et al. 2015], we employed CÉU in the context of wireless sensor networks, developing a number of applications, protocols, and drivers. We evaluated the expressiveness of CÉU in comparison to event-driven code in *C* and attested a reduction in source code size (around 25%) with a small increase in memory usage (around 5–10% for *text* and *data*) [Sant’Anna et al. 2013]. For the *VM* back end, a simple application that blinks three LEDs periodically occupies less than 100 bytes and can be transmitted over the air in a couple of packets.

The rest of the paper is organized as follows: Section 2 discusses the design of CÉU, focusing on the fundamental differences to Esterel. Section 3 presents a formal semantics for the control primitives of CÉU. Section 4 presents the *C* and *VM* implementation back ends. Section 5 discusses other synchronous languages targeting embedded systems. Section 6 concludes the paper.

2. THE DESIGN OF CÉU

CÉU is a synchronous reactive language inspired by Esterel with support for multiple concurrent lines of execution known as *trails*. By reactive, we mean that programs are stimulated by the environment through input events, which are broadcast to all awaiting trails. By synchronous, we mean that trails at any given moment are either reacting to the current event or are awaiting another event; in other words, trails never react to different events simultaneously.

In the sections that follow, we discuss the main differences between CÉU and Esterel: queue-based external events and stack-based internal events (Section 2.1), shared-memory concurrency and determinism (Section 2.2), safe abortion with finalization (Section 2.3), and first-class timers (Section 2.4).

Regarding the similarities, Figure 1 shows side-by-side the implementations in Esterel (a) and CÉU (b) for the following control specification: “*Emit an output O as soon as two inputs A and B have occurred. Reset this behavior each time the input R occurs*” [Berry 2000]. The first phrase of the specification, awaiting and emitting the events, is translated almost identically in the two languages (ln. 3–8, in both implementations), given that Esterel’s ‘||’ and CÉU’s *par/and* constructs are equivalent. For the second phrase, the reset behavior, the Esterel version uses a *abort-when* (ln. 2–9), which, in this case, serves the same purpose of CÉU’s *par/or* (ln. 2–11): the occurrence of event *R* aborts the awaiting statements in parallel and restarts the enclosing loop.

¹Arduino: <https://www.arduino.cc/en/Main/arduinoBoardUno>

MICAz: http://www.memsic.com/userfiles/files/Datasheets/WSN/micaz_datasheet-t.pdf

(Both use the *ATmega328* microcontroller with 32 Kbytes of FLASH and 2 Kbytes of SRAM.)

<pre> 1 loop 2 abort 3 [4 await A 5 6 await B 7]; 8 emit O 9 when R 10 end 11 12 . </pre>	<pre> 1 loop do 2 par/or do 3 par/and do 4 await A; 5 with 6 await B; 7 end 8 emit O; 9 with 10 await R; 11 end 12 end </pre>
(a) Esterel	(b) CÉU

Fig. 1. A control specification implemented in Esterel and CÉU: “Emit *O* after *A* and *B*, resetting each *R*”

CÉU employs the synchronous model, in which programs advance in a sequence of discrete reactions to external events. It also has a strong imperative flavor, with explicit control flow through sequences, loops, and parallels, and also assignments. Being designed for control-intensive applications, CÉU provides support for concurrent lines of execution and broadcast communication through events. Internal computations within a reaction (e.g. expressions, assignments, and system calls) are considered to take no time in accordance with the synchronous hypothesis [Potop-Butucaru et al. 2005]. An `await` is the only statement that halts a running reaction and allows a program to advance in this notion of time. To ensure that reactions run in bounded time and programs always progress, loops are statically required to contain at least one `await` statement in all possible paths [Sant’Anna et al. 2013; Berry 2000]. CÉU shares the same limitations with Esterel and synchronous languages in general: computations that run in unbounded time (e.g., cryptography, image processing) do not fit the zero-delay hypothesis, and cannot be directly implemented.

2.1. Queue-Based External Events and Stack-Based Internal Events

Esterel makes no semantic distinctions between internal and external signals. In particular, programs can emit different external input signals simultaneously, with all coexisting during a reaction. In CÉU, a reaction starts from an external input event and programs cannot emit inputs at all. Therefore, the occurring input event is unique during the entire reaction, resulting in intrinsic queue-based handling. In contrast, programs can emit internal events but these follow a stack-based execution policy, similar to subroutine calls in typical programming languages. Figure 2 illustrates the use of internal signals (events) in Esterel (a) and CÉU (b). In Esterel, when *A* occurs, *B* is emitted (ln. 4–5) and both events become active, resulting in the invocation of `f()` and `g()` in no particular order. In CÉU, the occurrence of *A* makes the program behave as follows:

- (1) 1st trail awakes (ln. 4), emits *b*, and pauses.
- (2) 2nd trail awakes (ln. 8), calls `_g()`, and terminates.
- (3) 1st trail (on top of the stack) resumes, calls `_f()`, and terminates.
- (4) Both trails have terminated, so the `par/and` rejoins, and the program also terminates.

Internal events provide a fine-grained intra-reaction control mechanism. For instance, it brings a limited form of subroutines, as depicted in Figure 3. The subroutine `inc` is defined as a loop (ln. 3–6) that continuously awaits its identifying event (ln. 4), incrementing the value passed as reference (ln. 5). A trail in parallel (ln. 8–11) invokes the subroutine in reaction to event *A* through an `emit` (ln. 10). Given the stacked execution for internal events, the calling trail pauses, the subroutine awakes (ln. 4), runs its body (yielding `v=2`), loops, and awaits the next “call” (ln. 4, again). Only after this sequence the calling trail resumes and passes the assertion test (ln. 11).

1	<code>input A; // external</code>	1	<code>input void A; // external (in uppercase)</code>
2	<code>signal B; // internal</code>	2	<code>event void b; // internal (in lowercase)</code>
3	<code>[[</code>	3	<code>par/and do</code>
4	<code> await A;</code>	4	<code> await A;</code>
5	<code> emit B;</code>	5	<code> emit b;</code>
6	<code> call f();</code>	6	<code> -f();</code>
7	<code> </code>	7	<code>with</code>
8	<code> await B;</code>	8	<code> await b;</code>
9	<code> call g();</code>	9	<code> -g();</code>
10	<code>]]</code>	10	<code>end</code>

(a) Esterel

(b) CÉU

Fig. 2. Internal signals (events) in Esterel and CÉU: similar syntax, but different semantics.

On the one hand, this form of subroutines has a significant limitation that it cannot express recursive calls: an `emit` to itself is always ignored, given that a running body cannot be awaiting itself. On the other hand, this very same limitation brings some important safety properties to subroutines: first, they are guaranteed to react in bounded time; second, memory for locals is also bounded, not requiring data stacks. Also, this form of subroutines can use the other primitives of CÉU, such as parallel compositions and the `await` statement. In particular, they await keeping context information such as locals and the program counter, similarly to coroutines [Moura and Ierusalimsky 2009].

Another distinction regarding event handling in comparison to CÉU is that Esterel supports same-cycle bi-directional communication [Edwards 1999], i.e., two threads can react to one another during the same cycle due to mutual signal dependency. CÉU has a different take, posing a tradeoff that an `await` is only valid for the next reaction, i.e., if an `await` and `emit` occur simultaneously in parallel trails, the `await` does not awake. These *delayed awaits* avoid corner cases of instantaneous termination and re-execution of statements in the same reaction (known as *schizophrenic statements* [Tardieu and De Simone 2003; Yun et al. 2013]).

The example in Figure 4 illustrates delayed awaits, which prevents infinite execution by design. Both sides of the `par/or` have an `await` statement to avoid instantaneous termination (ln. 4,7). However, if the `emit` (ln. 6) could awake the `await` (ln. 4) in the same reaction that reaches them, the `par/or` would terminate and restart the loop instantaneously, resulting in infinite execution.

In atypical scenarios requiring immediate awake, delayed awaits can be circumvented by manually copying or transforming the code to execute on awake. For instance, sometimes we need to execute a block of code immediately, and then, periodically from internal event requests, as illustrated in Figure 5. In this case, the `await` moved to the end of the loop (ln. 10) makes the periodic code to also execute immediately (ln. 9), and then in reactions to each `emit` request (ln. 5). If the periodic `emit`

```

1 event int* inc; // subroutine 'inc'
2 par/or do
3   loop do // definitions are loops
4     var int* p = await inc;
5     *p = *p + 1;
6   end
7 with
8   var int v = 1;
9   await A;
10  emit inc => &v; // call 'inc'
11  _assert(v==2); // after return
12 end

```

Fig. 3. Subroutine `inc` is defined in a loop (ln. 3–6), in parallel with the caller (ln. 8–11).

```

1 event void e,f;
2 loop do
3   par/or do
4     await e;
5     with
6       emit e;      // w/o delayed awaits, the emit awakes 1st trail
7       await f;     // and restarts the loop instantaneously
8     end
9   end
end

```

Fig. 4. Delayed awaits prevents re-execution of statements by design.

```

1 event void e;
2 par do
3   loop do
4     <...>      // code that awaits some period
5     emit e;    // periodic request
6   end
7 with
8   loop do
9     <...>      // code to execute immediately and then periodically
10    await e;    // await after
11  end
12 end

```

Fig. 5. An example that circumvents the *delayed await* restriction by post-fixing the await inside the loop.

depends on a condition, then the code transformation becomes more intricate, requiring an extra condition test around the periodic code to prevent its immediate execution. On the one hand, we transfer the burden of dealing with these corner cases to the programmer. On the other hand, we simplify the semantics of the language and eliminate the need for analysis to deal with schizophrenic statements.

2.2. Shared-Memory Concurrency and Determinism

Embedded applications make extensive use of global memory and shared resources, such as through memory-mapped registers and system calls to drivers. Hence, an important goal of CÉU is to ensure a reliable behavior for programs with concurrent lines of execution sharing memory and interacting with the environment.

Esterel is only deterministic with respect to reactive control: “the same sequence of inputs always produces the same sequence of outputs” [Berry 2000]. However, the execution order for operations within a reaction is non-deterministic: “if there is no control dependency, as in `<<call f1() || call f2()>>`, the order is unspecified and it would be an error to rely on it” [Berry 2000]. For this reason, Esterel assumes that external calls to the host language do not perform side effects. Also, Esterel syntactically forbids sharing memory between lines of execution to preserve determinism: “*if a variable is written by some thread, then it can neither be read nor be written by concurrent threads*” [Berry 2000].

Concurrency in CÉU is characterized when two or more trail segments in parallel execute during the same reaction. A trail segment is a sequence of statements followed by an await (or termination). In the example in Figure 6.a, the two assignments to `x` (ln. 5,8) can never run concurrently, because each trail segment reacts to different input events (ln. 4,7), which cannot occur simultaneously (according to Section 2.1). However, the example in Figure 6.b is non-deterministic, because the two assignments to `y` (ln. 5,8) occur in the same reaction to input `A` (ln. 4,7).

CÉU performs a temporal analysis at compile time and detects concurrent accesses to shared variables, as follows: if a variable is written in a trail segment, then a concur-

```

1 input void A, B;
2 var int x = 1;
3 par/and do
4   await A;
5   x = x + 1;
6 with
7   await B;
8   x = x * 2;
9 end

```

(a)

```

1 input void A;
2 var int y = 1;
3 par/and do
4   await A;
5   y = y + 1;
6 with
7   await A;
8   y = y * 2;
9 end

```

(b)

Fig. 6. Shared-memory concurrency in CÉU: Example (a) is safe because the trails access x atomically in different reactions; Example (b) is unsafe because both trails access y in the same reaction.

rent trail segment cannot read or write to that variable, nor dereference a pointer of that variable type. An analogous policy is applied for pointers *vs* variables and pointers *vs* pointers, as well as for system calls with side effects (e.g., `printf`).²

Regardless of the temporal analysis of CÉU, when multiple trails are active during the same reaction, they are scheduled in the order they appear in the program source code. Therefore, even though the example in Figure 6.b is suspicious, the assignments to y are both atomic and deterministic, i.e., after the reaction to A terminates, the value of y is 4 $((1+1)*2)$. On the one hand, enforcing an execution order for concurrent operations may seem arbitrary and also precludes true parallelism. On the other hand, it provides a priority scheme for trails, and makes shared-memory concurrency more tractable. For constrained embedded development, we believe that deterministic shared-memory concurrency is beneficial, given the extensive use of memory mapped ports for *I/O* and the lack of hardware support for real parallelism. Other synchronous embedded languages, such as *SOL* [Karpinski and Cahill 2007] and *PRET-C* [Andalam et al. 2010], made a similar design choice.

Figure 7 compares two syntactically equivalent code fragments in Esterel (a) and CÉU (b) to summarize the semantic differences regarding (non-)determinism. Even though the program in CÉU executes deterministically, the compiler still issues a warning, because an apparently innocuous reordering of trails modifies the semantics of the program. Note that in Esterel multiple external events can coexist in the same reaction, which disallows a similar temporal analysis.

2.3. Safe Abortion with Finalization

The introductory example in Figure 1 illustrates how synchronous languages can abort awaiting lines of execution without tweaking them with synchronization primitives. In contrast, traditional (asynchronous) multi-threaded languages cannot express thread termination safely [Berry 1993; ORACLE 2011]. Still, handling abortion when deal-

²CÉU assumes that all system calls perform side effects, unless they are annotated as “@pure”.

```

input A;
[
  await A;
  call f1();
||
  await A;
  call f2();
];

```

(a) Esterel

```

input void A;
par/and do
  await A;
  _f1();
with
  await A;
  _f2();
end

```

(b) CÉU

Fig. 7. In Esterel, the execution order between $f1$ and $f2$ is unspecified, whereas in CÉU, $_f1$ executes before $_f2$ due to deterministic scheduling based on lexical order.

<pre> 1 input void STOP, RETRANSMIT, SEND_ACK; 2 output _pkt.t* SEND_ENQUEUE, SEND_CANCEL; 3 par/or do 4 await STOP; 5 with 6 loop do 7 par/or do 8 await RETRANSMIT; 9 with 10 par/and do 11 await 1min; 12 with 13 var _pkt.t buffer; 14 <fill-buffer-info> 15 emit SEND_ENQUEUE => &buffer; 16 await SEND_ACK; 17 end 18 end 19 end 20 end </pre>	<pre> 12 <...> 13 var _pkt.t buffer; 14 <fill-buffer-info> 15 finalize 16 emit SEND_ENQUEUE => &buffer; 17 with 18 emit SEND_CANCEL => &buffer; 19 end 20 await SEND_ACK; 21 <...> 22 23 24 25 26 27 28 29 30 31 . </pre>
(a)	(b)

Fig. 8. (a) The unsafe network protocol does not compile. (b) The finalization clause extends the protocol to handle abortion properly.

ing with external resources is challenging because they are not subject to the same synchronous execution discipline.

To illustrate the risks related to abortion, consider the unsafe example of Figure 8.a, which does not compile in CÉU. It describes the state machine of the data collection protocol *CTP* for sensor networks [Gnawali et al. 2009]. The input and output events represent the external interface of the protocol (ln. 1–2). The protocol has to transmit a packet every minute (with `SEND_ENQUEUE`), unless it receives a `RETRANSMIT` request to immediately re-transmit it, or a `STOP` request to terminate. The protocol is implemented with two main trails: one simply monitors the stopping event (ln. 4); the other periodically transmits the packet (ln. 6–19). The periodic transmission is a loop that starts two other trails (ln. 7–18): one handles the immediate retransmission request (ln. 8); the other transmits the packet and waits for a confirmation (ln. 10–17). The actual transmission (ln. 13–16) is enclosed with a `par/and` that takes at least one minute before looping (ln. 11), in accordance with the specification. Note that the `emit SEND_ENQUEUE` (ln. 15) is *asynchronous*, handing to the radio driver a pointer to the lexically-scoped packet (ln. 13). The driver makes the transmission in the background, holding the packet until it signals the application with a `SEND_ACK` event to acknowledge completion (ln. 16). At any time, the client may request a retransmission (ln. 8), which terminates the `par/or` (ln. 7) and restarts the loop (ln. 6), but does not abort the ongoing transmission initiated with the `emit SEND_ENQUEUE` (ln. 15). The client may also request to stop the whole protocol at any time (ln. 4). Therefore, if the sending trail is aborted by the `STOP` or `RETRANSMIT` requests, the packet buffer goes out of scope (ln. 13), leaving behind a *dangling pointer* in the radio driver, which will possibly transmit corrupted data.

The unsafe example of Figure 8.a does not compile because CÉU tracks the interaction of `par/or` compositions with local variables and stateful output events calls in order to preserve safe abortion of trails. CÉU enforces a *finalization* clause to accompany the output request. The code in Figure 8.b properly cancels the packet transmission when the block of buffer goes out of scope, i.e., the finalization clause (ln. 18) executes automatically on external abortion.³

³Note that the compiler only enforces the programmer to write the finalization clause, but cannot check if it actually handles the resource properly.

```

var int v;
await 10ms;
v = 1;
await 1ms;
v = 2;
.

```

(a)

```

par/or do
  await 10ms;
  <...> // any non-awaiting sequence
  await 1ms;
  v = 1;
with
  await 12ms;
  v = 2;
end

```

(b)

Fig. 9. First-class timers in CÉU.

2.4. First-Class Timers

Activities that involve reactions to *wall-clock time*⁴ appear in typical patterns of embedded development, such as timeout watchdogs and sensor samplings. However, support for wall-clock time is somewhat low-level in existing languages, usually through timer callbacks or “sleep” blocking calls. Furthermore, in any concrete timer implementation, a requested timeout does not expire precisely without delays, a fact that is usually ignored in the development process. We define the difference between the requested timeout and the actual expiring time as the *residual delta time (delta)*. Without explicit manipulation, the recurrent use of timed activities in sequence (or in a loop) may accumulate a considerable amount of deltas that can lead to incorrect behavior in programs.

The `await` statement of CÉU supports wall-clock time and handles deltas automatically, resulting in more robust applications. For the example in Figure 9.a, suppose that after the first `await` request, the underlying system gets busy and takes 15ms to check for expiring awaits. The CÉU scheduler will notice that the `await 10ms` has not only already expired, but is delayed with `delta=5ms`. Then, the awaiting trail awakes, sets `v=1`, and invokes `await 1ms`. As the current delta is higher than the requested timeout (i.e. $5ms > 1ms$), the trail is rescheduled for execution, now with `delta=4ms`.

CÉU also takes into account the fact that time is a physical quantity that can be added and compared. For instance, for the example in Figure 9.b, although the scheduler cannot guarantee that the first trail terminates exactly in 11ms, it can at least ensure that the program always terminates with `v=1`. Given that any non-awaiting sequence is considered to take no time in the synchronous model, the first trail is guaranteed to terminate before the second trail, because $10 + 1 < 12$. A similar program in a language without first-class support for timers would depend on the execution timings for the code marked as `<...>`, making the reasoning about the execution behavior more difficult.

3. FORMAL SEMANTICS

In this section, we introduce a reduced syntax of CÉU and propose an operational semantics to formally describe the language. We describe a small synchronous kernel with broadcast communication highlighting the peculiarities of CÉU, in particular the stack-based execution for internal events. For the sake of simplicity, we focus on the control aspects of the language, leaving out side effects and system calls (which behave like in conventional imperative languages).

3.1. Abstract Syntax

Figure 10 shows the syntax for a subset of CÉU that is sufficient to describe all semantic peculiarities of the language. Except for *fin* and the expressions used internally by

⁴By wall-clock time we mean the passage of time from the real world, measured in hours, minutes, etc.

<code>p ::= mem(id)</code>	<code>// primary expressions</code>
<code> await(id)</code>	<code>(any memory access to 'id')</code>
<code> emit(id)</code>	<code>(await event 'id')</code>
<code> break</code>	<code>(emit event 'id')</code>
	<code>(loop escape)</code>
	<code>// compound expressions</code>
<code> if mem(id) then p else p</code>	<code>(conditional)</code>
<code> p ; p</code>	<code>(sequence)</code>
<code> loop p</code>	<code>(repetition)</code>
<code> p and p</code>	<code>(par/and)</code>
<code> p or p</code>	<code>(par/or)</code>
<code> fin p</code>	<code>(finalization)</code>
	<code>// derived by semantic rules</code>
<code> awaiting(id,n)</code>	<code>(awaiting 'id' since sequence number 'n')</code>
<code> emitting(n)</code>	<code>(emitting on stack level 'n')</code>
<code> p @ loop p</code>	<code>(unwinded loop)</code>
<code> nop</code>	<code>(terminated expression)</code>

Fig. 10. Reduced syntax of CÉU.

the semantics (i.e., *awaiting*, *emitting*, *p @ loop*; *p*, and *nop*), all other expressions are equivalent to their counterparts in the concrete language.

The *mem(id)* primitive represents all accesses, assignments, system calls, and output events that affect a memory location identified by *id*. According to the synchronous hypothesis of CÉU, *mem* expressions are considered to be atomic and instantaneous. As the challenging parts of CÉU reside on its control structures, we are not concerned here with a precise semantics for side effects, but only with their occurrences in programs. Note that *mem* and *await/emit* expressions do not share identifiers, i.e., an identifier is either a variable or an event.

3.2. Operational Semantics

The core of our semantics describes how a program reacts to a single external input event, i.e., starting from the input event, how the program behaves and becomes idle again to proceed to the subsequent reaction. We use a set of small-step operational rules, which are built in such a way that at most one transition is possible at any time, resulting in deterministic reactions. Each reaction is identified by a ever-increasing *n* that remains constant during the entire reaction. The transition rules map a program *p* and a stack of events *S* in a single step to a modified program and stack:

$$\langle S, p \rangle \xrightarrow[n]{} \langle S', p' \rangle \quad \textbf{(rule-inner)}$$

where

$$\begin{array}{ll} S, S' \in id^* & (\text{stack of event identifiers : } [id_{top}, \dots, id_{bottom}]) \\ p, p' \in P & (\text{program as described in Figure 10}) \\ n \in \mathbb{N} & (\text{unique identifier for the entire reaction}) \end{array}$$

At the beginning of a reaction, the stack is initialized with the occurring external event *ext* ($S = [ext]$), but *emit* expressions can push new events on top of it (we discuss how they are popped further). The sequence number *n*, which is incremented each reaction, prevents that awaiting expressions awake in the same reaction they are reached (the *delayed awaits* as explained in Section 2.1).

The transition rules for the primary expressions are as follows:

$$\langle S, \text{mem}(id) \rangle \xrightarrow[n]{} \langle S, \text{nop} \rangle \quad \textbf{(mem)}$$

$$\langle S, \text{await}(id) \rangle \xrightarrow[n]{} \langle S, \text{awaiting}(id, n+1) \rangle \quad \textbf{(await)}$$

$$\langle id : S, \text{awaiting}(id, m) \rangle \xrightarrow[n]{} \langle id : S, \text{nop} \rangle, \text{ if } m \leq n \quad \textbf{(awake)}$$

$$\langle S, \text{emit}(id) \rangle \xrightarrow[n]{} \langle id : S, \text{emitting}(|S|) \rangle \quad \textbf{(emit)}$$

$$\langle S, \text{emitting}(k) \rangle \xrightarrow[n]{} \langle S, \text{nop} \rangle, \text{ if } k = |S| \quad \textbf{(pop)}$$

A *mem* operation executes immediately and becomes a *nop* to indicate termination (rule **mem**). An *await* is transformed into an *awaiting* (rule **await**) as an artifice to remember the external sequence number $n+1$ it can awake: an *awaiting* can only transit to a *nop* (rule **awake**) if its referred event *id* matches the top of the stack and it was reached in a previous reaction (i.e., sequence number $m \leq n$). An *emit* transits to an *emitting* holding the current stack level ($|S|$ stands for the stack length), and pushing the referred event on the stack (rule **emit**). With the new stack level $|S|+1$, the *emitting*($|S|$) itself cannot transit, as rule **pop** expects its parameter *k* to match the current stack level. This trick provides the desired stack-based semantics for internal events.

Proceeding to compound expressions, the rules for conditionals and sequences are straightforward:

$$\frac{\text{val}(id, n) \neq 0}{\langle S, (\text{if mem}(id) \text{ then } p \text{ else } q) \rangle \xrightarrow[n]{} \langle S, p \rangle} \quad \textbf{(if-true)}$$

$$\frac{\text{val}(id, n) = 0}{\langle S, (\text{if mem}(id) \text{ then } p \text{ else } q) \rangle \xrightarrow[n]{} \langle S, q \rangle} \quad \textbf{(if-false)}$$

$$\frac{\langle S, p \rangle \xrightarrow[n]{} \langle S', p' \rangle}{\langle S, (p ; q) \rangle \xrightarrow[n]{} \langle S', (p' ; q) \rangle} \quad \textbf{(seq-adv)}$$

$$\langle S, (\text{nop} ; q) \rangle \xrightarrow[n]{} \langle S, q \rangle \quad \textbf{(seq-nop)}$$

$$\langle S, (\text{break} ; q) \rangle \xrightarrow[n]{} \langle S, \text{break} \rangle \quad \textbf{(seq-brk)}$$

Given that our semantics focuses on control, rules **if-true** and **if-false** are the only to query *mem* expressions. The “magical” function *val* receives a memory identifier and the current reaction sequence number, returning the current memory value. Although the value here is arbitrary, it is unique in a reaction, because a given expression can

execute only once within it (remember that *loops* must contain *awaits* which, from rule **await**, cannot awake in the same reaction they are reached).

The rules for loops are analogous to sequences, but use '@' as separators to properly bind breaks to their enclosing loops:

$$\langle S, (\text{loop } p) \rangle \xrightarrow{n} \langle S, (p @ \text{loop } p) \rangle \quad \textbf{(loop-expd)}$$

$$\frac{\langle S, p \rangle \xrightarrow{n} \langle S', p' \rangle}{\langle S, (p @ \text{loop } q) \rangle \xrightarrow{n} \langle S', (p' @ \text{loop } q) \rangle} \quad \textbf{(loop-adv)}$$

$$\langle S, (\text{nop} @ \text{loop } p) \rangle \xrightarrow{n} \langle S, \text{loop } p \rangle \quad \textbf{(loop-nop)}$$

$$\langle S, (\text{break} @ \text{loop } p) \rangle \xrightarrow{n} \langle S, \text{nop} \rangle \quad \textbf{(loop-brk)}$$

When a program encounters a *loop*, it first expands its body in sequence with itself (rule **loop-expd**). Rules **loop-adv** and **loop-nop** are similar to rules **seq-adv** and **seq-nop**, advancing the loop until they reach a *nop*. However, what follows the loop is the loop itself (rule **loop-nop**). Note that if we used ';' as a separator in loops, rules **loop-brk** and **seq-brk** would conflict. Rule **loop-brk** escapes the enclosing loop, transforming everything into a *nop*.

Proceeding to parallel compositions, the semantic rules for *and* and *or* always force transitions on their left branches *p* to occur before their right branches *q*:

$$\frac{\langle S, p \rangle \xrightarrow{n} \langle S', p' \rangle}{\langle S, (p \text{ and } q) \rangle \xrightarrow{n} \langle S', (p' \text{ and } q) \rangle} \quad \textbf{(and-adv1)}$$

$$\frac{\text{isBlocked}(n, S, p), \langle S, q \rangle \xrightarrow{n} \langle S', q' \rangle}{\langle S, (p \text{ and } q) \rangle \xrightarrow{n} \langle S', (p \text{ and } q') \rangle} \quad \textbf{(and-adv2)}$$

$$\frac{\langle S, p \rangle \xrightarrow{n} \langle S', p' \rangle}{\langle S, (p \text{ or } q) \rangle \xrightarrow{n} \langle S', (p' \text{ or } q) \rangle} \quad \textbf{(or-adv1)}$$

$$\frac{\text{isBlocked}(n, S, p), \langle S, q \rangle \xrightarrow{n} \langle S', q' \rangle}{\langle S, (p \text{ or } q) \rangle \xrightarrow{n} \langle S', (p \text{ or } q') \rangle} \quad \textbf{(or-adv2)}$$

The deterministic behavior of the semantics relies on the *isBlocked* predicate, which is defined in Figure 11 and used in rules **and-adv2** and **or-adv2**. These rules require the left branch *p* to be blocked in order to allow the right transition from *q* to *q'*. Basically, the *isBlocked* predicate determines that an expression becomes blocked when all of its trails in parallel hang in *awaiting* and *emitting* expressions that cannot advance.

For a parallel *and*, if one of the sides terminates, the composition is simply substituted by the other side (rules **and-nop1** and **and-nop2**, as follows). For a parallel *or*, if

$$\begin{aligned}
isBlocked(n, a : S, awaiting(b, m)) &= (a \neq b \vee m > n) \\
isBlocked(n, S, emitting(s)) &= (|S| \neq s) \\
isBlocked(n, S, (p ; q)) &= isBlocked(n, S, p) \\
isBlocked(n, S, (p @ loop q)) &= isBlocked(n, S, p) \\
isBlocked(n, S, (p and q)) &= isBlocked(n, S, p) \wedge isBlocked(n, S, q) \\
isBlocked(n, S, (p or q)) &= isBlocked(n, S, p) \wedge isBlocked(n, S, q) \\
isBlocked(n, S, _) &= false \quad (nop, await, \\
&\quad emit, break, if, loop)
\end{aligned}$$

Fig. 11. The recursive predicate *isBlocked* is true only if all branches in parallel are hanged in *awaiting* or *emitting* expressions that cannot transit.

one of the sides terminates, the whole composition terminates, also applying the *clear* function to properly finalize the aborted side (rules **or-nop1** and **or-nop2**):

$$\begin{aligned}
\langle S, (nop and q) \rangle &\xrightarrow{n} \langle S, q \rangle \quad \textbf{(and-nop1)} \\
\langle S, (p and nop) \rangle &\xrightarrow{n} \langle S, p \rangle \quad \textbf{(and-nop2)} \\
\langle S, (nop or q) \rangle &\xrightarrow{n} \langle S, clear(q) \rangle \quad \textbf{(or-nop1)} \\
\frac{isBlocked(n, S, p)}{\langle S, (p or nop) \rangle \xrightarrow{n} \langle S, clear(p) \rangle} &\quad \textbf{(or-nop2)}
\end{aligned}$$

The *clear* function, defined in Figure 12, concatenates all active *fin* bodies of the side being aborted, so that they execute before the composition rejoins. Note that there are no transition rules for *fin* expressions. This is because once reached, a *fin* expression halts and will only execute when it is aborted by a trail in parallel and is expanded by the *clear* function. In Section 3.3.3, we show how to map a finalization block in the concrete language to a *fin* in the formal semantics. Note also that there is a syntactic restriction that a *fin* body can only contain *mem* expressions in sequence, i.e., they are guaranteed to execute entirely within a reaction.

$$\begin{aligned}
clear(fin p) &= p \\
clear(p ; q) &= clear(p) \\
clear(p @ loop q) &= clear(p) \\
clear(p and q) &= clear(p) ; clear(q) \\
clear(p or q) &= clear(p) ; clear(q) \\
clear(_) &= nop
\end{aligned}$$

Fig. 12. The function *clear* extracts *fin* expressions in parallel and put their bodies in sequence.

Finally, a *break* in one of the sides in parallel escapes the closest enclosing *loop*, properly aborting the other side by applying the *clear* function:

$$\langle S, (break \text{ and } q) \rangle \xrightarrow[n]{} \langle S, (clear(q) ; break) \rangle \quad \textbf{(and-brk1)}$$

$$\frac{isBlocked(n, S, p)}{\langle S, (p \text{ and } break) \rangle \xrightarrow[n]{} \langle S, (clear(p) ; break) \rangle} \quad \textbf{(and-brk2)}$$

$$\langle S, (break \text{ or } q) \rangle \xrightarrow[n]{} \langle S, (clear(q) ; break) \rangle \quad \textbf{(or-brk1)}$$

$$\frac{isBlocked(n, S, p)}{\langle S, (p \text{ or } break) \rangle \xrightarrow[n]{} \langle S, (clear(p) ; break) \rangle} \quad \textbf{(or-brk2)}$$

A reaction eventually blocks in *awaiting* and *emitting* expressions in parallel trails. If all trails hangs only in *awaiting* expressions, it means that the program cannot advance in the current reaction. However, *emitting* expressions are pending in lower stack indexes and should eventually resume in the ongoing reaction (see rule **pop**). Therefore, we define another rule that behaves as **rule-inner** (presented above) if the program can advance, and, otherwise, pops the stack to resume the lower level:

$$\frac{\langle S, p \rangle \xrightarrow[n]{} \langle S', p' \rangle}{\langle S, p \rangle \xRightarrow[n]{} \langle S', p' \rangle} \quad \frac{isBlocked(n, s : S, p)}{\langle s : S, p \rangle \xRightarrow[n]{} \langle S, p \rangle} \quad \textbf{(rule-outer)}$$

To describe a *reaction* in CÉU, i.e., how a program behaves in reaction to a single external event, we use the reflexive transitive closure of **rule-outer**:

$$\langle S, p \rangle \xRightarrow[n]{*} \langle S', p' \rangle$$

Finally, to describe the complete execution of a program, we trigger multiple “invocations” of reaction chains in sequence:

$$\begin{aligned} \langle [e1], p \rangle &\xRightarrow[1]{*} \langle [], p' \rangle \\ \langle [e2], p' \rangle &\xRightarrow[2]{*} \langle [], p'' \rangle \\ \langle [e3], p'' \rangle &\xRightarrow[3]{*} \langle [], p''' \rangle \\ &\dots \end{aligned}$$

Each invocation starts with an external event at the top of the stack and finishes with a modified program and an empty stack. After each invocation, we increment the sequence number.

3.3. Concrete Language Mapping

Most statements from CÉU (“concrete CÉU”) map directly to those presented in the reduced syntax in Figure 10 (“abstract CÉU”). For instance, the *if* in the concrete language behaves exactly like the *if* in the formal semantics. However, there are some significant mismatches between the concrete and abstract CÉU, and we (informally) present appropriate mappings in this section. Again, we are not considering side-effects, which are all mapped to the *mem* semantic construct.




Fig. 13. Two-step translation from concrete to abstract emit and await expressions. The concrete code in (a) communicates the value 1 from the emit to the await. The abstract code in (c) uses a shared variable to hold the value.

3.3.1. *await and emit.* The concrete `await` and `emit` primitives support communication of values between them. In the two-step translation in Figure 13, we start with the concrete program in CÉU (a), which communicates the value 1 between the `emit` and `await` in parallel. In the intermediate translation (b), we include the shared variable `e_` to hold the value being communicated between the two trails in order to simplify the `emit`. Finally, we convert the program into the equivalent in the abstract syntax (c), translating side-effect statements into `mem` expressions. External events have a similar translation, i.e., each external event requires a corresponding variable that is explicitly set by the environment before each reaction.

3.3.2. *First-class Timers.* To encompass first-class timers, we introduce a special external event `DT` that is intercalated with each other event occurrence in an application (e.g. `e1`, `e2`):

$$\begin{aligned}
 \langle [DT], p \rangle &\xrightarrow[1]{*} \langle [], p' \rangle \\
 \langle [e1], p' \rangle &\xrightarrow[2]{*} \langle [], p'' \rangle \\
 \langle [DT], p'' \rangle &\xrightarrow[3]{*} \langle [], p''' \rangle \\
 \langle [e2], p''' \rangle &\xrightarrow[4]{*} \langle [], p'''' \rangle \\
 &\dots
 \end{aligned}$$

The event `DT` has an associated variable `DT_` carrying the wall-clock time elapsed between two occurrences in sequence, as depicted by the two-step translation in Figure 14. In the concrete program (a), the variable `dt` holds the residual delta time (as described in Section 2.4) after awaking from the timer. In the first step of the translation (b), we expand the `await 10ms` to a loop that decrements the elapsed number of microseconds for each occurrence of `DT`. When the variable `tot` reaches zero, we escape the loop setting the variable `dt` to contain the appropriate delta. In the last step (c), we convert the program to the abstract syntax.

3.3.3. *Finalization Blocks.* The biggest mismatch between concrete and abstract CÉU is regarding the `finalize` blocks, which require more complex modifications in the program for a proper mapping using `fin` expressions. In the three-step translation in Figure 15, we start with a concrete program (a) that uses a `finalize` to safely `release` the reference to `ptr` kept after the call to `hold`. In the translation, we first need to catch the outermost `do-end` termination to run the finalization code. For this, we translate the block into a `par/or` (b) with the original body in parallel with a `fin` expression to run the finalization code. Note that the `fin` has no transition rules in the seman-

<pre>dt = await 10ms; .</pre>	<pre>var int tot = 10000; // 10ms loop do await DT; tot = tot - DT; if tot <= 0 then dt = -tot; break; end end .</pre>	<pre>mem; loop (await (DT); mem; if mem then mem; break else nop)</pre>
(a)	(b)	(c)

Fig. 14. Two-step translation from concrete to abstract timer.

<pre>do var int* ptr = <...>; await A; finalize _hold(ptr); with _release(ptr); end await B; end .</pre>	<pre>par/or do var int* ptr = <...>; await A; _hold(ptr); await B; with { fin _release(ptr); } end .</pre>	<pre>f_ = 0; par/or do var int* ptr = <...>; await A; _hold(ptr); f_ = 1; await B; with { fin if f_ then _release(ptr); end } end .</pre>	<pre>mem; (mem; await (A); mem; mem; await (B); or fin if mem then mem else nop)</pre>
(a)	(b)	(c)	(d)

Fig. 15. Three-step translation from concrete to abstract finalization.

tics, keeping the *par/or* alive. This way, the *fin* body only executes when the *par/or* terminates either normally (after the *await B*), or aborted from an outer composition. However, the *fin* still (incorrectly) executes even if the call to *_hold* is not reached in the body due to an abort before awakening from the *await A*. To deal with this issue, for each *fin* we need a corresponding flag to keep track of code that needs to be finalized (c). The flag is initially set to false, avoiding the finalization code to execute. Only after the call to *_hold* that we set the flag to true and enable the *fin* body to execute. The complete translation substitutes the side-effect operations with *mem* expressions (d).

4. IMPLEMENTATION

The compilation process of a program in CÉU for the original *C* back end is composed of three main phases: the *parsing phase* converts the source code in CÉU to an *abstract syntax tree (AST)*; the *temporal analysis phase* detects inconsistencies in programs, such as unbounded loops and suspicious accesses to shared memory; the *code generation phase* converts the *AST* to standard *C* code and packs it with platform-dependent functionality (e.g., system calls) and the runtime of CÉU, compiling everything with *gcc* to generate the final binary.

In the sections that follow, we discuss the implementation insights related to the peculiarities of CÉU: the temporal analysis for determinism (Section 4.1), static memory allocation for data and trails (Sections 4.2 and 4.3), static scheduling and trail finalization (Section 4.4), interaction with the environment (Section 4.5), and the *VM* back end (Section 4.6),

<pre> 1 input void A, B; 2 var int y; 3 par/or do 4 await A; 5 y = 1; 6 with 7 await B; 8 y = 2; 9 end 10 await A; 11 y = 3; </pre>	<pre> 1 Stmts I={.} O={A} 2 Dcl_y I={.} O={.} 3 ParOr I={.} O={A,B} 4 Stmts I={.} O={A} 5 Await_A I={.} O={A} 6 Set_y I={A} O={A} 7 Stmts I={.} O={B} 8 Await_B I={.} O={B} 9 Set_y I={B} O={B} 10 Await_A I={A,B} O={A} 11 Set_y I={A} O={A} </pre>
(a)	(b)

Fig. 16. A program with a corresponding AST describing the sets I and O . The program is safe because accesses to y in parallel have no intersections for I .

4.1. Temporal Analysis for Shared-Memory Concurrency

The compile-time *temporal analysis* phase detects inconsistencies in CÉU programs. Here, we focus on the algorithm that detects suspicious access to shared variables, as discussed in Section 2.2.

For each node representing a statement in the program AST, we keep the set of events I (for *incoming*) that can lead to the execution of the node, and also the set of events O (for *outgoing*) that can terminate the node.

A node inherits the set I from its direct parent and calculates O according to its type:

- Nodes that represent expressions, assignments, C calls, and declarations simply reproduce $O = I$, as they do not await;
- An `await e` statement has $O = \{e\}$.
- A `break` statement has $O = \{\}$ as it escapes the innermost loop and never terminates, i.e., never proceeds to the statement immediately following it (see also `loop` below);
- A *sequence node* (`:`) modifies each of its children to have $I_n = O_{n-1}$. The first child inherits I from its parent node, and the set O for the sequence node is copied from its last child, i.e., $O = O_n$.
- A `loop` node includes its body's O on its own I ($I = I \cup O_{body}$), as the loop is also reached from its own body. The union of all `break` statements' O forms the set O for a loop.
- An `if` node has $O = O_{true} \cup O_{false}$.
- A parallel composition may terminate from any of its branches, hence $O = O_1 \cup \dots \cup O_n$.

With all sets calculated, we take all pairs of nodes that perform side effects and are in parallel branches, comparing their sets I for intersections. For each pair, if the intersection is not the empty set, we mark both nodes as suspicious.

The example in Figure 16.a has a corresponding AST, in Figure 16.b, with the sets I and O for each node. The event `.` (dot) represents the “boot” reaction. The assignments to y in parallel (ln. 5,8 in the code) have an empty intersection of I (ln. 6,9 in the AST), hence, they do not conflict. Note that although the accesses in ln. 5,11 in the code (ln. 6,11 in the AST) do have an intersection, they are not in parallel and are also safe.

4.2. Static Memory Layout

CÉU favors a fine-grained use of trails, being common to use trails that await a single event. For this reason, CÉU does not allocate per-trail stacks; instead, all data resides in fixed memory slots—this is true for the program variables as well as for temporary values and runtime flags. Memory for trails in parallel must coexist, while statements in sequence can reuse it. Translating this idea to C is straightforward [Kasten and


```

input int A, B, C;
do
  var int a = await A;
end
do
  var int b = await B;
end
par/and do
  await B;
with
  await C;
end

union {           // sequence
  int a_1;        // do_1
  int b_2;        // do_2
  struct {        // par/and
    int _and_3: 1;
    int _and_4: 1;
  };
} MEM ;

```

Fig. 17. A program with blocks in sequence and in parallel, with corresponding memory layout generated by the compiler.

Römer 2005; Bernauer and Römer 2013]: memory for blocks in sequence are packed in a struct, while blocks in parallel, in a union. CÉU reserves a single static block of memory to hold all memory slots, whose size is the maximum the program uses at a given time. A given position in the memory may hold different data (with variable sizes) during runtime. As an example, Figure 17 shows a program with corresponding memory layout. Each variable is assigned a unique *id* (e.g. *a_1*) so that variables with the same name can be distinguished. The *do-end* blocks in sequence are packed in a union, given that their variables cannot be in scope at the same time, e.g., *MEM.a_1* and *MEM.b_2* can safely share the same memory slot. The example also illustrates the presence of runtime flags related to the parallel composition, which also reside in reusable slots in the static memory.

4.3. Static and Lightweight Trail Allocation

Each line of execution in CÉU needs to carry associated data, such as which event it is awaiting and which code to execute when it awakes. The compiler statically infers the maximum number of trails a program can have at the same time and creates a static vector to hold the runtime information about them. Like normal variables, trails that cannot be active at the same time can share slots in the static memory vector.

At any given moment, a trail can be awaiting in one of the following states: *INACTIVE*, *STACKED*, *FINALIZE*, or in any of the events defined in the program:

```

enum {
  INACTIVE = 0,
  STACKED,
  FINALIZE,
  EVT_A,      // input void A;
  EVT_e,      // event int e;
  <...>       // other events
}

```

All terminated or not-yet-started trails stay in the *INACTIVE* state and are ignored by the scheduler. A *STACKED* trail holds an associated stack level and is delayed until the scheduler runtime reaches that level again. A *FINALIZE* trail represents a hanged finalization block which is only scheduled when its corresponding block goes out of scope. A trail waiting for an event stays in the state of the corresponding event, also holding the minimum sequence number (*seqno*) in which it can awake. In concrete terms, a trail is represented by the following struct:

```

struct trail_t {
  state_t evt;
  label_t lbl;
  union {
    unsigned char seqno;
    stack_t stk;
  };
};

```

```
};
};
```

The field `evt` holds the state of the trail (or the event it is awaiting); the field `lbl` holds the entry point in the code to execute when the trail segment is scheduled; the third field depends on the `evt` field and may hold the `seqno` for an event, or the stack level `stk` for a `STACKED` state.

The size of `state_t` depends on the number of events in the application; for an application with less than 253 events (plus the 3 states), one byte is enough. The size of `label_t` depends primarily on the number of `await` statements in the application—each `await` splits the code in two segments and requires a unique entry point in the code for its continuation. Additionally, split & join points for parallel compositions, `emit` continuations, and finalization blocks also require labels. The `seqno` could eventually overflow during execution (i.e., every 256 reactions). However, given that the scheduler traverses all trails on every reaction, it can adjust them to properly handle overflows (actually, 2 bits to hold the `seqno` is already enough). The size of `stack_t` depends on the maximum depth of nested emissions and is bounded to the maximum number of trails. In the worst case, a trail emits an event that awakes another trail, which emits an event that awakes another trail, and so on. The last trail cannot awake any trail, because they are all hanged in the `STACKED` state.

In the context of embedded systems, the size of `trail_t` is typically only 3 bytes (1 byte for each field), imposing a negligible memory overhead even for trails that only await a single event and terminate. For instance, the *CTP* collection protocol ported to CÉU reaches eight simultaneous lines of execution with an overhead of 2% in comparison to the original version in *C* [Sant’Anna et al. 2013].

4.4. Static Scheduling and Trail Finalization

In the final generated code in *C*, each trail segment label representing an entry point becomes a *switch case* with the associated code to execute. Figure 18 illustrates the generation process. For the program in (a), the compiler extracts the entry points and associated trails, e.g., the label `Awake.e` will execute on `TRAIL-0` (ln. 7). For each statement that pauses (`emit` and `await`), resumes (`par/and`, `par/or`, and `finalize`), or aborts (`par/or` and `break`), the compiler splits the trail into segments with associated entry points. The entry points translate to an `enum` in the generated code (ln. 1–10, in (b)). The state of trails translates to a vector of type `trail_t` with the maximum number of simultaneous trails (ln. 12–15). On initialization, `TRAIL-0` is set to execute the `Main` entry point (ln. 13), while all others are set to `INACTIVE` (in the example, only one, in ln. 14).

The scheduler executes in two passes: in the *broadcast* pass, it sets all trails that are waiting for the current event to `STACKED` in the current stack level; in the *dispatch* pass, it executes each trail that is `STACKED` to run in the current level, setting it immediately to `INACTIVE` (the trail segment may reset it in sequence if it doesn’t terminate).

During the dispatch pass, if a trail executes and emits an internal event, the scheduler increments the stack level and re-executes the two passes. After all trails are properly dispatched, the scheduler decrements the stack level and resumes the previous execution. For the first reaction, the scheduler starts from the *dispatch* pass, given that the `Main` label is the only one that can be active at the stack level 0 (ln. 13, in Figure 18.b).

The code in Figure 18.c dispatches a trail segment according to the current label to execute. For the first reaction, it executes the `Main` label in `TRAIL-0`. When the `Main` label reaches the `par/and`, it stacks `TRAIL-1` (ln. 4–7) and proceeds to the code in `TRAIL-0` (ln. 9–14), respecting the deterministic execution order. The code sets the running `TRAIL-0` to await `EVT.e` on label `Awake.e`, and then halts with a `break`. The next iteration of *dispatch*

<pre> 1 input void A; 2 event void e; 3 // TRAIL 0 - lbl Main 4 par/and do 5 // TRAIL 0 - lbl Main 6 await e; 7 // TRAIL 0 - lbl Awake.e 8 // TRAIL 0 - lbl And.chk 9 with 10 // TRAIL 1 - lbl And.sub.2 11 await A; 12 // TRAIL 1 - lbl Awake.A.1 13 emit e; 14 // TRAIL 1 - lbl Emit.cont 15 // TRAIL 1 - lbl And.chk 16 end 17 // TRAIL 0 - lbl And.out 18 await A; 19 // TRAIL 0 - lbl Awake.A.2 20 21 22 23 24 25 </pre>	<pre> 1 enum { 2 Main = 1, // ln 3 3 Awake.e, // ln 7 4 And.chk, // ln 8,15 5 And.sub.2, // ln 10 6 Awake.A.1, // ln 12 7 Emit.cont, // ln 14 8 And.out, // ln 17 9 Awake.A.2 // ln 19 10 }; 11 12 trail.t TRLS[2] = { 13 { STACKED, Main, 0 }; 14 { INACTIVE, 0, 0 }; 15 }; 16 17 18 19 20 21 22 23 24 25 </pre>	<pre> 1 void dispatch (trail.t* t) { 2 switch (t->lbl) { 3 case Main: 4 // activate TRAIL 1 5 TRLS[1].evt = STACKED; 6 TRLS[1].lbl = And.sub.2; 7 TRLS[1].stk = cur.stack; 8 9 // code in the 1st trail 10 // await e; 11 TRLS[0].evt = EVT.e; 12 TRLS[0].lbl = Awake.e; 13 TRLS[0].seq = cur.seqno; 14 break; 15 16 case And.sub.2: 17 // await A; 18 TRLS[1].evt = EVT.A; 19 TRLS[1].lbl = Awake.A.1; 20 TRLS[1].seq = cur.seqno; 21 break; 22 23 <...> // other labels 24 } 25 } </pre>
(a)	(b)	(c)

Fig. 18. (a) Static allocation of trails: the comments identify the trail indexes inferred by the compiler; (b) Entry-point labels: each trail segment has an associated numeric identifier generated by the compiler. (c) Dispatch function: uses a switch to associate each segment identifier with the corresponding code to execute.

takes TRAIL-1 and executes its registered label `And.sub.2` (ln. 16–21), which sets TRAIL-1 to await `EVT.A` and also halts.

Regarding abortion and finalization, when a `par/or` terminates, the scheduler makes a *broadcast* pass for the `FINALIZE` event, but limited to the range of trails covered by the terminating `par/or`. Trails that do not match the `FINALIZE` are set to `INACTIVE`, as they have to be aborted. Given that trails in parallel are allocated in subsequent slots in the static vector `TRLS`, this pass only aborts the desirable trails. The subsequent *dispatch* pass executes the finalization code. Escaping a loop that contains parallel compositions also triggers the same abortion process.

4.5. Interaction with the Environment

As a reactive language, the execution of programs in CÉU is guided entirely by the occurrence of external input events. The binding for a specific platform (environment) is responsible for calling hook functions in the API of the runtime of CÉU whenever an external event occurs. However, the binding must never interleave or run multiple API calls in parallel. This would break the CÉU sequential/discrete semantics of time.

As an example, Figure 19 shows our binding for *TinyOS* [Hill et al. 2000], which maps system callbacks to input events in CÉU. The file `ceu_app.h` (ln. 3) contains all definitions for the compiled CÉU program, which are further queried through `#ifdef`'s. The file `ceu_app.c` (ln. 4) contains the runtime of CÉU with the scheduler and dispatcher pointing to the labels defined in the program. The callback `Boot.booted` (ln. 6–11) is called by *TinyOS* on startup, so we initialize CÉU inside it (ln. 7). If the CÉU program uses timers, we also start a periodic timer (ln. 8–10) that triggers callback `Timer.fired` (ln. 13–17) every 10 milliseconds and advances the wall-clock time of CÉU (ln. 15)⁵. The

⁵We also offer a mechanism to start the underlying timer on demand to avoid the “battery unfriendly” 10ms polling.

```

1 implementation
2 {
3     #include "ceu_app.h"
4     #include "ceu_app.c"
5
6     event void Boot.booted () {
7         ceu_init();
8     #ifdef CEU_WCLOCKS
9         call Timer.startPeriodic(10);
10    #endif
11    }
12
13    #ifdef CEU_WCLOCKS
14        event void Timer.fired () {
15            ceu_wclock(10000);
16        }
17    #endif
18
19    #ifdef EVT_PHOTO_READDONE
20        event void Photo.readDone (int val) {
21            ceu_go(EVT_PHOTO_READDONE, &val);
22        }
23    #endif
24
25    #ifdef EVT_RADIO_SENDDONE
26        event void RadioSend.sendDone (message_t* msg) {
27            ceu_go(EVT_RADIO_SENDDONE, &msg);
28        }
29    #endif
30
31    #ifdef EVT_RADIO_RECEIVE
32        event message_t* RadioReceive.receive (message_t* msg) {
33            ceu_go(EVT_RADIO_RECEIVE, &msg);
34            return msg;
35        }
36    #endif
37
38    <...>    // other events
39 }

```

Fig. 19. The *TinyOS* binding for CÉU. This platform-dependent template includes the *C* files generated from the original application in CÉU (*ceu_app.h* and *ceu_app.c*) for the *Final generation* phase of Figure ??.

remaining lines map pre-defined *TinyOS* events that can be used in CÉU programs, such as the light sensor (ln. 19–23) and the radio transceiver (ln. 25–36). The scheduler of the *TinyOS* is already synchronous by default and always execute event handlers atomically, hence, the API calls to CÉU are properly serialized.

4.6. The Terra Virtual Machine

Terra is a system for programming wireless sensor network applications which uses CÉU as its scripting language [Branco et al. 2015]. Figure 20 shows the three basic elements of Terra: CÉU as the scripting language, a set of customized pre-built components, and the embedded virtual-machine engine which can disseminate and install bytecode images dynamically. This approach aims to combine the flexibility of remotely uploading code with the expressiveness and safety guarantees of CÉU.

The main difference between the standard *C* back end and the Terra VM is the *code generation phase*, which here outputs assembly instructions for the VM, instead of statements in *C*. To reduce the memory footprint of applications, the VM includes special instructions for complex and recurrent operations from the runtime of CÉU, such as handling events and trails.

In Terra, CÉU scripts cannot execute arbitrary *C* code, instead, they rely on pre-built components that can be customized for different application domains. Consider-




Fig. 20. Terra programming system basic elements.

<pre> // Output events output void REQUEST_TEMPERATURE; output int REQUEST_SEND; // sends int value // Input events input int TEMPERATURE_DONE; // recvs int value input void SEND_DONE; // System calls function int getRadioID (void); </pre>	<pre> 1 // Output events 2 void VM.out(int evt_id, void* args) { 3 switch (id){ 4 case O_REQUEST_TEMPERATURE: 5 call TINYOS_TEMP.read(); 6 <...>; // O_REQUEST_SEND 7 } 8 } 9 10 // Input events 11 event TINYOS_TEMP.done (int val) { 12 VM.enqueue(I_TEMPERATURE_DONE, &val); 13 } 14 <...> // TINYOS_SEND.done 15 16 // System calls 17 void VM.function(int id, void* params) { 18 switch (id) { 19 case F_GET_RADIO_ID: 20 VM.push(TINYOS_NODE_ID); 21 } 22 } </pre>
(a)	(b)

Fig. 21. (a) Céu interface with customized VM. (b) The routine VM.out redirects all output events to the corresponding OS calls (ln. 1–8). Each *TinyOS* event callback calls VM.enqueue for the corresponding input event (ln 10–14). System calls use VM.push for immediate return values (ln. 16–22).

ing the domain of sensor networks, Terra already provides components organized in four areas: radio communication, group management, data aggregation, and local operations (e.g., access to sensors and actuators). When creating an instance of the VM, the programmer can choose whether or not to include each component, setting different abstraction boundaries for scripts. The generated VM has to be preloaded into the embedded devices before they are physically distributed.

The communication between scripts in Céu and the components in the VM is mostly through events: scripts emit requests through output events and await answers through input events. Terra also provides system calls for initialization and configuration of components (e.g., *getters* and *setters*). Figure 21.a shows a Céu interface with the available functionality for a customized VM (with temperature and radio components). Figure 21.b shows the associated bindings for output events (ln. 1–8), input events (ln. 10–14), and system calls (ln. 16–22). Note that all applications for the customized VM must comply with the same interface. In contrast, the template-based C back end (illustrated in Figure 19) allows applications to choose all possible combinations of functionalities from the underlying platform at compile time.

5. RELATED WORK

CÉU has a strong influence from Esterel and embraces the disciplined synchronous-reactive model with support for lexical composition of lines of execution. However, there are fundamental semantic differences that prevents the design of CÉU as pure extensions to Esterel. In particular, Esterel has a notion of time similar to digital circuits in which multiple signals can be active at a clock tick. In fact, Esterel is also used in hardware design. In CÉU, instead of clock ticks, the occurrence of a single external event that defines a time unit. CÉU also distinguishes external events from stack-based internal events, which provide a limited form of coroutines supporting reactive statements.

The event-driven approach of CÉU is widespread [Ousterhout 1996] and popular in many software communities, such as web frameworks (e.g., *jQuery* [Chaffer 2009] and *Node.js* [Tilkov and Vinoski 2010]), GUI toolkits (e.g., *Tcl/Tk* [Ousterhout 1991] and *Java Swing* [Eckstein et al. 1998]), and Games [Nystrom 2014]. Like CÉU, event-driven programming is essentially synchronous, i.e., events go through a queue and are dispatched sequentially and atomically to prevent race conditions. We believe that for software design, this approach is more familiar to programmers and simplifies the reasoning about concurrency. For instance, the uniqueness of external events in CÉU is a prerequisite for the temporal analysis that enables safe shared-memory concurrency.

Many synchronous languages have been designed to interoperate with C, such as *Reactive C* [Boussinot 1991], Protothreads [Dunkels et al. 2006], *PRET-C* [Andalam et al. 2010] and *Sync-C* [Von Hanxleden 2009]. They offer Esterel-like parallel compositions with communication via shared variables, relying on deterministic scheduling to preserve determinism. However, it is the responsibility of the programmer to specify the execution order for threads, based on either explicit priorities, or source code lexical order (similar to CÉU). These languages have a tick-based notion of time similar to Esterel, which prevents the event-based temporal analysis of CÉU.

URBI [Baillie 2005] is a reactive scripting language with a rich set of control constructs for time management, event-driven communication, and concurrency. Concurrency is based on stackful coroutines, diverging from our goals regarding resource efficiency and static bounds for memory and execution time.

Esterel has different compilation back ends that synthesizes to software and also to hardware circuits [Dayaratne et al. 2005; Edwards 2003]. Among the software-based approaches, *SAXO-RT* [Closse et al. 2002] is the closest to our implementation with respect to trail allocation and scheduling: the compiler slices programs into “control points” (analogous to our “entry points”) and rearranges them into a directed acyclic graph respecting the constructive semantics of Esterel. Then, it flattens the graph into sequential code in C suitable for static scheduling.

A number of virtual machines have been proposed for embedded systems. The *Sun SPOT* platform with the *Squawk JVM* brings Java for the embedded domain [Simon et al. 2006], but requires a much powerful hardware.⁶ *Darjeeling* [Brouwers et al. 2008] and *TakaTuka* [Aslam et al. 2010] are complete *JVMs* targeting constrained embedded systems with support for multithreading and garbage collection. Java has antagonistic design choices in comparison to CÉU: it does not impose static bounds on memory usage and execution time, and provides preemptive multithreading with synchronization primitives for accessing shared memory. Plummer et al. [Plummer et al. 2006] propose a Esterel-based *VM* with similar design choices to our work. To reduce code size, the *VM* has a specialized instruction set to deal with events and concurrency constructs that are particular to Esterel. However, the proposed *VM* is only a proof of concept,

⁶The *Sun SPOT* uses a 32-bit CPU with 4 Mbytes of FLASH and 512 KBytes of SRAM.

with no support for arithmetic operations, external system calls, or remote reprogramming.

6. CONCLUSION

We presented the design, semantics, and implementation of CÉU, a synchronous reactive language inspired by Esterel targeting constrained embedded systems.

CÉU is a concurrency-safe language, employing a static analysis that encompass all control constructs and ensures that the high degree of concurrency in embedded systems does not pose safety threats to applications. As a summary, the following safety properties hold for all programs that successfully compile in CÉU: time and memory-bounded reactions to the environment (except for external system calls), no race conditions in shared memory, reliable abortion for activities handling resources, and automatic synchronization for timers. These properties are usually desirable in embedded applications and are guaranteed as preconditions in CÉU by design.

CÉU is a resource-efficient language suitable for constrained embedded systems. The reference implementation compiles to portable event-driven code in C, with no special requirements for OS threads or per-trail data stacks. The VM implementation uses the same front end and imposes no extra restrictions, being equally suitable for constrained systems.

CÉU is a practical language with expressive control constructs, such as lexically scoped parallel compositions, convenient first-class timers, and a unique stack-based signaling mechanism. Programs interoperate seamlessly with C, and can take advantage of existing libraries, lowering the entry barrier for adoption. CÉU has an open source implementation and bindings for *TinyOS*, *Arduino*, and the *SDL* graphical library.⁷

For the past three years, we have been teaching CÉU for undergraduate and graduate students in research projects and two hands-on courses on *distributed systems* and *reactive programming*. Our experience shows that students take advantage of the sequential-imperative style of CÉU and can implement non-trivial concurrent applications in a few weeks.

REFERENCES

- Sidharta Andalam and others. 2010. Predictable multithreading of embedded applications using PRET-C. In *Proceeding of MEMOCODE'10*. IEEE, 159–168.
- Faisal Aslam and others. 2010. Optimized java binary and virtual machine for tiny motes. In *Distributed Computing in Sensor Systems*. Springer, 15–30.
- Jean-Christophe Baillie. 2005. Urbi: Towards a universal robotic low-level programming language. In *International Conference on Intelligent Robots and Systems*. IEEE, 820–825.
- Albert Benveniste and others. 2003. The synchronous languages twelve years later. In *Proceedings of the IEEE*, Vol. 91. 64–83.
- Alexander Bernauer and Kay Römer. 2013. A Comprehensive Compiler-Assisted Thread Abstraction for Resource-Constrained Systems. In *Proceedings of IPSN'13*. Philadelphia, USA.
- Gérard Berry. 1993. Preemption in Concurrent Systems.. In *FSTTCS (LNCS)*, Vol. 761. Springer, 72–93.
- G. Berry. 1996. The Constructive Semantics of Pure Esterel. (1996).
- G. Berry. 2000. *The Esterel-V5 Language Primer*. CMA and Inria, Sophia-Antipolis, France. Version 5.10, Release 2.0.
- Frédéric Boussinot. 1991. Reactive C: An extension of C to program reactive systems. *Software: Practice and Experience* 21, 4 (1991), 401–428.
- Frédéric Boussinot. 1998. SugarCubes implementation of causality. (1998).
- F. Boussinot and R. de Simone. 1991. The Esterel language. *Proc. IEEE* 79, 9 (Sep 1991), 1293–1304.

⁷Website of CÉU: <http://www.ceu-lang.org/>

- Adriano Branco, Francisco Sant'anna, Roberto Ierusalimsky, Noemi Rodriguez, and Silvana Rossetto. 2015. Terra: Flexibility and Safety in Wireless Sensor Networks. *ACM Trans. Sen. Netw.* 11, 4, Article 59 (Sept. 2015), 27 pages. DOI: <http://dx.doi.org/10.1145/2811267>
- Niels Brouwers, Peter Corke, and Koen Langendoen. 2008. Darjeeling, a Java compatible virtual machine for microcontrollers. In *Proceedings of the ACM/IFIP/USENIX Middleware'08 Conference Companion*. ACM, 18–23.
- Jonathan Chaffer. 2009. *Learning JQuery 1.3: Better Interaction and Web Development with Simple JavaScript Techniques*. Packt Publishing Ltd.
- Etienne Closse and others. 2002. Saxo-RT: Interpreting esterel semantic on a sequential execution structure. *Electronic Notes in Theoretical Computer Science* 65, 5 (2002), 80–94.
- Sajeewa Dayaratne and others. 2005. Direct Execution of Esterel Using Reactive Microprocessors. In *Proceedings of SLAP'05*.
- Dunkels and others. 2006. Protothreads: simplifying event-driven programming of memory-constrained embedded systems. In *Proceedings of SenSys'06*. ACM, 29–42.
- Robert Eckstein, Marc Loy, and Dave Wood. 1998. *Java swing*. O'Reilly & Associates, Inc.
- Stephen A Edwards. 1999. Compiling Esterel into sequential code. In *7th International Workshop on Hardware/Software Codesign*. ACM, 147–151.
- Stephen A Edwards. 2003. Tutorial: Compiling concurrent languages for sequential processors. *ACM Transactions on Design Automation of Electronic Systems* 8, 2 (2003), 141–187.
- Stephen A. Edwards. 2005. Using and Compiling Esterel. MEMOCODE'05 Tutorial. (July 2005).
- Omrakash Gnawali and others. 2009. Collection tree protocol. In *Proceedings of SenSys'09*. ACM, 1–14.
- Hill and others. 2000. System architecture directions for networked sensors. *SIGPLAN Notices* 35 (November 2000), 93–104. Issue 11.
- Marcin Karpinski and Vinny Cahill. 2007. High-Level Application Development is Realistic for Wireless Sensor Networks. In *Proceedings of SECON'07*. 610–619.
- Oliver Kasten and Kay Römer. 2005. Beyond Event Handlers: Programming Wireless Sensors with Attributed State Machines. In *Proceedings of IPSN '05*. 45–52.
- Ana Lúcia De Moura and Roberto Ierusalimsky. 2009. Revisiting coroutines. *ACM TOPLAS* 31, 2 (Feb. 2009), 6:1–6:31.
- Robert Nystrom. 2014. *Game Programming Patterns*. Genever Benning.
- ORACLE. 2011. Java Thread Primitive Deprecation. <http://docs.oracle.com/javase/6/docs/technotes/guides/concurrency/threadPrimitiveDeprecation.html> (accessed in Aug-2014). (2011).
- John Ousterhout. 1996. Why Threads Are A Bad Idea (for most purposes). (January 1996).
- John K Ousterhout. 1991. An X11 Toolkit Based on the Tcl Language.. In *USENIX Winter*. 105–116.
- Becky Plummer, Mukul Khajanchi, and Stephen A Edwards. 2006. An Esterel virtual machine for embedded systems. In *International Workshop on Synchronous Languages, Applications, and Programming (SLAP'06)*. Citeseer, Vienna, Austria.
- Dumitru Potop-Butucaru and others. 2005. The Synchronous Hypothesis and Synchronous Languages. In *Embedded Systems Handbook*, R. Zurawski (Ed.).
- Francisco Sant'Anna and others. 2013. Safe System-level Concurrency on Resource-Constrained Nodes. In *Proceedings of SenSys'13*. ACM.
- Klaus Schneider, Jens Brandt, and Tobias Schuele. 2004. Causality analysis of synchronous programs with delayed actions. In *Proceedings of the 2004 international conference on Compilers, architecture, and synthesis for embedded systems*. ACM, 179–189.
- Klaus Schneider, Jens Brandt, and Tobias Schuele. 2006. A verified compiler for synchronous programs with local declarations. *Electronic Notes in Theoretical Computer Science* 153, 4 (2006), 71–97.
- Klaus Schneider and Michael Wenz. 2001. A new method for compiling schizophrenic synchronous programs. In *Proceedings of the 2001 international conference on Compilers, architecture, and synthesis for embedded systems*. ACM, 49–58.
- Ellen M Sentovich. 1997. Quick conservative causality analysis. In *System Synthesis, 1997. Proceedings, Tenth International Symposium on*. IEEE, 2–8.
- Thomas R Shiple, Gerard Berry, and Hemé Touati. 1996. Constructive analysis of cyclic circuits. In *European Design and Test Conference, 1996. ED&TC 96. Proceedings*. IEEE, 328–333.
- Doug Simon, Cristina Cifuentes, Dave Cleal, John Daniels, and Derek White. 2006. Java on the bare metal of wireless sensor devices: the Squawk Java Virtual Machine. In *Proceedings of the 2nd international conference on Virtual execution environments*. ACM, 78–88.

- Olivier Tardieu and Robert De Simone. 2003. Instantaneous termination in pure Esterel. In *Static Analysis*. Springer, 91–108.
- Olivier Tardieu and Robert De Simone. 2004. Curing schizophrenia by program rewriting in Esterel. In *Formal Methods and Models for Co-Design, 2004. MEMOCODE'04. Proceedings. Second ACM and IEEE International Conference on*. IEEE, 39–48.
- Stefan Tilkov and Steve Vinoski. 2010. Node.js: Using JavaScript to build high-performance network programs. *IEEE Internet Computing* 6 (2010), 80–83.
- Reinhard Von Hanxleden. 2009. SyncCharts in C: a proposal for light-weight, deterministic concurrency. In *Proceedings of the seventh ACM international conference on Embedded software*. ACM, 225–234.
- Jeong-Han Yun, Chul-Joo Kim, Seonggun Kim, Kwang-Moo Choe, and Taisook Han. 2013. Detection of harmful schizophrenic statements in esterel. *ACM Transactions on Embedded Computing Systems (TECS)* 12, 3 (2013), 80.