A

The Design and Implementation of the Synchronous Language CÉU

Francisco Sant'Anna, Departamento de Informática, PUC–Rio
Roberto Ierusalimschy, Departamento de Informática, PUC–Rio
Noemi Rodriguez, Departamento de Informática, PUC–Rio
Silvana Rossetto, Departamento de Ciência da Computação, UFRJ
Adriano Branco, Departamento de Informática, PUC–Rio

CÉU is a synchronous language inspired by Esterel which aims to offer a simpler semantics and more fine-grained control over the program execution. CÉU employs an event-triggered notion of time that allows for a simple reasoning about reactions in parallel and enables compile-time analysis to produce deterministic and concurrency-safe programs. We discuss the particularities of our design in comparison to Esterel, such as stack-based internal events, temporal analysis for concurrent statements, safe integration with *C*, and first-class timers. We also present two implementation back ends: one aiming for resource efficiency and interoperability with *C*, and another based on a virtual machine that allows remote reprogramming.

Additional Key Words and Phrases: Concurrency, Determinism, Embedded Systems, Esterel, Synchronous, Reactivity

1. INTRODUCTION

An established alternative to *C* in the field of embedded systems is the family of reactive synchronous languages [Benveniste et al. 2003]. Two major styles of synchronous languages have evolved: in the *control–imperative* style, programs are structured with control flow primitives, such as parallelism, repetition, and preemption; in the *dataflow–declarative* style, programs can be seen as graphs of values, in which a change to a value is propagated through its dependencies without explicit programming. Considering the control-based languages, Esterel [Boussinot and de Simone 1991] was the first to appear and succeed, influencing a number of embedded languages, such as *Reactive-C* [Boussinot 1991], *OSM* [Kasten and Römer 2005], *Sync-C* [Von Hanxleden 2009], and *PRET-C* [Andalam et al. 2010].

Despite its success and influence, Esterel has an overly complex semantics that requires careful static analysis such as to detect and refuse programs with *causality* and *schizophrenia* problems. Both problems have an extensive coverage and debate in the literature [Berry 1996; Shiple et al. 1996; Sentovich 1997; Boussinot 1998; Schneider and Wenz 2001; Tardieu and De Simone 2004; Edwards 2005; Yun et al. 2013]. The complex semantics not only challenges the analysis and compilation of programs, but also results in incompatible and non-compliant implementations. Above all, it also affects the programmer's understanding about the code, which, ultimately, has to solve the errors when facing corner cases. Another drawback of the Esterel semantics consists of the loose and non-deterministic execution for intra-reaction statements, which prevents threads to interact with stateful system calls safely and makes shared-memory concurrency not as straightforward as reading and writing to shared variables.

In this work, we present CÉU, a new programming language that inherits the synchronous and imperative mindset of Esterel but diverges in fundamental semantic aspects. Overall, CÉU has a simple semantics with fine-grained control for intra-reaction execution which is amenable to a simple temporal analysis that improves safety. The list that follows summarizes the contributions behind the design of CÉU:

— Unique and queue-based external events, which define the notion of time in CÉU.
— Stack-based internal events for intra-reaction communication, which also provides a limited form of coroutines.

— Static temporal analysis to detect suspicious concurrent statements.
— Safe integration with *C* which enforces finalization for external resources.
— First-class synchronized timers with a dedicated syntax.

We also present a lightweight single-threaded implementation of CÉU with two back ends: one aiming for resource efficiency and interoperability with *C*, and another as a virtual machine that allows remote reprogramming. Our implementations target resource-constrained devices, such as *Arduino* and *MICAz* sensor nodes based on 8-bit microcontrollers[1], showing the practical aspect of our clean semantics.

In previous work [Sant'Anna et al. 2013; Branco et al. 2015], we employed CÉU in the context of wireless sensor networks, developing a number of applications, protocols, and device drivers. We evaluated the expressiveness of CÉU in comparison to hand-crafted event-driven code in *C* and attested a reduction in source code size (around 25%) with a small increase in memory usage (around 5–10% for *text* and *data*) [Sant'Anna et al. 2013]. For the *VM* back end, applications have a bytecode footprint in the order of hundreds of bytes which can be transmitted over the air in a few packets [Branco et al. 2015].

The rest of the paper is organized as follows: Section 2 discusses the design of CÉU, focusing on the fundamental differences to Esterel. Section 3 presents the *C* and *VM* implementation back ends. Section 4 discusses other synchronous languages targeting embedded systems. Section 5 concludes the paper.

2. THE DESIGN OF CÉU

CÉU is a synchronous reactive language inspired by Esterel in which programs advance in a sequence of discrete reactions to external events. Like Esterel, CÉU is designed for control-intensive applications, supporting concurrent lines of execution, known as *trails*, and broadcast communication through events. Internal computations within a reaction (e.g. expressions, assignments, and system calls) are considered to take no time in accordance with the synchronous hypothesis [Potop-Butucaru et al. 2005]. An `await` is the only statement that halts a running reaction and allows a program to advance in this notion of time. To ensure that reactions run in bounded time and programs always progress, loops are statically required to contain at least one `await` statement in all possible paths [Sant'Anna et al. 2013; Berry 2000]. CÉU shares the same limitations with Esterel and synchronous languages in general: computations that run in unbounded time (e.g., cryptography, image processing) do not fit the zero-delay hypothesis, and cannot be directly implemented.

Figure 1 shows side-by-side the implementations in Esterel `[a]` and CÉU`[b]` for the following control specification: *"Emit an output O as soon as two inputs A and B have occurred. Reset this behavior each time the input R occurs"* [Berry 2000]. The first phrase of the specification, awaiting and emitting the events, is translated almost identically in the two languages (ln. 5–10, in both implementations), i.e., Esterel's '||' and CÉU's `par/and` constructs are equivalent. For the second phrase, the reset behavior, the Esterel version uses a `abort-when` (ln. 4–11), which, in this case, serves the same purpose of CÉU's `par/or` (ln. 4–13): the occurrence of event `R` aborts the awaiting statements in parallel and restarts the enclosing `loop`.

In the subsections that follow, we discuss the main differences between CÉU and Esterel: Unique and queue-based external events (2.1); Stack-based internal events (2.2); Static temporal analysis for concurrent statements (2.3); Safe integration with *C* (2.4); and First-class synchronized timers (2.5).

[1]Both *Arduino* and *MICAz* use the 8-bit *ATmega328* microcontroller with 32K of FLASH and 2K of SRAM.

```
input  A, B;
output O;
loop
   abort
      [
         await A
      ||
         await B
      ];
      emit O
   when R
end

.
```

[a] Esterel

```
input  void A, B;
output void O;
loop do
   par/or do
      par/and do
         await A;
      with
         await B;
      end
      emit O;
   with
      await R;
   end
end
```

[b] CÉU

Fig. 1. A control specification implemented in Esterel and CÉU: *"Emit O after A and B, resetting each R."*

2.1. Unique and Queue-Based External Events

Esterel defines time as a discrete sequence of logical unit instants or "ticks". At each tick, the program reacts to an arbitrary number of simultaneous input events, depending on external stimuli the environment provides. In contrast, CÉU defines time as a discrete sequence of reactions to unique input events. At each input event, which constitutes a logical unit of time, the program reacts exclusively to it. The execution of a program in CÉU is as follows [Sant'Anna et al. 2013]:

(1) The program initiates the "boot reaction" in a single trail (but parallel constructs may create new trails).
(2) Active trails execute until they await or terminate. This step is named a *reaction chain*, and always runs in bounded time.
(3) The program goes idle and the environment takes control.
(4) On the occurrence of a new external input event, the environment awakes *all* trails awaiting that event. It then goes to step 2.

A program must react to an occurring event completely before handling the next. Based on the synchronous hypothesis, a program takes no time on step 2 and is always idle on step 3. In practice, if a new external input event occurs while a reaction chain is running (step 2), it is enqueued to occur in the next reaction.

Figure 2 compares the discrete notions of time in Esterel and CÉU. The box `Real World` shows event occurrences over a continuous timeline divided in units of 10 milliseconds. Esterel and CÉU use discrete logical units of time in which the occurring events (which are the same in all boxes) have to fit somehow:

— `[Box-1]`: "Esterel with fixed-length ticks", within which reactions to occurring inputs have to fit [Li et al. 2005]. We assume a `R(boot)` reaction at `tick-0` which happens before any input. The input `A` "physically" occurs during the boot reaction but, because time is discrete, its corresponding reaction only executes in the next tick. Note that `R(A)` takes more time than `tick-1`, causing a *timing violation* [Li et al. 2005] that invades `tick-2`. The events `B` and `C` occur during `tick-1` and are delayed to happen *simultaneously* at `tick-2` with `R(B+C)`. Since no new events occur during `tick-2`, the CPU stays idle during the whole `tick-3`. Finally, one instance of event `D` and two instances of event `E` occur "simultaneously" during the idle `tick-3`. However, the program can only detect one occurrence of `E`, which is considered in `R(D+E)`.
— `[Box-2]`: "Esterel with variable-length ticks", which are adjusted to the CPU time the corresponding reaction takes to complete [Roop et al. 2004]. This approach avoids the time violation for `R(A)` and also results in smaller idle periods. For instance,

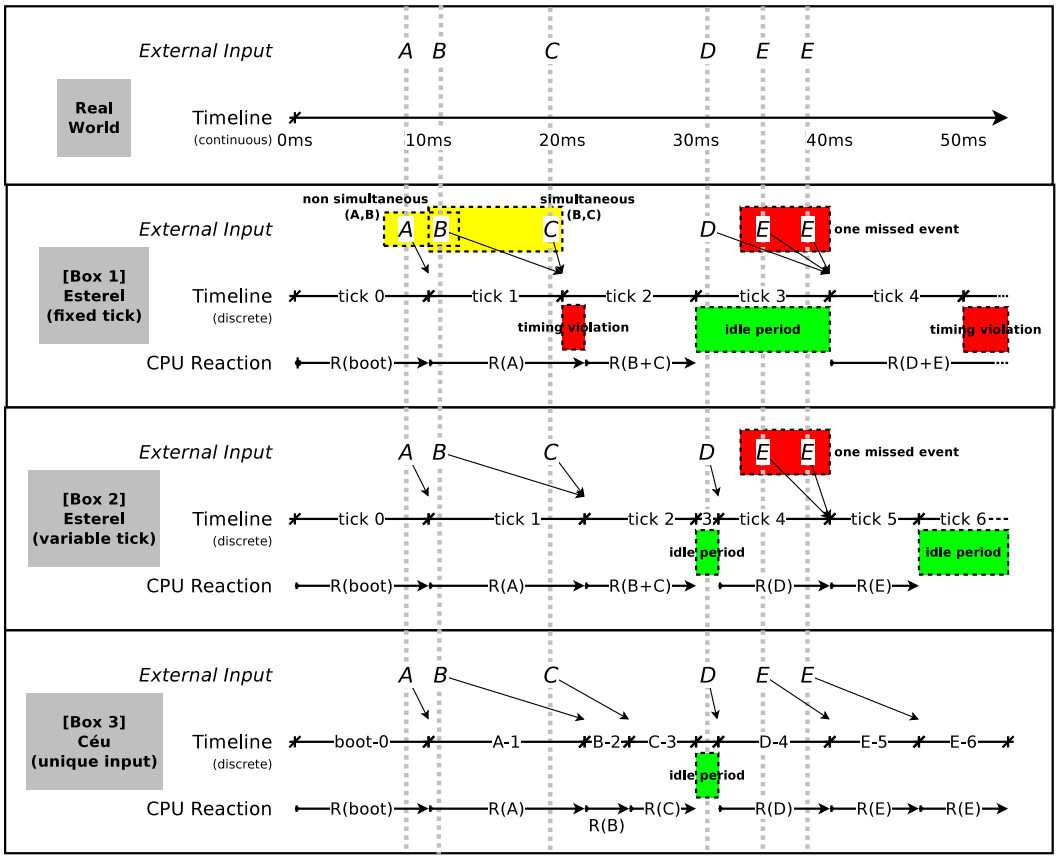

Fig. 2. The discrete notions of time in Esterel and CÉU.

the occurrence of `D` interrupts the idle `tick-3` to react alone as `R(D)` on `tick-4`. Similarly to the fixed-tick approach, only one of the two simultaneous occurrences of `E` is considered on `tick-5`, now because `R(D)` takes too long.

— `[Box-3]`: "CÉU with unique and queue-based input events". We also assume a `R(boot)` reaction before any input. Because the occurrence of event `A` is unique during `tick-0`, the behavior in CÉU is similar to `Box-2` for the first two units of time (`boot-0` and `A-1`). However, CÉU does not consider the events `B` and `C` as simultaneous, and handles each in subsequent reactions `R(B)` and `R(C)`. We assume the CPU times for `R(B+C)` in Esterel and `R(B)+R(C)` in CÉU to be roughly the same. This way, the first idle period in `Box-2` and `Box-3` coincide. Finally, CÉU recognizes and reacts to the two instances of `E` independently, which are handled in sequence.

We decided for the unique and queue-based semantics in CÉU for the reasons that follow:

— **A "tick" is too abstract and imprecise:** Outside the domain of hardware specification, a tick has no natural counterpart in the real world. Also, since ticks require no time regularity [Berry and Sentovich 2001], the two approaches for Esterel in Figure 2 are legitimate, but lead to different behaviors for the same sequence of inputs.

— **Events are never absolutely simultaneous:** From a rigorous point of view, event occurrences are infinitesimal, having zero probability of being simultaneous. This way, we believe that the notion of simultaneity should not be imposed by the language, but defined explicitly for each use case. In the case of Esterel, simultaneity depends on the imprecise length of discrete ticks. For instance, in Figure 2 (`box-1` and `box-2`), the events `B` and `C` are simultaneous, even though `A` and `B` actually happen much closer to one another. In Section 2.5, after introducing internal events and first-class timers, we show an example that detects simultaneous button clicks.
— **Unique input events imply mutual exclusion:** Reactions to multiple events never overlap because they are atomic. Automated mutual exclusion simplifies the reasoning about the program and is a prerequisite for the temporal analysis to be discussed in Section 2.3.

The synchronous hypothesis for CÉU holds if the reactions run faster than the rate of incoming input events. Otherwise, the application continuously accumulates delays between the real occurrence and actual reaction of a given event. This is also the case for the variable-length-tick approach of Esterel, since the more inputs to handle, the longer the reaction takes, and the more inputs accumulates for subsequent ticks. For the fixed-length-tick approach of Esterel, a breach in the synchronous hypothesis causes timing violations, which can be avoided with *worst case reaction time* analysis to infer an appropriate value for the tick length [Li et al. 2005].

2.2. Stack-Based Internal Events

Esterel makes no semantic distinctions between internal and external signals. In particular, programs can emit external and internal signals simultaneously, with all coexisting during the entire reaction. In CÉU, subsequent external events coming from the environment define a timeline, which programs cannot interfere by emitting new external events. To cope with deterministic intra-reaction communication, CÉU supports internal events which programs can `emit` and `await`. In contrast with external events, internal events follow a stack-based execution policy similar to subroutine calls in typical programming languages. Figure 3 illustrates the use of internal signals (events) in Esterel `[a]` and CÉU`[b]`. In Esterel, when `A` occurs, `B` is emitted (ln. 4–5) and both events become active, resulting in the invocation of `f()` and `g()` in no particular order. In CÉU, the occurrence of `A` makes the program behave as follows:

(1) 1st trail awakes (ln. 4), broadcasts `b`, and pauses.
(2) 2nd trail awakes (ln. 8), calls `_g()`, and terminates. *(No other trails awake to `b`.)*
(3) 1st trail (on top of the stack) resumes, calls `_f()`, and terminates.
(4) Both trails have terminated, so the `par/and` rejoins, and the program also terminates.

Internal events provide a fine-grained intra-reaction control mechanism. For instance, it brings a limited form of subroutines, as depicted in Figure 4. The subroutine

```
1  input  A;   // external
2  signal B;   // internal
3  [[
4      await A;
5      emit B;
6      call f();
7  ||
8      await B;
9      call g();
10 ]]
```

[a] Esterel

```
1  input void A;  // external (in uppercase)
2  event void b;  // internal (in lowercase)
3  par/and do
4      await A;
5      emit b;
6      _f();
7  with
8      await b;
9      _g();
10 end
```

[b] CÉU

Fig. 3. Internal signals (events) in Esterel and CÉU: similar syntax, but different semantics.

```
1  event int* inc; // subroutine 'inc'
2  par/or do
3      loop do     // definitions are loops
4          var int* p = await inc;
5          *p = *p + 1;
6      end
7  with
8      var int v = 1;
9      await A;
10     emit inc => &v; // call 'inc'
11     _assert(v==2);  // after return
12 end
```

Fig. 4. Subroutine `inc` is defined in a loop (ln. 3–6), in parallel with the caller (ln. 8–11).

`inc` is defined as a loop (ln. 3–6) that continuously awaits its identifying event (ln. 4), incrementing the value passed as reference (ln. 5). A trail in parallel (ln. 8–11) invokes the subroutine in reaction to event `A` through an `emit` (ln. 10). Given the stacked execution for internal events, the calling trail pauses, the subroutine awakes (ln. 4), runs its body (yielding `v=2`), loops, and awaits the next "call" (ln. 4, again). Only after this sequence the calling trail resumes and passes the assertion test (ln. 11).

On the one hand, this form of subroutines has a significant limitation that it cannot express recursive calls: an `emit` to itself is always ignored, given that a running body cannot be awaiting itself. On the other hand, this very same limitation brings some important safety properties to subroutines: first, they are guaranteed to react in bounded time; second, memory for locals is also bounded, not requiring data stacks. Also, this form of subroutines can use the other primitives of CÉU, such as parallel compositions and the `await` statement. In particular, they await keeping context information such as locals and the program counter, similarly to coroutines [Moura and Ierusalimschy 2009].

Another distinction regarding event handling in comparison to CÉU is that Esterel supports same-cycle bi-directional communication [Edwards 1999], i.e., two threads can react to one another during the same cycle due to mutual signal dependency. CÉU has a different take, posing a tradeoff that an `await` is only valid for the next reaction, i.e., if an `await` and `emit` occur simultaneously in parallel trails, the `await` does not awake. These *delayed awaits* avoid corner cases of instantaneous termination and re-execution of statements in the same reaction (known as *schizophrenic statements* [Berry 1996]).

The example in Figure 5 illustrates delayed awaits, which prevents infinite execution by design. Both sides of the `par/or` have an `await` statement to avoid instantaneous termination (ln. 4,7). However, if the `emit` (ln. 6) could awake the `await` (ln. 4) in the same reaction that reaches them, the `par/or` would terminate and restart the `loop` instantaneously, resulting in infinite execution.

```
event void e,f;
loop do
    par/or do
        await e;
    with
        emit e;     // w/o delayed awaits, the emit awakes 1st trail
        await f;    // and restarts the loop instantaneously
    end
end
```

Fig. 5. Delayed awaits prevents re-execution of statements by design.

```
1  event void e;
2  par do
3      loop do
4          <...>     // code that awaits some period
5          emit e;   // periodic request
6      end
7  with
8      loop do
9          <...>     // code to execute immediately and then periodically
10         await e;  // await after
11     end
12 end
```

Fig. 6. An example that circumvents the *delayed await* by post-fixing the `await` inside the `loop`.

```
1  input void A, B;
2  var int x = 1;
3  par/and do
4      await A;
5      x = x + 1;
6  with
7      await B;
8      x = x * 2;
9  end
```

[a] Accesses to `x` are safe

```
1  input void A;
2  var int y = 1;
3  par/and do
4      await A;
5      y = y + 1;
6  with
7      await A;
8      y = y * 2;
9  end
```

[b] Accesses to `y` are unsafe

Fig. 7. Shared-memory concurrency in CÉU: Example [a] is safe because the trails access `x` atomically in different reactions; Example [b] is unsafe because both trails access `y` in the same reaction.

In atypical scenarios requiring immediate awake, delayed awaits can be circumvented by manually copying or transforming the code to execute on awake. For instance, sometimes we need to execute a block of code immediately, and then periodically from internal event requests, as illustrated in Figure 6. In this case, the `await` moved to the end of the loop (ln. 10) makes the periodic code to also execute immediately (ln. 9), and then in reactions to each `emit` request (ln. 5). If the periodic `emit` depends on a condition, then the code transformation becomes more intricate, requiring an extra condition test around the periodic code to prevent its immediate execution. On the one hand, we transfer the burden of dealing with these specific corner cases to the programmer. On the other hand, we simplify the semantics of the language and eliminate the need for complex analysis to deal with schizophrenic statements.

2.3. Static Temporal Analysis for Concurrent Statements

Embedded applications make extensive use of global memory and shared resources, such as through memory-mapped registers and system calls to device drivers. Hence, an important goal of CÉU is to ensure a reliable behavior for programs with concurrent lines of execution sharing memory and interacting with the environment.

Esterel is only deterministic with respect to reactive control: "the same sequence of inputs always produces the same sequence of outputs" [Berry 2000]. However, the execution order for operations within a reaction is non-deterministic: "if there is no control dependency, as in `<<call f1() || call f2()>>`, the order is unspecified and it would be an error to rely on it" [Berry 2000]. A number of Esterel-based synchronous languages, such as *SOL* [Karpinski and Cahill 2007], *SC* [Von Hanxleden 2009], and *PRET-C* [Andalam et al. 2010], enforce intra-reaction determinism with an arbitrary execution order for statements in multiple lines of execution. CÉU also takes the deterministic approach and when multiple trails are active during the same reaction, they are scheduled in the order they appear in the program source code.

Even so, we consider that enforcing an arbitrary execution order can be misleading in some cases. For instance, consider the two examples in Figure 7, both of which define a shared variable (ln. 2), and assign to it in parallel trails (ln. 5, 8). In the example `[a]`, the two assignments to `x` execute from reactions to different events `A` and `B` which, by definition, cannot occur simultaneously (Section 2.1). Hence, for the sequence `A->B`, `x` becomes `(1+1)*2->4`, while for `B->A`, `x` becomes `(1*2)+1->3`. In the example `[b]`, the two assignments to `y` are simultaneous because they execute in reaction to the same event `A`. Since CÉU employs lexical order for intra-reaction statements, the execution is still deterministic, and `y` always becomes `(1+1)*2->4`. However, an (apparently innocuous) change in the order of trails changes the semantics of the program, which we consider unsafe.

To mitigate this threat, CÉU performs a temporal analysis at compile time and detects concurrent accesses to shared variables, as follows: if a variable is written in a trail segment, then a concurrent trail segment cannot read or write to that variable, nor dereference a pointer of that variable type. A trail segment is a sequence of statements followed by an `await` (or termination). Concurrency in CÉU is characterized when two or more trail segments in parallel react to the same input event. Considering the examples in Figure 7:

— The assignments to `x` in lines 2 and 5 in `[a]` **cannot** be concurrent because they are not in parallel trails.
— The assignments to `x` in lines 5 and 8 in `[a]` **cannot** be concurrent because they **cannot** execute during the same reaction.
— The assignments to `y` in lines 5 and 8 in `[b]` **can** be concurrent because they are in parallel trails and **can** execute during the same reaction.

The algorithm for the analysis, which is depicted in Section 3.1, inspects all possible `await` statements that precede a variable access and keep a list with all corresponding awaking events. Then, it checks all accesses in parallel trails to see if they share an awaking event. If it is the case, the compiler warns about the suspicious accesses.

Note that this analysis is only possible due to the uniqueness of input events within reactions. Otherwise, any two trail segments in parallel can be concurrent, even if they react to different input events.

2.4. Safe Integration with *C*

In CÉU, any identifier prefixed with an underscore is repassed *as is* to the *C* compiler that generates the final binary. Therefore, access to *C* is seamless and, more importantly, easily trackable. Similarly to Esterel, which supports the `call` primitive for external code, calls are assumed to be instantaneous [Berry 2000]. This way, programs should only resort to *C* for asynchronous functionality[2], such as non-blocking I/O, or simple `struct` accessors, but never for control purposes.

The temporal analysis of Section 2.3 also considers calls and accesses to external symbols in *C*. As an example, the program in Figure 8.a defines four external symbols inside a `native` block with standard declarations in *C* (ln. 1–6). The `par/and` (ln. 7–11) creates two trails that react concurrently during the boot reaction (ln. 8,10): the first trail calls symbol `_f`, while the second calls `_g` and `_id`, and also reads `_NUM`. Since CÉU does not inspect any code in *C*, it complains about suspicious concurrent accesses between `_f` and all symbols in the second trail.

[2]In CÉU, it is possible to restrict the available *C* symbols as a compile-time option.

```
1  native do
2      #define NUM 10
3      void f  (void)  { <...> }
4      void g  (int v) { <...> }
5      int  id (int v) { <...> }
6  end
7  par/and do
8      _f();
9  with
10     _g(_id(_NUM));
11 end
```

[a] Definitions and uses of symbols

```
native @const _NUM;
native @pure  _id;
native @safe  _f with _g;
.
```

[b] Annotations for the symbols in [a]

Fig. 8. The unsafe program in [a] only compiles with the annotations in [b].

The code in Figure 8.b uses annotations to provide hints to the compiler about the semantics of the *C* symbols in use in [a]. With the annotations in [b], the program in [a] compiles successfully:

— _NUM is a constant symbol, meaning that it is safe to use it concurrently with any other symbol in the program.
— _id is a pure function, also meaning that it is safe to call it concurrently with any other symbol in the program.
— Both _f and _g are impure, but their effects do not conflict and they can be safely called concurrently.

Esterel's abort and CÉU's par/or statements distinguish synchronous from asynchronous languages with respect to orthogonal abortion of lines of execution (i.e., without tweaking them with synchronization primitives) [Berry 1993]. In contrast, traditional (asynchronous) multi-threaded languages cannot express thread termination safely [ORACLE 2011]. Still, aborting lines of execution that deal with external resources is challenging because they may end up in an inconsistent state. For this reason, Esterel and CÉU provide a finalize construct to unconditionally execute a series of statements even if the enclosing block is aborted and does not terminate normally.

The example in Figure 9 in Esterel [a] and CÉU[b] uses the lock and unlock calls which represent accesses to an external resource. The normal behavior is to lock the resource (ln. [a]:4 and [b]:3), perform some operations in subsequent reactions to input A (ln. [a]:5–8 and [b]:7–10), and then unlock the resource (ln. [a]:10 and [b]:5). However, if the aborting input B (ln. [a]:12 and [b]:12) occurs after the lock but before the reactions to A, we still want to call unlock to safely release the resource. In Esterel and CÉU, the finalize clauses (ln. [a]:10 and [b]:5) are automatically called if the enclosing blocks (ln. [a]:3–1 and [b]:3–10) are externally aborted (ln. [a]:12 and [b]:12).

CÉU goes one step further and enforces the use of finalize for system calls that deal with pointers representing resources:

— If a system call **receives** a pointer from CÉU, the pointer represents a **local** resource that requires finalization.
— If a system call **returns** a pointer to CÉU, the pointer represents an **external** resource that requires finalization.

CÉU tracks the interaction of system calls with pointers and requires finalization clauses to accompany them. In the example in Figure 10.a, the local variable msg (ln. 2) is an internal resource passed as a pointer to _send_request (ln. 5), which is an asynchronous call that transmits the buffer in the background. If the block aborts (ln. 11) before receiving an acknowledge from the environment (ln. 9), the local msg goes out

```
1  input A, B;
2  abort
3      finalize
4          call lock();
5          await A;
6          <...>;      // do something
7          await A;
8          <...>;      // do something
9      with
10         call unlock();
11     end
12 when B
13 .
```

[a] Esterel

```
1  input void A, B;
2  par/or do
3      _lock();
4      finalize with
5          _unlock();
6      end
7      await A;
8      <...>;      // do something
9      await A;
10     <...>;      // do something
11 with
12     await B;
13 end
```

[b] CÉU

Fig. 9. Finalization in Esterel and CÉU: after the call to `lock`, both languages guarantee to call `unlock` if the enclosing block aborts when `B` occurs.

```
1  par/or do
2     var _buffer_t msg;
3     <...> // prepare msg
4     finalize
5        _send_request(&msg);
6     with
7        _send_cancel(&msg);
8     end
9     await SEND_ACK;
10 with
11    <...>
12 end
13 .
```

[a] Local resource finalization

```
1  par/or do
2     var _FILE* f;
3     finalize
4        f = _fopen(...);
5     with
6        _fclose(f);
7     end
8     _fwrite(..., f);
9     await A;
10    _fwrite(..., f);
11 with
12    <...>
13 end
```

[b] External resource finalization

Fig. 10. CÉU enforces the use of finalization to prevent *dangling pointers* for local resources and *memory leaks* for external resources.

of scope and the external transmission now holds a *dangling pointer*. The finalization ensures that the transmission also aborts (ln. 7). In the example in Figure 10.b, the call to `_fopen` returns an external file resource as a pointer. If the block aborts (ln. 12) during the `await A` (ln. 9), the file remains open as a *memory leak*. The finalization ensures that the file closes properly (ln. 6). In both cases, the code does not compile if the programmer forgets to use the `finalize` construct.

Note that the illustrative example of Figure 9 does not manipulate pointers (i.e., the resource is a *singleton*). That case is an example of a bad and unsafe API to expose to CÉU because the compiler will not enforce the use of finalization.

2.5. First-Class Timers

Activities that involve reactions to *wall-clock time*[3] appear in typical patterns of embedded development, such as timeout watchdogs and sensor samplings. However, support for wall-clock time is somewhat low-level in existing languages, usually through timer callbacks or "sleep" blocking calls. Furthermore, in any concrete timer implementation, a requested timeout does not expire precisely without delays, a fact that is usually ignored in the development process. We define the difference between the requested timeout and the actual expiring time as the *residual delta time (delta)*. Without explicit manipulation, the recurrent use of timed activities in sequence (or in a loop) may accumulate a considerable amount of deltas that can lead to incorrect behavior in programs.

[3]By wall-clock time we mean the passage of time from the real world, measured in hours, minutes, etc.

```
var int v;
await 10ms;
v = 1;
await 1ms;
v = 2;



.
```

[a]

```
par/or do
    await 10ms;
    <...>          // any non-awaiting sequence
    await  1ms;
    v = 1;
with
    await 12ms;
    v = 2;
end
```

[b]

Fig. 11. First-class timers in CÉU.

The `await` statement of CÉU supports wall-clock time and handles deltas automatically, resulting in more robust applications. For the example in Figure 11.a, suppose that after the first `await` request, the underlying system gets busy and takes 15ms to check for expiring awaits. The CÉU scheduler will notice that the `await 10ms` has not only already expired, but is delayed with `delta=5ms`. Then, the awaiting trail awakes, sets `v=1`, and invokes `await 1ms`. As the current delta is higher than the requested timeout (i.e. $5ms > 1ms$), the trail is rescheduled for execution, now with `delta=4ms`.

CÉU also takes into account the fact that time is a physical quantity that can be added and compared. For instance, for the example in Figure 11.b, although the scheduler cannot guarantee that the first trail terminates exactly in 11ms, it can at least ensure that the program always terminates with `v=1`. Given that any non-awaiting sequence is considered to take no time in the synchronous model, the first trail is guaranteed to terminate before the second trail, because $10+1 < 12$. A similar program in a language without first-class support for timers would depend on the execution timings for the code marked as `<...>`, making the reasoning about the execution behavior more difficult.

In Section 2.1, we argue that events occurrences are infinitesimal and can never be absolutely simultaneous. However, the "sensation of simultaneity" is not infinitesimal, but increases with the inaccuracy of the observer (e.g., a human being). Therefore, simultaneity should be defined with respect to the observer, which varies from case to case.

First-class timers simplify the implementation of application-defined simultaneity. Figure 12 emulates a `milddle_click` event (ln. 3) in terms of "simultaneous" occurrences of `LEFT_CLICK` and `RIGHT_CLICK` (ln. 1–2). If both events occur, we emit the internal event `middle_click` (ln. 6–7). However, if one of them occurs and the *200ms* timer expires (ln. 9–10), we abort the whole behavior with the `par/or` (ln. 5–11) and try again with the enclosing `loop` (ln. 4). In this specification, "simultaneous" means "within 200 milliseconds", which is a huge amount of time for a language-defined tick, and which would break the synchronous hypothesis. Note that a similar implementation for Esterel would not rely on the tick notion of simultaneity either. Note also that the application in CÉU could define an external input `TICK` and change the expiring event from `200ms` to `TICK` (ln. 10) to recover the abstract notion of simultaneity.

3. IMPLEMENTATION

The compilation process of a program in CÉU for the original *C* back end is composed of three main phases: the *parsing phase* converts the source code in CÉU to an *abstract syntax tree (AST)*; the *temporal analysis phase* detects inconsistencies in programs, such as unbounded loops and suspicious accesses to shared memory; the *code generation phase* converts the *AST* to standard *C* code and packs it with platform-dependent functionality (e.g., system calls) and the runtime of CÉU, compiling everything with *gcc* to generate the final binary.

```
1  input void LEFT_CLICK;
2  input void RIGHT_CLICK;
3  event void middle_click;
4  loop do
5      par/or do
6          AWAIT_AND(LEFT_CLICK, RIGHT_CLICK);
7          emit middle_click;
8      with
9          AWAIT_OR(LEFT_CLICK, RIGHT_CLICK);
10         await 200 ms;
11     end
12 end
```

```
13 #define AWAIT_AND(e1, e2)    \
14     par/and do               \
15         await e1;            \
16     with                     \
17         await e2;            \
18     end
19 #define AWAIT_OR(e1, e2)     \
20     par/or do                \
21         await e1;            \
22     with                     \
23         await e2;            \
24     end
```

Fig. 12. Application defines that a `middle_click` event occurs whenever both `LEFT_CLICK` and `RIGHT_CLICK` occur within 200 milliseconds. The macros `AWAIT_AND` (ln. 13–18) and `AWAIT_OR` (ln. 19–24) are simple expansions to a `par/and` and `par/or` for better readability.

In the sections that follow, we discuss the implementation highlights related to the peculiarities of CÉU: the temporal analysis for determinism (Section 3.1), static memory allocation for data and trails (Sections 3.2 and 3.3), static scheduling and trail finalization (Section 3.4), interaction with the environment (Section 3.6), and the *VM* back end (Section 3.7),

3.1. Temporal Analysis for Concurrent Statements

TODO: internal events, emits are propagated from preceding awaits to input

The compile-time *temporal analysis* phase detects inconsistencies in CÉU programs. Here, we focus on the algorithm that detects suspicious access to shared variables, as discussed in Section 2.3.

For each node representing a statement in the program AST, we keep the set of events I (for *incoming*) that can lead to the execution of the node, and also the set of events O (for *outgoing*) that can terminate the node.

A node inherits the set I from its direct parent and calculates O according to its type:

— Nodes that represent expressions, assignments, *C* calls, and declarations simply reproduce $O = I$, as they do not await;
— An `await e` statement has $O = \{e\}$.
— A `break` statement has $O = \{\}$ as it escapes the innermost `loop` and never terminates, i.e., never proceeds to the statement immediately following it (see also `loop` below);
— A *sequence node (;)* modifies each of its children to have $I_n = O_{n-1}$. The first child inherits I from its parent node, and the set O for the sequence node is copied from its last child, i.e., $O = O_n$.
— A `loop` node includes its body's O on its own I ($I = I \cup O_{body}$), as the loop is also reached from its own body. The union of all `break` statements' O forms the set O for a `loop`.
— An `if` node has $O = O_{true} \cup O_{false}$.
— A parallel composition may terminate from any of its branches, hence $O = O_1 \cup ... \cup O_n$.

With all sets calculated, we take all pairs of nodes that perform side effects and are in parallel branches, comparing their sets I for intersections. For each pair, if the intersection is not the empty set, we mark both nodes as suspicious.

The example in Figure 13.a has a corresponding AST, in Figure 13.b, with the sets I and O for each node. The event . (dot) represents the "boot" reaction. The assignments to `y` in parallel (ln. 5,8 in the code) have an empty intersection of I (ln. 6,9 in the AST), hence, they do not conflict. Note that although the accesses in ln. 5,11 in the code (ln. 6,11 in the AST) do have an intersection, they are not in parallel and are also safe.

```
1  input void A, B;
2  var int y;
3  par/or do
4    await A;
5    y = 1;
6  with
7    await B;
8    y = 2;
9  end
10 await A;
11 y = 3;
```

[a]

```
1  Stmts I={.} O={A}
2      Dcl_y I={.} O={.}
3      ParOr I={.} O={A,B}
4          Stmts I={.} O={A}
5              Await_A I={.} O={A}
6              Set_y I={A} O={A}
7          Stmts I={.} O={B}
8              Await_B I={.} O={B}
9              Set_y I={B} O={B}
10     Await_A I={A,B} O={A}
11     Set_y I={A} O={A}
```

[b]

Fig. 13. A program with a corresponding AST describing the sets I and O. The program is safe because accesses to y in parallel have no intersections for I.

```
input int A, B, C;
do
    var int a = await A;
end
do
    var int b = await B;
end
par/and do
    await B;
with
    await C;
end
```

```
union {                 // sequence
    int a_1;            //   do_1
    int b_2;            //   do_2
    struct {            //   par/and
        int _and_3: 1;
        int _and_4: 1;
    };
} MEM ;
```

Fig. 14. A program with blocks in sequence and in parallel, with corresponding memory layout generated by the compiler.

3.2. Static Memory Layout

CÉU favors a fine-grained use of trails, being common to use trails that await a single event. For this reason, CÉU does not allocate per-trail stacks; instead, all data resides in fixed memory slots—this is true for the program variables as well as for temporary values and runtime flags. Memory for trails in parallel must coexist, while statements in sequence can reuse it. Translating this idea to *C* is straightforward [Kasten and Römer 2005; Bernauer and Römer 2013]: memory for blocks in sequence are packed in a `struct`, while blocks in parallel, in a `union`. CÉU reserves a single static block of memory to hold all memory slots, whose size is the maximum the program uses at a given time. A given position in the memory may hold different data (with variable sizes) during runtime. As an example, Figure 14 shows a program with corresponding memory layout. Each variable is assigned a unique id (e.g. `a_1`) so that variables with the same name can be distinguished. The `do-end` blocks in sequence are packed in a `union`, given that their variables cannot be in scope at the same time, e.g., `MEM.a_1` and `MEM.b_2` can safely share the same memory slot. The example also illustrates the presence of runtime flags related to the parallel composition, which also reside in reusable slots in the static memory.

3.3. Static and Lightweight Trail Allocation

Each line of execution in CÉU needs to carry associated data, such as which event it is awaiting and which code to execute when it awakes. The compiler statically infers the maximum number of trails a program can have at the same time and creates a static vector to hold the runtime information about them. Like normal variables, trails that cannot be active at the same time can share slots in the static memory vector.

At any given moment, a trail can be awaiting in one of the following states: INACTIVE, STACKED, FINALIZE, or in any of the events defined in the program:

```
enum {
    INACTIVE = 0,
    STACKED,
    FINALIZE,
    EVT_A,      // input void A;
    EVT_e,      // event int e;
    <...>       // other events
}
```

All terminated or not-yet-started trails stay in the INACTIVE state and are ignored by the scheduler. A STACKED trail holds an associated stack level and is delayed until the scheduler runtime reaches that level again. A FINALIZE trail represents a hanged finalization block which is only scheduled when its corresponding block goes out of scope. A trail waiting for an event stays in the state of the corresponding event, also holding the minimum sequence number (*seqno*) in which it can awake. In concrete terms, a trail is represented by the following struct:

```
struct trail_t {
    state_t evt;
    label_t lbl;
    union {
        unsigned char seqno;
        stack_t       stk;
    };
};
```

The field evt holds the state of the trail (or the event it is awaiting); the field lbl holds the entry point in the code to execute when the trail segment is scheduled; the third field depends on the evt field and may hold the seqno for an event, or the stack level stk for a STACKED state.

The size of state_t depends on the number of events in the application; for an application with less than 253 events (plus the 3 states), one byte is enough. The size of label_t depends primarily on the number of await statements in the application—each await splits the code in two segments and requires a unique entry point in the code for its continuation. Additionally, split & join points for parallel compositions, emit continuations, and finalization blocks also require labels. The seqno could eventually overflow during execution (i.e., every 256 reactions). However, given that the scheduler traverses all trails on every reaction, it can adjust them to properly handle overflows (actually, 2 bits to hold the seqno is already enough). The size of stack_t depends on the maximum depth of nested emissions and is bounded to the maximum number of trails. In the worst case, a trail emits an event that awakes another trail, which emits an event that awakes another trail, and so on. The last trail cannot awake any trail, because they are all hanged in the STACKED state.

In the context of embedded systems, the size of trail_t is typically only 3 bytes (1 byte for each field), imposing a negligible memory overhead even for trails that only await a single event and terminate. For instance, the *CTP* collection protocol ported to CÉU reaches eight simultaneous lines of execution with an overhead of 2% in comparison to the original version in *C* [Sant'Anna et al. 2013].

3.4. Static Scheduling and Trail Finalization

In the final generated code in *C*, each trail segment label representing an entry point becomes a *switch case* with the associated code to execute. Figure 15 illustrates the generation process. For the program in [a], the compiler extracts the entry points and associated trails, e.g., the label Awake_e will execute on TRAIL-0 (ln. 7). For each statement that pauses (emit and await), resumes (par/and, par/or, and finalize), or aborts

(par/or and break), the compiler splits the trail into segments with associated entry points. The entry points translate to an enum in the generated code (ln. 1–10, in [b]). The state of trails translates to a vector of type trail_t with the maximum number of simultaneous trails (ln. 12–15). On initialization, TRAIL-0 is set to execute the Main entry point (ln. 13), while all others are set to INACTIVE (in the example, only one, in ln. 14).

The scheduler executes in two passes: in the *broadcast* pass, it sets all trails that are waiting for the current event to STACKED in the current stack level; in the *dispatch* pass, it executes each trail that is STACKED to run in the current level, setting it immediately to INACTIVE (the trail segment may reset it in sequence if it doesn't terminate).

During the dispatch pass, if a trail executes and emits an internal event, the scheduler increments the stack level and re-executes the two passes. After all trails are properly dispatched, the scheduler decrements the stack level and resumes the previous execution. For the first reaction, the scheduler starts from the *dispatch* pass, given that the Main label is the only one that can be active at the stack level 0 (ln. 13, in Figure 15.b).

The code in Figure 15.c dispatches a trail segment according to the current label to execute. For the first reaction, it executes the Main label in TRAIL-0. When the Main label reaches the par/and, it stacks TRAIL-1 (ln. 4–7) and proceeds to the code in TRAIL-0 (ln. 9–14), respecting the deterministic execution order. The code sets the running TRAIL-0 to await EVT_e on label Awake_e, and then halts with a break. The next iteration of dispatch takes TRAIL-1 and executes its registered label And_sub_2 (ln. 16–21), which sets TRAIL-1 to await EVT_A and also halts.

Regarding abortion and finalization, when a par/or terminates, the scheduler makes a *broadcast* pass for the FINALIZE event, but limited to the range of trails covered by the terminating par/or. Trails that do not match the FINALIZE are set to INACTIVE, as they have to be aborted. Given that trails in parallel are allocated in subsequent slots in the

```
input void A;
event void e;
// TRAIL 0 — lbl Main
par/and do
  // TRAIL 0 — lbl Main
  await e;
  // TRAIL 0 — lbl Awake_e
  // TRAIL 0 — lbl And_chk
with
  // TRAIL 1 — lbl And_sub_2
  await A;
  // TRAIL 1 — lbl Awake_A_1
  emit e;
  // TRAIL 1 — lbl Emit_cont
  // TRAIL 1 — lbl And_chk
end
// TRAIL 0 — lbl And_out
await A;
// TRAIL 0 — lbl Awake_A_2
.
```

[a]

```
1  enum {
2    Main = 1,  // ln 3
3    Awake_e,   // ln 7
4    And_chk,   // ln 8,15
5    And_sub_2, // ln 10
6    Awake_A_1, // ln 12
7    Emit_cont, // ln 14
8    And_out,   // ln 17
9    Awake_A_2  // ln 19
10 };
11
12 trail_t TRLS[2] = {
13   { STACKED,  Main, 0 };
14   { INACTIVE, 0,    0 };
15 };
16
17
18
19
20
21
22
23
24
25 .
```

[b]

```
1  void dispatch (trail_t* t) {
2    switch (t->lbl) {
3      case Main:
4        // activate TRAIL 1
5        TRLS[1].evt = STACKED;
6        TRLS[1].lbl = And_sub_2;
7        TRLS[1].stk = cur_stack;
8
9        // code in the 1st trail
10       // await e;
11       TRLS[0].evt = EVT_e;
12       TRLS[0].lbl = Awake_e;
13       TRLS[0].seq = cur_seqno;
14       break;
15
16     case And_sub_2:
17       // await A;
18       TRLS[1].evt = EVT_A;
19       TRLS[1].lbl = Awake_A_1;
20       TRLS[1].seq = cur_seqno;
21       break;
22
23     <...>  // other labels
24   }
25 }
```

(c)

Fig. 15. [a] Static allocation of trails: the comments identify the trail indexes inferred by the compiler; [b] Entry-point labels: each trail segment has an associated numeric identifier generated by the compiler. (c) Dispatch function: uses a switch to associate each segment identifier with the corresponding code to execute.

static vector TRLS, this pass only aborts the desirable trails. The subsequent *dispatch* pass executes the finalization code. Escaping a `loop` that contains parallel compositions also triggers the same abortion process.

3.5. Single-Threaded Dispatching

The implementation of CÉU dispatches active trails sequentially in a single thread, taking no advantage of multi-core CPUs. This decision comes not only from the fact that CÉU targets constrained embedded systems with a single CPU, but also because CÉU imposes deterministic execution for intra-reaction statements.

Nonetheless, as discussed in Section 2.3, the temporal analysis of CÉU infers precisely trails that are concurrent yet do not share resources. Hence, these non-conflicting trails could potentially execute with real parallelism in multiple cores. However, our experiments with a multi-threaded implementation in multi-core CPUs execute slower than the single-threaded implementation in the same CPUs. Considering that we use CÉU primarily in control-dominated applications, this result is not surprising [Yuan et al. 2011; Haribi 2012]. We believe that the continuous fork-and-rejoin overhead due to small reactions as well as the excessive locality of data due to stackless threads sharing contiguous static memory seem to prevent any speed-up gains.

If we consider data-dominated applications, multi-core implementations can offer considerable speed-up gains. However, data-intensive computations do not typically require a disciplined step-wise execution and can actually execute in asynchronous calls. Esterel provides a `task` primitive for this purpose [Berry 2000], and CÉU provides an equivalent `async/thread` primitive (which are out of the scope of this paper).

Since loops in CÉU must contain `await` statements, reactions run in bounded time, which guarantees that successive calls to `dispatch` never block the scheduler for long. However, the code generation phase does not inspect C calls and also has no extra analysis such as for worst-case reaction times [Boldt et al. 2008; Li et al. 2005].

Single-thread dispatching may not be suitable for hard real-time activities. In a previous work [Sant'Anna et al. 2013], we measure how synchronous lengthy computations in C (e.g., hashing and compression) can block the scheduler and affect higher-priority activities such as a radio driver. In such cases, the system requires careful testing to avoid undersized hardware deployment.

3.6. Interaction with the Environment

As a reactive language, the execution of programs in CÉU is guided entirely by the occurrence of external input events. The binding for a specific platform (environment) is responsible for calling hook functions in the API of the runtime of CÉU whenever an external event occurs. However, the binding must never interleave or run multiple API calls in parallel. This would break the CÉU sequential/discrete semantics of time.

As an example, Figure 16 shows our binding for *TinyOS* [Hill et al. 2000], which maps system callbacks to input events in CÉU. The file `ceu_app.h` (ln. 3) contains all definitions for the compiled CÉU program, which are further queried through `#ifdef`'s. The file `ceu_app.c` (ln. 4) contains the runtime of CÉU with the scheduler and dispatcher pointing to the labels defined in the program. The callback `Boot.booted` (ln. 6–11) is called by TinyOS on startup, so we initialize CÉU inside it (ln. 7). If the CÉU program uses timers, we also start a periodic timer (ln. 8–10) that triggers callback `Timer.fired` (ln. 13–17) every 10 milliseconds and advances the wall-clock time of CÉU (ln. 15)[4]. The remaining lines map pre-defined TinyOS events that can be used in CÉU programs, such as the light sensor (ln. 19–23) and the radio transceiver (ln. 25–36). The scheduler

[4]We also offer a mechanism to start the underlying timer on demand to avoid the "battery unfriendly" 10ms polling.

```
implementation
{
    #include "ceu_app.h"
    #include "ceu_app.c"

    event void Boot.booted () {
        ceu_init();
#ifdef CEU_WCLOCKS
        call Timer.startPeriodic(10);
#endif
    }

#ifdef CEU_WCLOCKS
    event void Timer.fired () {
        ceu_wclock(10000);
    }
#endif

#ifdef EVT_PHOTO_READDONE
    event void Photo.readDone (int val) {
        ceu_go(EVT_PHOTO_READDONE, &val);
    }
#endif

#ifdef EVT_RADIO_SENDDONE
    event void RadioSend.sendDone (message_t* msg) {
        ceu_go(EVT_RADIO_SENDDONE, &msg);
    }
#endif

#ifdef EVT_RADIO_RECEIVE
    event message_t* RadioReceive.receive (message_t* msg) {
        ceu_go(EVT_RADIO_RECEIVE, &msg);
        return msg;
    }
#endif

    <...>   // other events
}
```

Fig. 16. The *TinyOS* binding for CÉU. This platform-dependent template includes the *C* files generated from the original application in CÉU (`ceu_app.h` and `ceu_app.c`) for the *code generation phase*.

of the TinyOS is already synchronous by default and always execute `event` handlers atomically, hence, the API calls to CÉU are properly serialized.

3.7. The Terra Virtual Machine

Terra is a system for programming wireless sensor network applications which uses CÉU as its scripting language [Branco et al. 2015]. Figure 17 shows the three basic elements of Terra: CÉU as the scripting language, a set of customized pre-built components, and the embedded virtual-machine engine which can disseminate and install bytecode images dynamically. This approach aims to combine the flexibility of remotely uploading code with the expressiveness and safety guarantees of CÉU.

The main difference between the standard *C* back end and the Terra *VM* is the *code generation phase*, which here outputs assembly instructions for the *VM*, instead of statements in *C*. To reduce the memory footprint of applications, the *VM* includes special instructions for complex and recurrent operations from the runtime of CÉU, such as handling events and trails.

In Terra, CÉU scripts cannot execute arbitrary *C* code, instead, they rely on pre-built components that can be customized for different application domains. Considering the domain of sensor networks, Terra already provides components organized in four areas: radio communication, group management, data aggregation, and local op-

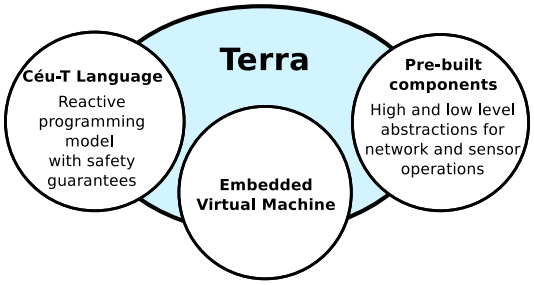

Fig. 17. Terra programming system basic elements.

```
// Output events
output void REQUEST_TEMPERATURE;
output int  REQUEST_SEND;    // sends int value

// Input events
input int  TEMPERATURE_DONE; // recvs int value
input void SEND_DONE;

// System calls
function int getRadioID (void);
```

[a]

```
1  // Output events
2  void VM.out(int evt_id, void* args) {
3      switch (id){
4          case O_REQUEST_TEMPERATURE:
5              call TINYOS_TEMP.read();
6          <...>; // O_REQUEST_SEND
7      }
8  }
9
10 // Input events
11 event TINYOS_TEMP.done (int val) {
12     VM.enqueue(I_TEMPERATURE_DONE, &val);
13 }
14 <...> // TINYOS_SEND.done
15
16 // System calls
17 void VM.function(int id, void* params) {
18     switch (id) {
19         case F_GET_RADIO_ID:
20             VM.push(TINYOS_NODE_ID);
21     }
22 }
```

[b]

Fig. 18. [a] CÉU interface with customized *VM*. [b] The routine `VM.out` redirects all output events to the corresponding OS calls (ln. 1–8). Each *TinyOS* event callback calls `VM.enqueue` for the corresponding input event (ln 10–14). System calls use `VM.push` for immediate return values (ln. 16–22).

erations (e.g., access to sensors and actuators). When creating an instance of the *VM*, the programmer can choose whether or not to include each component, setting different abstraction boundaries for scripts. The generated *VM* has to be preloaded into the embedded devices before they are physically distributed.

The communication between scripts in CÉU and the components in the *VM* is mostly through events: scripts `emit` requests through `output` events and `await` answers through `input` events. Terra also provides system calls for initialization and configuration of components (e.g., *getters* and *setters*). Figure 18.a shows a CÉU interface with the available functionality for a customized *VM* (with temperature and radio components). Figure 18.b shows the associated bindings for output events (ln. 1–8), input events (ln. 10–14), and system calls (ln. 16–22). Note that all applications for the customized *VM* must comply with the same interface. In contrast, the template-based *C* back end (illustrated in Figure 16) allows applications to choose all possible combinations of functionalities from the underlying platform at compile time.

4. RELATED WORK

CÉU has a strong influence from Esterel, embracing its disciplined synchronous-reactive model with support for lexical composition of lines of execution. However,

there are fundamental semantic differences that prevents the design of CÉU as pure extensions to Esterel. In particular, Esterel has a notion of time similar to digital circuits in which multiple signals can be active at a clock tick. In CÉU, instead of clock ticks, the occurrence of a single external event that defines a time unit. CÉU also distinguishes external events from stack-based internal events, which provide a limited form of coroutines supporting reactive statements.

The event-driven approach of CÉU is widespread [Ousterhout 1996] and popular in many software communities, such as web frameworks (e.g., *jQuery* [Chaffer 2009] and *Node.js* [Tilkov and Vinoski 2010]), GUI toolkits (e.g., *Tcl/Tk* [Ousterhout 1991] and *Java Swing* [Eckstein et al. 1998]), and Games [Nystrom 2014]. Like CÉU, event-driven programming is essentially synchronous, i.e., events go through a queue and are dispatched sequentially and atomically to prevent race conditions. We believe that for software design, this approach is more familiar to programmers and simplifies the reasoning about concurrency. For instance, the uniqueness of external events in CÉU is a prerequisite for the temporal analysis that enables safe shared-memory concurrency.

Many synchronous languages have been designed to interoperate with *C*, such as *Reactive C* [Boussinot 1991], Protothreads [Dunkels et al. 2006], *PRET-C* [Andalam et al. 2010] and *SC* [Von Hanxleden 2009]. They offer Esterel-like parallel compositions with communication via shared variables, relying on deterministic scheduling to preserve determinism. However, it is the responsibility of the programmer to specify the execution order for threads, based on either explicit priorities, or source code lexical order (similar to CÉU). These languages have a tick-based notion of time similar to Esterel, which prevents the event-based temporal analysis of CÉU.

URBI [Baillie 2005] is a reactive scripting language with a rich set of control constructs for time management, event-driven communication, and concurrency. Concurrency is based on stackful coroutines, diverging from our goals regarding resource efficiency and static bounds for memory and execution time.

Esterel has different compilation back ends that synthesizes to software and also to hardware circuits [Dayaratne et al. 2005; Edwards 2003]. Among the software-based approaches, *SAXO–RT* [Closse et al. 2002] is the closest to our implementation with respect to trail allocation and scheduling: the compiler slices programs into "control points" (analogous to our "entry points") and rearranges them into a directed acyclic graph respecting the constructive semantics of Esterel. Then, it flattens the graph into sequential code in *C* suitable for static scheduling.

A number of virtual machines have been proposed for embedded systems. *Darjeeling* [Brouwers et al. 2008] and *TakaTuka* [Aslam et al. 2010] are complete *Java VMs* targeting constrained embedded systems with support for multithreading and garbage collection. Java has antagonistic design choices in comparison to CÉU: it does not impose static bounds on memory usage and execution time, and provides preemptive multithreading which requires synchronization primitives for accessing shared memory. Plummer et al. [Plummer et al. 2006] propose a Esterel-based *VM* with similar design choices to our work. To reduce code size, the *VM* has a specialized instruction set to deal with events and concurrency constructs that are particular to Esterel. However, the proposed *VM* is only a proof of concept, with no support for arithmetic operations, external system calls, or remote reprogramming.

5. CONCLUSION

- queue in the semantics - permeates all aspects - As far as we know, CÉU is the first Esterel-based language to HAVE? queue-based/unique in the semantics of the language and not in the library - stack - deterministic intra-reaction execution - imposed

intra-reaction execution - for inter-reaction commm - for intra-reaction commm - for stmts

We presented the design and implementation of CÉU, a synchronous reactive language inspired by Esterel targeting constrained embedded systems.

CÉU is a concurrency-safe language, employing a static analysis that encompass all control constructs and ensures that the high degree of concurrency in embedded systems does not pose safety threats to applications. As a summary, the following safety properties hold for all programs that successfully compile in CÉU: time and memory-bounded reactions to the environment (except for external system calls), no race conditions in shared memory, reliable abortion for activities handling resources, and automatic synchronization for timers. These properties are usually desirable in embedded applications and are guaranteed as preconditions in CÉU by design.

CÉU is a resource-efficient language suitable for constrained embedded systems. The reference implementation compiles to portable event-driven code in *C*, with no special requirements for OS threads or per-trail data stacks. The *VM* implementation uses the same front end and imposes no extra restrictions, being equally suitable for constrained systems.

CÉU is a practical language with expressive control constructs, such as lexically scoped parallel compositions, convenient first-class timers, and a unique stack-based signaling mechanism. Programs interoperate seamlessly with *C*, and can take advantage of existing libraries, lowering the entry barrier for adoption. CÉU has an open source implementation and bindings for *TinyOS*, *Arduino*, and the *SDL* graphical library.[5]

For the past three years, we have been teaching CÉU for undergraduate and graduate students in research projects and two hands-on courses on *distributed systems* and *reactive programming*. Our experience shows that students take advantage of the sequential-imperative style of CÉU and can implement non-trivial concurrent applications in a few weeks.

REFERENCES

Sidharta Andalam and others. 2010. Predictable multithreading of embedded applications using PRET-C. In *Proceeding of MEMOCODE'10*. IEEE, 159–168.

Faisal Aslam and others. 2010. Optimized java binary and virtual machine for tiny motes. In *Distributed Computing in Sensor Systems*. Springer, 15–30.

Jean-Christophe Baillie. 2005. Urbi: Towards a universal robotic low-level programming language. In *International Conference on Intelligent Robots and Systems*. IEEE, 820–825.

Albert Benveniste and others. 2003. The synchronous languages twelve years later. In *Proceedings of the IEEE*, Vol. 91. 64–83.

Alexander Bernauer and Kay Römer. 2013. A Comprehensive Compiler-Assisted Thread Abstraction for Resource-Constrained Systems. In *Proceedings of IPSN'13*. Philadelphia, USA.

Gérard Berry. 1993. Preemption in Concurrent Systems.. In *FSTTCS (LNCS)*, Vol. 761. Springer, 72–93.

Gérard Berry. 1996. The Constructive Semantics of Pure Esterel. (1996).

Gérard Berry. 2000. *The Esterel-V5 Language Primer*. CMA and Inria, Sophia-Antipolis, France. Version 5.10, Release 2.0.

Gérard Berry and Ellen Sentovich. 2001. Multiclock esterel. In *Correct Hardware Design and Verification Methods*. Springer, 110–125.

Marian Boldt, Claus Traulsen, and Reinhard von Hanxleden. 2008. Worst case reaction time analysis of concurrent reactive programs. *Electronic Notes in Theoretical Computer Science* 203, 4 (2008), 65–79.

Frédéric Boussinot. 1991. Reactive C: An extension of C to program reactive systems. *Software: Practice and Experience* 21, 4 (1991), 401–428.

Frédéric Boussinot. 1998. SugarCubes implementation of causality. (1998).

[5]Website of CÉU: http://www.ceu-lang.org/

F. Boussinot and R. de Simone. 1991. The Esterel language. *Proc. IEEE* 79, 9 (Sep 1991), 1293–1304.

Adriano Branco, Francisco Sant'anna, Roberto Ierusalimschy, Noemi Rodriguez, and Silvana Rossetto. 2015. Terra: Flexibility and Safety in Wireless Sensor Networks. *ACM Trans. Sen. Netw.* 11, 4, Article 59 (Sept. 2015), 27 pages. DOI:http://dx.doi.org/10.1145/2811267

Niels Brouwers, Peter Corke, and Koen Langendoen. 2008. Darjeeling, a Java compatible virtual machine for microcontrollers. In *Proceedings of the ACM/IFIP/USENIX Middleware'08 Conference Companion*. ACM, 18–23.

Jonathan Chaffer. 2009. *Learning JQuery 1.3: Better Interaction and Web Development with Simple JavaScript Techniques*. Packt Publishing Ltd.

Etienne Closse and others. 2002. Saxo–RT: Interpreting esterel semantic on a sequential execution structure. *Electronic Notes in Theoretical Computer Science* 65, 5 (2002), 80–94.

Sajeewa Dayaratne and others. 2005. Direct Execution of Esterel Using Reactive Microprocessors. In *Proceedings of SLAP'05*.

Dunkels and others. 2006. Protothreads: simplifying event-driven programming of memory-constrained embedded systems. In *Proceedings of SenSys'06*. ACM, 29–42.

Robert Eckstein, Marc Loy, and Dave Wood. 1998. *Java swing*. O'Reilly & Associates, Inc.

Stephen A. Edwards. 1999. Compiling Esterel into sequential code. In *7th International Workshop on Hardware/Software Codesign*. ACM, 147–151.

Stephen A. Edwards. 2003. Tutorial: Compiling concurrent languages for sequential processors. *ACM Transactions on Design Automation of Electronic Systems* 8, 2 (2003), 141–187.

Stephen A. Edwards. 2005. Using and Compiling Esterel. MEMOCODE'05 Tutorial. (July 2005).

Wahbi Haribi. 2012. Compiling Esterel for Multi-Core Execution. *Synchrone Sprachen* (2012), 45.

Hill and others. 2000. System architecture directions for networked sensors. *SIGPLAN Notices* 35 (November 2000), 93–104. Issue 11.

Marcin Karpinski and Vinny Cahill. 2007. High-Level Application Development is Realistic for Wireless Sensor Networks. In *Proceedings of SECON'07*. 610–619.

Oliver Kasten and Kay Römer. 2005. Beyond Event Handlers: Programming Wireless Sensors with Attributed State Machines. In *Proceedings of IPSN '05*. 45–52.

Xin Li, Jan Lukoschus, Marian Boldt, Michael Harder, and Reinhard Von Hanxleden. 2005. An Esterel processor with full preemption support and its worst case reaction time analysis. In *Proceedings of the 2005 international conference on Compilers, architectures and synthesis for embedded systems*. ACM, 225–236.

Ana Lúcia De Moura and Roberto Ierusalimschy. 2009. Revisiting coroutines. *ACM TOPLAS* 31, 2 (Feb. 2009), 6:1–6:31.

Robert Nystrom. 2014. *Game Programming Patterns*. Genever Benning.

ORACLE. 2011. Java Thread Primitive Deprecation. http://docs.oracle.com/javase/6/docs/technotes/guides/concurrency/threadPrimitiveDeprecation.html (accessed in Aug-2014). (2011).

John K. Ousterhout. 1991. An X11 Toolkit Based on the Tcl Language.. In *USENIX Winter*. 105–116.

John K. Ousterhout. 1996. Why Threads Are A Bad Idea (for most purposes). (January 1996).

Becky Plummer, Mukul Khajanchi, and Stephen A. Edwards. 2006. An Esterel virtual machine for embedded systems. In *International Workshop on Synchronous Languages, Applications, and Programming (SLAP'06)*. Citeseer, Vienna, Austria.

Dumitru Potop-Butucaru and others. 2005. The Synchronous Hypothesis and Synchronous Languages. In *Embedded Systems Handbook*, R. Zurawski (Ed.).

Partha S Roop, Zoran Salcic, and MW Dayaratne. 2004. Towards direct execution of Esterel programs on reactive processors. In *Proceedings of the 4th ACM international conference on Embedded software*. ACM, 240–248.

Francisco Sant'Anna and others. 2013. Safe System-level Concurrency on Resource-Constrained Nodes. In *Proceedings of SenSys'13*. ACM.

Klaus Schneider and Michael Wenz. 2001. A new method for compiling schizophrenic synchronous programs. In *Proceedings of the 2001 international conference on Compilers, architecture, and synthesis for embedded systems*. ACM, 49–58.

Ellen M Sentovich. 1997. Quick conservative causality analysis. In *System Synthesis, 1997. Proceedings., Tenth International Symposium on*. IEEE, 2–8.

Thomas R Shiple, Gerard Berry, and Hemé Touati. 1996. Constructive analysis of cyclic circuits. In *European Design and Test Conference, 1996. ED&TC 96. Proceedings*. IEEE, 328–333.

Olivier Tardieu and Robert De Simone. 2004. Curing schizophrenia by program rewriting in Esterel. In *Formal Methods and Models for Co-Design, 2004. MEMOCODE'04. Proceedings. Second ACM and IEEE International Conference on*. IEEE, 39–48.

Stefan Tilkov and Steve Vinoski. 2010. Node.js: Using JavaScript to build high-performance network programs. *IEEE Internet Computing* 6 (2010), 80–83.

Reinhard Von Hanxleden. 2009. SyncCharts in C: a proposal for light-weight, deterministic concurrency. In *Proceedings of the seventh ACM international conference on Embedded software*. ACM, 225–234.

Simon Yuan, Li Hsien Yoong, and Partha S Roop. 2011. Compiling Esterel for multi-core execution. In *Digital System Design (DSD), 2011 14th Euromicro Conference on*. IEEE, 727–735.

Jeong-Han Yun, Chul-Joo Kim, Seonggun Kim, Kwang-Moo Choe, and Taisook Han. 2013. Detection of harmful schizophrenic statements in esterel. *ACM Transactions on Embedded Computing Systems (TECS)* 12, 3 (2013), 80.