

The Design and Implementation of the Synchronous Language Céu

Francisco Sant'Anna, Departamento de Informática e Ciência da Computação, UERJ

Roberto Ierusalimsky, Departamento de Informática, PUC-Rio

Noemi Rodriguez, Departamento de Informática, PUC-Rio

Silvana Rossetto, Departamento de Ciência da Computação, UFRJ

Adriano Branco, Departamento de Informática, PUC-Rio

Céu is a synchronous language inspired by Esterel with simpler semantics and more fine-grained control over program execution. Céu adopts an event-triggered notion of time that enables compile-time temporal analysis to detect conflicting statements, resulting in deterministic and concurrency-safe programs. Using Esterel as a base, we discuss the particularities of our design, such as stack-based internal events, the temporal analysis, safe integration with *C*, and first-class timers. We also present two implementation back ends: one aiming for resource efficiency and interoperability with *C*, and another as a virtual machine that allows remote reprogramming.

Additional Key Words and Phrases: Concurrency, Determinism, Embedded Systems, Esterel, Synchronous, Reactivity

1. INTRODUCTION

An established alternative to *C* in the field of embedded systems is the family of reactive synchronous languages [Benveniste et al. 2003]. Two major styles of synchronous languages have evolved: in the *control-imperative* style, programs are structured with control flow primitives, such as parallelism, repetition, and preemption; in the *dataflow-declarative* style, programs can be seen as graphs of values, in which a change to a value is propagated through its dependencies without explicit programming. Considering the control-based languages, Esterel [Boussinot and de Simone 1991] was the first to appear and succeed, influencing a number of embedded languages, such as *Reactive-C* [Boussinot 1991], *OSM* [Kasten and Römer 2005], *Sync-C* [Von Hanxleden 2009], and *PRET-C* [Andalam et al. 2010].

Despite its success and influence, Esterel has an overly complex semantics that requires careful static analysis such as to detect and refuse programs with *causality* and *schizophrenia* problems [Berry 1996; Shiple et al. 1996; Sentovich 1997; Boussinot 1998; Schneider and Wenz 2001; Tardieu and De Simone 2004; Edwards 2005; Yun et al. 2013]. The complex semantics not only challenges the analysis and compilation of programs, but also results in incompatible and non-compliant implementations. Above all, it also affects the programmer's understanding about the code, which, ultimately, has to solve the errors when facing corner cases. Another drawback of the Esterel semantics consists of loose and non-deterministic execution for intra-reaction statements. Non determinism prevents threads to interact with stateful system calls safely and makes shared-memory concurrency not as straightforward as reading and writing to shared variables.

In this work, we present Céu, a new programming language that inherits the synchronous and imperative mindset of Esterel but diverges in fundamental semantic aspects. Overall, Céu has a simple semantics with fine-grained control for intra-reaction execution, and is amenable to a temporal analysis that improves safety. The list that follows summarizes the contributions behind the design of Céu:

- Unique and queue-based external events, which define the notion of time in Céu.
- Stack-based internal events for intra-reaction communication, which also provides a limited form of coroutines.
- Static temporal analysis to detect suspicious concurrent statements.

- Safe integration with *C*, which enforces finalization for external resources.
- First-class timers with dedicated syntax and automatic synchronization.

We also present a lightweight single-threaded implementation of CÉU with two back ends: one aiming for resource efficiency and interoperability with *C*, and another as a virtual machine that allows remote reprogramming. Our implementations target resource-constrained devices, such as *Arduino* and *MICAz* sensor nodes based on 8-bit microcontrollers¹, showing a practical aspect of our simple semantics.

In previous work [Sant’Anna et al. 2013b; Branco et al. 2015], we employed CÉU in the context of wireless sensor networks, developing a number of applications, protocols, and device drivers. We evaluated the expressiveness of CÉU in comparison to hand-crafted event-driven code in *C* and attested a reduction in source code size (around 25%) with a small increase in memory usage (around 5–10% for *text* and *data*) [Sant’Anna et al. 2013b]. For the *VM* back end, applications have a bytecode footprint in the order of hundreds of bytes which can be transmitted over the air in a few packets [Branco et al. 2015].

The rest of the paper is organized as follows: Section 2 discusses the design of CÉU, focusing on the fundamental differences to Esterel. Section 3 presents the *C* and *VM* implementation back ends. Section 4 discusses other synchronous languages targeting embedded systems. Section 5 concludes the paper.

2. THE DESIGN OF CÉU

CÉU is a synchronous reactive language inspired by Esterel in which programs advance in a sequence of discrete reactions to external events. Like Esterel, CÉU is designed for control-intensive applications, supporting concurrent lines of execution, known as *trails*, and broadcast communication through events. Internal computations within a reaction (e.g. expressions, assignments, and system calls) are considered to take no time in accordance with the synchronous hypothesis [Potop-Butucaru et al. 2005]. An *await* is the only statement that halts a running reaction and allows a program to advance in this discrete notion of time. To ensure that reactions run in bounded time and programs always progress, loops are statically required to contain at least one *await* statement in all possible paths [Sant’Anna et al. 2013b; Berry 2000]. CÉU shares the same limitations with Esterel and synchronous languages in general: computations that run in unbounded time (e.g., cryptography, image processing) do not fit the zero-delay hypothesis, and cannot be directly implemented.

Figure 1 illustrates the syntactic similarities between the languages, showing side-by-side the implementations in Esterel [a] and CÉU [b] for the following control specification: “*Emit an output *O* as soon as two inputs *A* and *B* have occurred. Reset this behavior each time the input *R* occurs*” [Berry 2000]. The first phrase of the specification, awaiting and emitting the events, is translated almost identically in the two languages (ln. 5–10, in both implementations), as Esterel’s ‘||’ and CÉU’s *par/and* constructs are equivalent. For the second phrase, the reset behavior, the Esterel version uses a *abort-when* (ln. 4–11) which, in this case, serves the same purpose of CÉU’s *par/or* (ln. 4–13): the occurrence of event *R* aborts the awaiting statements in parallel and restarts the enclosing *loop*.

In the subsections that follow, we discuss the main differences between CÉU and Esterel: Unique and queue-based external events (2.1); Stack-based internal events (2.2); Static temporal analysis for concurrent statements (2.3); Safe integration with *C* (2.4); and First-class synchronized timers (2.5). We finish with a summary of our

¹Both *Arduino* and *MICAz* use the 8-bit *ATmega328* microcontroller with 32K of FLASH and 2K of SRAM.

<pre> 1 input A, B; 2 output O; 3 loop 4 abort 5 [6 await A 7 8 await B 9]; 10 emit O 11 when R 12 end 13 14 . </pre>	<pre> 1 input void A, B; 2 output void O; 3 loop do 4 par/or do 5 par/and do 6 await A; 7 with 8 await B; 9 end 10 emit O; 11 with 12 await R; 13 end 14 end </pre>
[a] Esterel	[b] CÉU

Fig. 1. A control specification implemented in Esterel and CÉU: “Emit *O* after *A* and *B*, resetting each *R*”

design (2.6). We present the formal specification of the semantics of CÉU in a separate work [Sant’Anna 2013].

2.1. Unique and Queue-Based External Events

Esterel defines time as a discrete sequence of logical unit instants or “ticks”. At each tick, the program reacts to an arbitrary number of simultaneous input events from the environment. In contrast, CÉU defines time as a discrete sequence of reactions to unique input events. At each input event, which constitutes a logical unit of time, the program reacts exclusively to it. The event-triggered execution of a program in CÉU is as follows [Sant’Anna et al. 2013b]:

- (1) The program initiates the “boot reaction” in a single trail (but parallel constructs may create new trails).
- (2) Active trails execute, one after the other, until they await or terminate. This step is named a *reaction chain*, and always runs in bounded time.
- (3) The program goes idle and the environment takes control.
- (4) On the occurrence of a new external input event, the environment awakes *all* trails awaiting that event. It then goes to step 2.

A program must react to an occurring event completely before handling the next. Based on the synchronous hypothesis, a program takes a negligible time on step 2 and is always idle on step 3. In practice, if a new external input event occurs while a reaction chain is running, it is enqueued to occur in the next reaction.

Figure 2 compares the discrete notions of time in two variations of Esterel and in CÉU. The box *Real World* shows event occurrences over a continuous timeline. The other boxes show discrete logical units of time in which the same occurring events fit in different ways.

- [Box-1]: *Esterel with fixed-length ticks*, in which reactions must fit [Li et al. 2005]. We assume a reaction $R(\text{boot})$ at tick-0 which happens before any input. The input *A* “physically” occurs during the boot reaction but, because time is discrete, its corresponding reaction only executes in the next tick. Note that $R(A)$ takes more time than tick-1, causing a *timing violation* [Li et al. 2005] that invades tick-2. The events *B* and *C* occur during tick-1 and are delayed to happen *simultaneously* at tick-2 with $R(B+C)$. Since no new events occur during tick-2, the CPU stays idle during the whole tick-3. Finally, one instance of event *D* and two instances of event *E* occur during the idle tick-3. However, only one occurrence of *E* is considered in $R(D+E)$.




Fig. 2. The discrete notions of time in Esterel and CÉU.

- [Box-2]: *Esterel with variable-length ticks*, which are adjusted to the CPU time the corresponding reaction takes to complete [Roop et al. 2004]. This approach avoids the time violation for $R(A)$ and also results in smaller idle periods. For instance, the occurrence of D interrupts the idle tick-3 to react alone as $R(D)$ on tick-4. Similarly to the fixed-tick approach, only one of the two simultaneous occurrences of E is considered on tick-5, now because $R(D)$ takes too long.
- [Box-3]: *CÉU with unique and queue-based input events*. We also assume a reaction $R(\text{boot})$ before any input. Because the occurrence of event A is unique during tick-0, the behavior in CÉU is similar to Box-2 for the first two reactions (boot-0 and $A-1$). However, CÉU does not consider the events B and C as simultaneous, and handles each in subsequent reactions $R(B)$ and $R(C)$. We assume the CPU times for $R(B+C)$ in Esterel and $R(B)+R(C)$ in CÉU to be roughly the same. This way, the first idle period in Box-2 and Box-3 coincide. Finally, CÉU recognizes and reacts to the two instances of E independently, which are handled in sequence.

We decided for the unique and queue-based semantics in CÉU for the reasons that follow:

- **A “tick” is too abstract and imprecise:** Outside the domain of hardware specification, a tick has no natural counterpart in the real world. Also, since ticks require no time regularity [Berry and Sentovich 2001], the two approaches for Esterel in

Figure 2 are legitimate, but lead to different behaviors for the same sequence of inputs.

- **Events are never absolutely simultaneous:** From a rigorous point of view, event occurrences are infinitesimal, having zero probability of being simultaneous. This way, we believe that the notion of simultaneity should not be imposed by the language, but defined explicitly for each use case. In the case of Esterel, simultaneity depends on the imprecise length of discrete ticks. For instance, in Figure 2 (box-1 and box-2), Esterel considers the events *B* and *C* to be simultaneous, even though *A* and *B* “physically” happen much closer to one another. In Section 2.5, after introducing internal events and first-class timers, we show an example to detect simultaneous button clicks in CÉU.
- **Unique input events imply mutual exclusion:** Reactions to multiple events in CÉU never overlap because they are atomic. Automated mutual exclusion simplifies the reasoning about concurrency and is a prerequisite for the temporal analysis to be discussed in Section 2.3.

The synchronous hypothesis for CÉU holds if the reactions run faster than the rate of incoming input events. Otherwise, the application continuously accumulates delays between the real occurrence and actual reaction of a given event. This is also the case for the variable-length-tick approach of Esterel, since the more inputs to handle, the longer the reaction takes, and the more inputs accumulates for subsequent ticks. For the fixed-length-tick approach of Esterel, a breach in the synchronous hypothesis causes timing violations, which can be avoided with *worst case reaction time* analysis to infer an appropriate value for the tick length [Li et al. 2005].

A limitation of event-triggered execution is that all program behavior must be purely reactive, given that no code can execute in the absence of inputs. Tick-triggered execution allows for active behavior, since code can execute regularly on every tick. Although CÉU supports asynchronous active execution [Sant’Anna et al. 2012], we still consider it as purely reactive since its synchronous core cannot express active behavior.

2.2. Stack-Based Internal Events

Esterel makes no semantic distinctions between internal and external signals. In CÉU, in contrast with queue-based external events, internal events follow a stack-based execution policy similar to subroutine calls in typical programming languages. Figure 3 illustrates the use of internal signals (events) in Esterel [a] and CÉU [b]. In Esterel, when *A* occurs, *B* is emitted (ln. 4–5) and both events become active, resulting in the invocation of *f()* and *g()* in no particular order (ln. 6,9). In CÉU, the occurrence of *A* makes the program behave as follows:

- (1) 1st trail awakes, broadcasts *b*, and pauses (ln. 4–5).
- (2) 2nd trail awakes, calls *g()*, and terminates (ln. 8–9). (*No other trails awake to b.*)
- (3) 1st trail (on top of the stack) resumes, calls *f()*, and terminates (ln. 5–6).
- (4) Both trails have terminated, so the *par*/and rejoins, and the program also terminates.

Internal events provide a limited form of subroutines, as depicted in Figure 4. The subroutine *inc* is defined as a loop (ln. 3–6) that continuously awaits its identifying event (ln. 4), incrementing the value passed as reference (ln. 5). A trail in parallel (ln. 8–11) invokes the subroutine in reaction to event *A* through an *emit* (ln. 10). Given the stacked execution for internal events, the calling trail pauses, the subroutine awakes (ln. 4), runs its body (yielding *v=2*), loops, and awaits the next “call” (ln. 4, again). Only after this sequence, the calling trail resumes and passes the assertion test (ln. 11).

CÉU also supports nested *emit* invocations for internal events. For instance, the code in Figure 4 that executes after the *await inc* (ln. 4) could *emit* another event, creating

<pre> 1 input A; // external 2 signal B; // internal 3 [[4 await A; 5 emit B; 6 call f(); 7 8 await B; 9 call g(); 10]] </pre>	<pre> 1 input void A; // external (in uppercase) 2 event void b; // internal (in lowercase) 3 par/and do 4 await A; 5 emit b; 6 -f(); 7 with 8 await b; 9 -g(); 10 end </pre>
[a] Esterel	[b] CÉU

Fig. 3. Internal signals (events) in Esterel and CÉU: similar syntax, but different semantics.

```

1 event int* inc; // subroutine 'inc'
2 par/or do
3   loop do // definitions are loops
4     var int* p = await inc;
5     *p = *p + 1;
6   end
7 with
8   var int v = 1;
9   await A;
10  emit inc => &v; // call 'inc'
11  _assert(v==2); // after return
12 end

```

Fig. 4. Subroutine inc is defined in a loop (ln. 3–6), in parallel with the caller (ln. 8–11).

a new level in the stack. The runtime stack constitutes fine-grained micro reactions, one on top of the other, all inside the same reaction to an external event.

On the one hand, this form of subroutines has a significant limitation that it cannot express recursive calls: an `emit` to itself is always ignored, given that a running body cannot be awaiting itself. On the other hand, this very same limitation brings some important safety properties to subroutines: first, they are guaranteed to react in bounded time; second, memory for locals is also bounded, not requiring data stacks. Also, this form of subroutines can use the other primitives of CÉU, such as parallel compositions and the `await` statement. In particular, they await keeping context information such as locals and the program counter, similarly to coroutines [Moura and Ierusalimsky 2009]. In a previous work, we build on top of internal events other advanced control mechanisms, such as resumable exceptions and reactive variables [Sant’Anna et al. 2013a].

Another distinction regarding event handling in comparison to CÉU is that Esterel supports same-cycle bi-directional communication [Edwards 1999], i.e., two threads can react to one another during the same cycle due to mutual signal dependency. CÉU has a different take, posing a tradeoff that an `await` is only valid for the next reaction, i.e., if an `await` and `emit` occur simultaneously in parallel trails, the `await` does not awake. These *delayed awaits* avoid corner cases of instantaneous termination and re-execution of statements in the same reaction (known as *schizophrenic statements* [Berry 1996]).

The example in Figure 5 illustrates delayed awaits, which prevents infinite execution by design. Both sides of the `par/or` have an `await` statement (ln. 4,7), which makes the enclosing `loop` non instantaneous (ln 2–9). However, if the `emit` (ln. 6) could awake the `await` (ln. 4) in the same reaction that reaches them, the `par/or` would terminate and restart the `loop` instantaneously, resulting in infinite execution. In atypical scenarios requiring immediate awake, delayed awaits can be circumvented by manually copying or transforming the code to execute on awake. On the one hand, we transfer

```

1 event void e,f;
2 loop do
3   par/or do
4     await e;
5     with
6       emit e;      // w/o delayed awaits, the emit awakes 1st trail
7       await f;     // and restarts the loop instantaneously
8     end
9   end

```

Fig. 5. Delayed awaits prevents re-execution of statements by design.

```

1 input void A, B;
2 var int x = 1;
3 par/and do
4   await A;
5   x = x + 1;
6 with
7   await B;
8   x = x * 2;
9 end

```

[a] Accesses to x are safe

```

1 input void A;
2 var int y = 1;
3 par/and do
4   await A;
5   y = y + 1;
6 with
7   await A;
8   y = y * 2;
9 end

```

[b] Accesses to y are unsafe

Fig. 6. Shared-memory concurrency in CÉU: example [a] is safe because the trails access x atomically in different reactions; example [b] is unsafe because both trails access y in the same reaction.

the burden of dealing with these corner cases to the programmer. On the other hand, we simplify the semantics of the language and eliminate the need for complex analysis to deal with schizophrenic statements.

2.3. Static Temporal Analysis for Concurrent Statements

Embedded applications make extensive use of global memory and shared resources, such as through memory-mapped registers and system calls to device drivers. Hence, an important goal of CÉU is to ensure a reliable behavior for programs with concurrent lines of execution sharing memory and interacting with the environment.

Esterel is only deterministic with respect to external behavior: “the same sequence of inputs always produces the same sequence of outputs” [Berry 2000]. However, the execution order for operations within a reaction is non-deterministic: “if there is no control dependency, as in `<<call f1() || call f2()>>`, the order is unspecified and it would be an error to rely on it” [Berry 2000]. A number of Esterel-based synchronous languages have intra-reaction determinism by enforcing an arbitrary execution order for statements in multiple lines of execution (*Reactive C* [Boussinot 1991], *Prothreads* [Dunkels et al. 2006], *SOL* [Karpinski and Cahill 2007], *SC* [Von Hanxleden 2009], and *PRET-C* [Andalam et al. 2010]). CÉU also takes the deterministic approach and, when multiple trails are active during the same reaction, they are scheduled in lexical order, i.e., in the order they appear in the program source code.

Even so, we consider that enforcing an arbitrary execution order can be misleading in some cases. For instance, consider the two examples in Figure 6, both defining a shared variable (ln. 2), and assigning to it in parallel trails (ln. 5, 8). In the example [a], the two assignments to x execute in reactions to different events A and B only, which by definition, cannot occur simultaneously (Section 2.1). Hence, for the sequence A→B, x becomes 4 ((1+1)*2), while for B→A, x becomes 3 ((1*2)+1). In the example [b], the two assignments to y are simultaneous because they execute in reaction to the same event A. Since CÉU employs lexical order for intra-reaction statements, the execution is still deterministic, and y always becomes 4 ((1+1)*2). However, an (apparently innocuous)

<pre> 1 native do 2 #define NUM 10 3 void f (void) { <...> } 4 void g (int v) { <...> } 5 int id (int v) { <...> } 6 end 7 par/and do 8 _f (); 9 with 10 _g (_id (_NUM)); 11 end </pre>	<pre> native @const _NUM; native @pure _id (); native @safe _f () with _g (); </pre>
[a] Definitions and uses of symbols	[b] Annotations for the symbols in [a]

Fig. 7. The unsafe program in [a] only compiles with the annotations in [b].

change in the order of trails modifies the semantics of the program, which we consider unsafe.

To mitigate this threat, CÉU performs a temporal analysis at compile time to detect concurrent accesses to shared variables: if a variable is written in a trail segment, then a concurrent trail segment cannot read or write to that variable, nor dereference a pointer of that variable type. Concurrency in CÉU is characterized when two or more trail segments in parallel react to the same input event. A trail segment is a sequence of statements followed by an `await` (or termination). Considering the examples in Figure 6:

- The assignments to `x` ([a]:2,5) **cannot** be concurrent because they are **not** in parallel trails.
- The assignments to `x` ([a]:5,8) **cannot** be concurrent because they **cannot** execute during the same reaction.
- The assignments to `y` ([b]:5,8) **can** be concurrent because they are in parallel trails and **can** execute during the same reaction.

The algorithm for the analysis, which is depicted in Section 3.1, inspects all possible `await` statements that precede a variable access and keep a list with all corresponding awaking events. Then, it checks all accesses in parallel trails to see if they share an awaking event. If it is the case, the compiler warns about the suspicious accesses.

Note that such analysis is only possible due to the uniqueness of input events within reactions. Otherwise, any two trail segments in parallel could be concurrent, even if they react to different input events. Note also that the analysis is optional and does not affect the semantics of the program.

2.4. Safe Integration with C

In CÉU, any identifier prefixed with an underscore is repassed *as is* to the C compiler that generates the final binary. Therefore, access to C is seamless and, more importantly, easily trackable. Similarly to Esterel with the `call` primitive, external calls are assumed to be instantaneous [Berry 2000]. This way, programs should only resort to C for asynchronous functionality, such as non-blocking I/O, or simple struct accessors, but never for control purposes.²

2.4.1. Temporal Analysis. As a safety measure, the temporal analysis of Section 2.3 also considers concurrent calls and accesses to external symbols in C. As an example, the program in Figure 7.a defines four external symbols inside a `native` block with standard declarations in C (ln. 1–6). During the boot reaction, two trails react concurrently inside the `par/and` (ln. 7–11): the first trail calls symbol `_f` (ln. 8), while the second calls

²In CÉU, it is possible to restrict the available C symbols as a compile-time option.

<pre> 1 par do 2 <...> // animate and redraw "background" 3 _redraw(background); 4 with 5 <...> // animate and redraw "foreground" 6 _redraw(foreground); 7 end </pre>	<pre> 1 native do 2 #define redraw_non_commutative redraw 3 end 4 @safe _redraw_non_commutative with 5 _redraw_non_commutative; 6 par do 7 <...> // animate and redraw "background" 8 _redraw_non_commutative(background); 9 with 10 <...> // animate and redraw "fore" 11 _redraw_non_commutative(foreground); 12 end </pre>
--	---

[a] `_redraw` **cannot** be concurrent[b] `_redraw_non_commutative` **can** be concurrent

Fig. 8. Making the non-commutative redrawing calls from [a] to compile in [b].

`_g` and `_id`, and also reads `_NUM` (ln. 10). Since CÉU does not inspect any code in *C*, it opts to complain about suspicious concurrent accesses between `_f` and all symbols in the second trail.

2.4.2. Annotations. The annotations in Figure 7.b provides hints to the compiler about the semantics of the *C* symbols in code [a], which now compiles successfully:

- `_NUM` is a constant symbol, meaning that it is safe to use it concurrently with any other symbol in the program.
- `_id` is a pure function, also meaning that it is safe to call it concurrently with any other symbol in the program.
- Both `_f` and `_g` are impure, but have non-conflicting commutative effects, and can be safely called concurrently.

From our experience, however, we find that programs often need non-commutative concurrent calls. This is the case for logging (e.g., calls to `_printf` in trails in parallel) and for redrawing objects in the screen. Figure 8.a shows an abstract code to animate and redraw the objects background and foreground in trails in parallel. In typical graphical APIs, consecutive calls to `_redraw` overwrites conflicting pixels, which makes them non commutative and prevent the code to compile. However, we want to rely on lexical order to always redraw the background object before the foreground object. Therefore, in Figure 8.b, we redefine `_redraw` to an “intimidating” name `_redraw_non_commutative` (ln. 2, to be explicit about its effect), and annotate it as safe (ln. 4–5) to make the code compile successfully.

2.4.3. Finalization. Esterel’s abort and CÉU’s `par/or` statements provide orthogonal abortion of lines of execution, which is a distinctive feature of synchronous languages in comparison to asynchronous languages [Berry 1993]. However, aborting lines of execution that deal with external resources is still error prone because they may end up in an inconsistent state. For this reason, Esterel and CÉU provide a `finalize` construct to unconditionally execute a series of statements even if the enclosing block is aborted and does not terminate normally.

In the example in Figure 9 in Esterel [a] and CÉU[b], the calls to `lock` and `unlock` represent accesses to an external resource. After we `lock` the resource (ln. [a]:4 and [b]:3), we perform some operations in subsequent reactions to input A (ln. [a]:5–8 and [b]:7–10), and then we `unlock` the resource (ln. [a]:10 and [b]:5). Note that if the aborting input B (ln. [a]:12 and [b]:12) occurs after the `lock` but before the reactions to A, we still want to call `unlock` to safely release the resource. In Esterel and CÉU, the `finalize` clause (ln. [a]:10 and [b]:5) execute automatically if the enclosing block (ln. [a]:3–1 and [b]:3–10) is externally aborted (ln. [a]:12 and [b]:12).

```

1 input A, B;
2 abort
3   finalize
4     call lock();
5     await A;
6     <...>;      // do something
7     await A;
8     <...>;      // do something
9   with
10    call unlock();
11  end
12 when B
13 .

```

[a] Esterel

```

1 input void A, B;
2 par/or do
3   _lock();
4   finalize with
5     _unlock(); // defer execution
6   end
7   await A;
8   <...>;      // do something
9   await A;
10  <...>;      // do something
11 with
12   await B;
13 end

```

[b] CÉU

Fig. 9. Finalization in Esterel and CÉU: after the call to lock, both languages guarantee to call unlock if the enclosing block aborts when B occurs.

```

1 par/or do
2   var _buffer_t msg;
3   <...> // prepare msg
4   finalize
5     _send_request(&msg);
6   with
7     _send_cancel(&msg);
8   end
9   await SEND_ACK;
10 with
11  <...>
12 end
13 .

```

[a] Local resource finalization

```

1 par/or do
2   var _FILE* f;
3   finalize
4     f = _fopen(...);
5   with
6     _fclose(f);
7   end
8   _fwrite(..., f);
9   await A;
10  _fwrite(..., f);
11 with
12  <...>
13 end

```

[b] External resource finalization

Fig. 10. CÉU enforces the use of finalization to prevent *dangling pointers* for local resources and *memory leaks* for external resources.

CÉU goes one step further and actually enforces the use of `finalize` for system calls that deal with pointers representing resources:

- If CÉU **passes** a pointer to a system call, the pointer represents a **local** resource that requires finalization.
- If CÉU **receives** a pointer from a system call, the pointer represents an **external** resource that requires finalization.

CÉU tracks the interaction of system calls with pointers and requires finalization clauses to accompany them. In the example in Figure 10.a, the local variable `msg` (ln. 2) is an internal resource passed as a pointer to `_send_request` (ln. 5), which is an asynchronous call that transmits the buffer in the background. If the block aborts (ln. 11) before receiving an acknowledge from the environment (ln. 9), the local `msg` goes out of scope and the external transmission now holds a *dangling pointer*. The finalization ensures that the transmission also aborts (ln. 7). In the example in Figure 10.b, the call to `_fopen` (ln. 4) returns an external file resource as a pointer. If the block aborts (ln. 12) during the `await A` (ln. 9), the file remains open as a *memory leak*. The finalization ensures that the file closes properly (ln. 6). In both cases, the code does not compile without the `finalize` construct.

Note that the illustrative example in Figure 9 does not manipulate pointers (i.e., the resource is a *singleton*). That case is an example of a bad and unsafe API to expose to CÉU because the compiler will not enforce the use of finalization.

<pre> 1 var int v; 2 await 10ms; 3 v = 1; 4 await 1ms; 5 v = 2; 6 7 8 9 . </pre>	<pre> 1 par/or do 2 await 10ms; 3 <...> // any non-awaiting sequence 4 await 1ms; 5 v = 1; 6 with 7 await 12ms; 8 v = 2; 9 end </pre>
[a]	[b]

Fig. 11. First-class timers in C  U.

2.5. First-Class Timers

Activities that involve reactions to *wall-clock time*³ appear in typical patterns of embedded development, such as timeout watchdogs and sensor samplings. However, the interaction between system clocks and programs is not 100% precise, a fact that is usually ignored in the development process. We define the difference between a requested timeout and the actual expiring time as the *residual delta time* (*delta*). Without explicit manipulation, the recurrent use of timed activities in sequence (or in a loop) may accumulate a considerable amount of deltas that can lead to incorrect behavior in programs.

The `await` statement of C  U supports wall-clock time and handles deltas automatically, resulting in more robust applications. In the example in Figure 11.a, suppose that after the first `await` request, the underlying system gets busy and takes 15ms to check for expiring awaits. The C  U scheduler will notice that the `await 10ms` (ln. 2) has not only already expired, but is delayed with `delta=5ms`. Then, the awaiting trail awakes, sets `v=1` (ln. 3), and invokes `await 1ms` (ln. 4). As the current delta is still higher than the requested timeout (i.e. $5ms > 1ms$), the trail is rescheduled for execution, now with `delta=4ms`.

C  U also considers that time is a physical quantity that can be added and compared. For instance, in the example in Figure 11.b, although the scheduler cannot guarantee that the first trail terminates exactly in 11ms (ln. 2,4), it can at least ensure that the program always terminates with `v=1`. Given that any non-awaiting sequence is considered to take no time in the synchronous model, the first trail (ln. 2–5) is guaranteed to terminate before the second trail (ln. 7–8), because $10 + 1 < 12$. A similar program in a language without first-class support for timers would depend on the execution timings for the code marked as `<...>`, making the reasoning about the execution behavior more difficult.

In Section 2.1, we argue that event occurrences are infinitesimal and can never be absolutely simultaneous. However, the “sensation of simultaneity” is not infinitesimal, but actually increases with the inaccuracy of the observer (e.g., a human being). Therefore, we consider that simultaneity should be defined case by case, and should not be imposed by the language. First-class timers simplify the implementation of application-defined simultaneity. Figure 12 emulates a `middle_click` event (ln. 3) in terms of “simultaneous” occurrences of `LEFT_CLICK` and `RIGHT_CLICK` (ln. 1–2). If both events occur, we emit the internal event `middle_click` (ln. 6–7). However, if one of them occurs and the `200ms` timer expires (ln. 9–10), we abort the whole behavior with the `par/or` (ln. 5) and try again with the enclosing `loop` (ln. 4). In this specification, “simultaneous” means “within 200 milliseconds”, which is a huge amount of time for a language-defined tick. Therefore, a similar implementation for Esterel would not rely on the tick notion of simultaneity either.

³By wall-clock time we mean the passage of time from the real world, measured in hours, minutes, etc.

```

1  input void LEFT_CLICK;
2  input void RIGHT_CLICK;
3  event void middle_click;
4  loop do
5      par/or do
6          AWAIT_AND(LEFT_CLICK, RIGHT_CLICK);
7          emit middle_click;
8      with
9          AWAIT_OR(LEFT_CLICK, RIGHT_CLICK);
10         await 200 ms;
11     end
12 end

```

```

13 #define AWAIT_AND(e1, e2) \
14     par/and do \
15         await e1; \
16     with \
17         await e2; \
18     end \
19 #define AWAIT_OR(e1, e2) \
20     par/or do \
21         await e1; \
22     with \
23         await e2; \
24     end

```

Fig. 12. Application defines that a `middle_click` event occurs whenever both `LEFT_CLICK` and `RIGHT_CLICK` occur within 200 milliseconds. The macros `AWAIT_AND` (ln. 13–18) and `AWAIT_OR` (ln. 19–24) are simple expansions to a `par/and` and `par/or` for better readability.

2.6. Summary

CÉU aims to offer a simpler semantics than Esterel with more determinism and fine-grained control over program execution. The list that follows summarizes the contributions of our design in this direction:

- *Event-triggered notion of time bound to the semantics of the language.* Event-driven programming is popular in many domains, such as server and GUI development. We believe that programmers are more familiar with dealing with events in isolation, which simplifies the reasoning about concurrency. In addition, the uniqueness of external events is a prerequisite for the temporal analysis of CÉU.
- *Deterministic intra-reaction execution and communication.* Determinism in CÉU is “all-inclusive” and does not depend on additional levels of static analysis. It encompasses the whole language, including memory accesses, system calls, and stack-based internal events. Programmers can always figure out which statement executes next, making runtime analysis and debugging easier.
- *Static temporal analysis for concurrent statements.* Although execution is deterministic, the compiler of CÉU still advises about suspicious statements that can react concurrently to the same event, appropriate serialization.
- *Safe integration with C.* CÉU enforces the programmer to deal with pointers representing resources through finalization clauses. When dealing with concurrent system calls, programmers can provide the compiler with annotations to remove false positives from the temporal analysis or to allow non commutative concurrent behavior.
- *First-class timers with dedicated syntax and automatic synchronization.* Given the omnipresence of timers in embedded systems, a dedicated syntax can simplify the development and readability of programs. Furthermore, automatic synchronization releases the programmer from the burden of adjusting timers in sequence and in parallel.

Our synchronous and deterministic approach also leads to some limitations as follows:

- *Execution is purely reactive as result of event-triggered reactions.* Since only event occurrences can start reactions, programs cannot execute proactively in the absence of events. In addition, `await` statements cannot awake in the same reaction they are reached.
- *Reactions must execute in bounded time due to the synchronous hypothesis.* As a synchronous language, CÉU requires CPU times for reactions to be negligible in comparison to the rate of incoming events.

Application	Language	tokens	C��u vs nesC	ROM	C��u vs nesC	RAM	C��u vs nesC
CTP	nesC	383	-23%	18896	9%	1295	2%
	C��u	295		20542		1319	
SRP	nesC	418	-30%	12266	5%	1252	-3%
	C��u	291		12836		1215	
DRIP	nesC	342	-25%	12708	8%	393	4%
	C��u	258		13726		407	
CC2420	nesC	519	-27%	10546	2%	283	3%
	C��u	380		10782		291	

Fig. 13. Resource usage for C    and nesC in the domain of sensor networks.

— *Execution is sequential because of intra-reaction determinism.* The deterministic semantics of C    does not make implicit parallelization easy (to be discussed in Section 3.5).

Nonetheless, we advocate for keeping a tractable synchronous reactive core with support for shared memory concurrency and deterministic execution. To deal with the limitations above (which is not in the scope of this paper), we recommend for memory-isolated parallelizable asynchronous primitives as separate extensions to the synchronous core [Berry et al. 1993].

3. IMPLEMENTATION

The compilation process of programs in C    is composed of three main phases: the *parsing phase* converts the source code in C    to an abstract syntax tree (AST); the *temporal analysis phase* detects inconsistencies in programs, such as unbounded loops and suspicious concurrent statements; the *code generation phase* converts the AST to standard C code and packs it with platform-dependent functionality (e.g., system calls) and the runtime of C   , compiling everything with *gcc* to generate the final binary.

In a previous work [Sant’Anna et al. 2013b], we evaluate the implementation of C    in comparison to a hand-crafted event-driven code base⁴ in nesC (a C variant) [Gay et al. 2003]. Figure 13 compares source size (*number of tokens*), binary size (*ROM*), and memory usage (*RAM*) for a number of standardized network protocols and a radio driver. The small overhead in resource usage shows that the gains in productivity and safety with C    make it a viable alternative to C in the context of constrained embedded systems. Unfortunately, most real applications in Esterel seem to be closed source. In a future work, we plan to cooperate with groups that have more experience and access to real code bases in Esterel for a comprehensive comparison with C   .

In the subsections that follow, we discuss C   -specific implementation details: temporal analysis for determinism (Section 3.1), static memory allocation for data and trails (Sections 3.2 and 3.3), static scheduling and trail finalization (Section 3.4), single-threaded dispatching (Section 3.5), interaction with the environment (Section 3.6), and the VM back end (Section 3.7),

3.1. Temporal Analysis for Concurrent Statements

The compile-time temporal analysis phase detects inconsistencies in C    programs. Here, we focus on the algorithm that detects suspicious concurrent statements, such as accesses to shared variables, as discussed in Section 2.3.

⁴TinyOS repository: <http://github.com/tinyos/tinyos-release/>

<pre> 1 input void A, B; 2 var int y; 3 par/or do 4 await A; 5 y = 1; 6 with 7 await B; 8 y = 2; 9 end 10 await A; 11 y = 3; </pre>	<pre> 1 Stmts I={.} O={A} 2 Decl_y I={.} O={.} 3 ParOr I={.} O={A,B} 4 Stmts I={.} O={A} 5 Await_A I={.} O={A} 6 Set_y I={A} O={A} 7 Stmts I={.} O={B} 8 Await_B I={.} O={B} 9 Set_y I={B} O={B} 10 Await_A I={A,B} O={A} 11 Set_y I={A} O={A} </pre>
[a] A program in CÉU...	[b] ...with corresponding sets I and O .

Fig. 14. A program with a corresponding AST describing the sets I and O . The program is safe because accesses to y in parallel have no intersections for I .

For each node representing a statement in the program AST, we keep the set of input events I (*incoming*) that can lead to the execution of the node, and also the set of input events O (*outgoing*) that can terminate the node. As an example, for the single-statement program `await A`, we have $I = \{\text{boot}\}$ and $O = \{A\}$.

A node inherits the set I from its immediate parent and calculates O according to its type, as follows:

- Nodes that represent expressions, assignments, C calls, and declarations simply reproduce $O = I$, as they do not await;
- An `await E` statement, where E is an external input event, has $O = \{E\}$ (see also internal events below).
- A `break` statement has $O = \{\}$ as it escapes the innermost `loop` and never proceeds to the statement immediately following it (see also `loop` below);
- A *sequence node* ($;$) modifies each of its children to have $I_n = O_{n-1}$, except for $n = 1$ which inherits I from the parent node. The set O for the whole node is copied from its last child, i.e., $O = O_n$.
- A `loop` node includes the output of its body on its own I ($I = I \cup O_{\text{body}}$), as the loop is also reached from its own body. The union of all O from nested `break` statements forms the set O for a loop.
- An `if` node has $O = O_{\text{true}} \cup O_{\text{false}}$, where `true` and `false` are the two clauses of the `if`.
- A parallel composition may terminate from any of its branches, hence $O = O_1 \cup \dots \cup O_n$.
- For internal events, an `await` awakes from any input that leads to any matching `emit` in a trail in parallel:
 - An `await e` has $O = I_{e1} \cup \dots \cup I_{eN}$, where $e1 \dots eN$ are `emit e` statements in trails in parallel.
 - An `emit e` terminates in the same reaction, having $O = I$.

With all sets calculated, we take all pairs of nodes that perform side effects and are in parallel branches, and compare their sets I for intersections. For each pair, if the intersection is not empty, we mark both nodes as suspicious.

The code in Figure 14.a has its corresponding AST and sets I and O in Figure 14.b. The event `.` (dot) represents the “boot” reaction. The assignments to y in parallel (ln. [a]:5,8) have an empty intersection of I (ln. [b]:6,9), hence, they do not conflict. Note that, although the accesses to y in sequence (ln. [a]:5,11) do have an intersection (ln. [b]:6,11), they are not in parallel branches and are also safe.

<pre> 1 input int A, B, C; 2 do 3 var int a = await A; 4 end 5 do 6 var int b = await B; 7 end 8 par/and do 9 await B; 10 with 11 await C; 12 end </pre>	<pre> 1 union { // sequence 2 int a_1; // do_1 3 int b_2; // do_2 4 struct { // par/and 5 int _and_3: 1; 6 int _and_4: 1; 7 }; 8 } MEM ; 9 10 11 12 </pre>
--	---

[a] A program in CÉU...

[b] ...with corresponding memory layout

Fig. 15. A program with blocks in sequence and in parallel, with corresponding memory layout generated by the compiler.

3.2. Static Memory Layout

CÉU favors a fine-grained use of trails, being common to use trails that await a single event and terminate. For this reason, CÉU does not allocate per-trail stacks; instead, all data resides in fixed memory slots—this is true for the program variables as well as for temporary values and runtime flags. Memory for trails in parallel must coexist, while statements in sequence can reuse it. Translating this idea to C is straightforward [Kasten and Römer 2005; Bernauer and Römer 2013]: memory for blocks in sequence are packed in a union, while blocks in parallel, in a struct. CÉU reserves a single static block of memory to hold all memory slots, whose size is the maximum the program uses at any given time. A position in the memory may hold different data (with variable sizes) during runtime. As an example, Figure 15 shows a program with its corresponding memory layout. Each variable is assigned a unique *id* so that variables with the same name can be distinguished (e.g., `a_1` in [b]:2 corresponding to `a` in [a]:3). The do-end blocks and par/and in sequence (ln. [a]:2–4,5–7,8–12) are packed in a union (ln. [b]:2,3,4–7), given that their variables cannot be in scope at the same time, e.g., `MEM.a_1` and `MEM.b_2` (ln. [b]:2,3) can safely share the same memory slot. The example also illustrates the presence of runtime flags (ln. [b]:4–7) related to the par/and and termination (ln. [a]:8–12), which also reside in reusable slots in the static memory.

3.3. Static and Lightweight Trail Allocation

Each line of execution in CÉU needs to carry associated data, such as which event it is currently awaiting and which code to execute when it awakes. The compiler statically infers the maximum number of trails a program can have at the same time and creates a static vector to hold the runtime information about them. Like normal variables, trails that cannot be active at the same time share slots in the static memory vector.

At any given moment, a trail can be awaiting in one of the following states: `INACTIVE`, `STACKED`, `FINALIZE`, or in any of the events defined in the program:

```

enum {
    INACTIVE = 0,
    STACKED,
    FINALIZE,
    EVT_A,      // input void A;
    EVT_e,      // event int e;
    <...>      // other events
}

```

All terminated or not-yet-started trails stay in the `INACTIVE` state and are ignored by the scheduler. A `STACKED` trail holds an associated numeric stack level and can only execute when scheduler runtime reaches that level. A `FINALIZE` trail represents a hanged finalization block which is scheduled only when its corresponding block goes out of

scope. A trail waiting for an event stays in the state of the corresponding event, also holding the minimum sequence number (*seqno*) in which it can awake (for *delayed awaits*).

to prevent from awaking in the same reaction it is reached. In concrete terms, a trail is represented by the following struct:

```
struct trail_t {
    state_t evt;           // awaiting event
    label_t lbl;           // awaking execution label
    union {
        unsigned char seqno; // if evt=EVT.*
        stack_t stk;         // if evt=STACKED
    };
};
```

The field `evt` holds the state of the trail (or the event it is awaiting); the field `lbl` holds the entry point in the code to execute when the trail segment is scheduled; the third field depends on the `evt` field and may hold the `seqno` for an event, or the stack level `stk` for a `STACKED` trail.

The size of `state_t` depends on the number of events in the application; for an application with less than 253 events (plus the 3 states), one byte is enough. The size of `label_t` depends primarily on the number of `await` statements in the application—each `await` splits the code in two segments and requires a unique entry point in the code for its continuation. Additionally, `split & join` points for parallel compositions, `emit` continuations, and finalization blocks also require labels. Since the scheduler traverses all trails on every reaction, it adjusts all `seqno` to avoid overflows (actually, 2 bits to hold the `seqno` is already enough). The size of `stack_t` depends on the maximum depth of nested emissions but is bounded to the maximum number of trails: in the worst case, a trail emits an event that awakes another trail, which emits an event that awakes another trail, and so on; the last trail cannot awake any trail, because they are all hanged in the `STACKED` state.

In the context of embedded systems, the size of `trail_t` is typically only 3 bytes (1 byte for each field), imposing a negligible memory overhead even for trails that only await a single event and terminate. For instance, the *CTP* collection protocol ported to CÉU reaches eight simultaneous lines of execution with memory overhead of only 2% in comparison to the original single-threaded version in *C* [Sant’Anna et al. 2013b].

3.4. Static Scheduling and Trail Finalization

In the final generated code in *C*, each trail segment label representing an entry point becomes a *switch case* with the associated code to execute. Figure 16 illustrates the generation process. For the program in [a], the compiler extracts the entry points and associated trails, e.g., the label `Awake_e` will execute on `TRAIL-0` (ln. [a]:7 and [b]:3). For each statement that pauses (`emit` and `await`), resumes (`par/and`, `par/or`, and `finalize`), or aborts (`par/or` and `break`), the compiler splits the trail in two segments with associated entry points. The entry points translate to an `enum` in the generated code (ln. [b]:1–10). The state of trails translates to a vector of type `trail_t` with the maximum number of simultaneous trails (ln. [b]:12–15). On initialization, `TRAIL-0` is set to execute the `Main` entry point (ln. [b]:13), while all others are set to `INACTIVE` (ln. [b]:14).

The scheduler executes in two passes: in the *broadcast* pass, it sets all trails that are waiting for the current event to the state `STACKED` in the current numeric stack level; in the *dispatch* pass, it executes each trail that is `STACKED` to run in the current level, setting it immediately to `INACTIVE`.

During the dispatch pass, if a trail executes and emits an internal event, the scheduler increments the stack level and re-executes the two passes. After all trails are properly dispatched, the scheduler decrements the stack level and resumes the previ-

<pre> 1 input void A; 2 event void e; 3 // TRAIL 0 - lbl Main 4 par/and do 5 // TRAIL 0 - lbl Main 6 await e; 7 // TRAIL 0 - lbl Awake_e 8 // TRAIL 0 - lbl And.chk 9 with 10 // TRAIL 1 - lbl And.sub.2 11 await A; 12 // TRAIL 1 - lbl Awake.A.1 13 emit e; 14 // TRAIL 1 - lbl Emit.cont 15 // TRAIL 1 - lbl And.chk 16 end 17 // TRAIL 0 - lbl And.out 18 await A; 19 // TRAIL 0 - lbl Awake.A.2 20 21 22 23 24 25 </pre>	<pre> 1 enum { 2 Main = 1, // ln 3 3 Awake_e, // ln 7 4 And.chk, // ln 8,15 5 And.sub.2, // ln 10 6 Awake.A.1, // ln 12 7 Emit.cont, // ln 14 8 And.out, // ln 17 9 Awake.A.2 // ln 19 10 }; 11 12 trail.t TRLS[2] = { 13 { STACKED, Main, 0 }; 14 { INACTIVE, 0, 0 }; 15 }; 16 17 18 19 20 21 22 23 24 25 </pre>	<pre> 1 void dispatch (trail.t* t) { 2 switch (t->lbl) { 3 case Main: 4 // activate TRAIL 1 5 TRLS[1].evt = STACKED; 6 TRLS[1].lbl = And.sub.2; 7 TRLS[1].stk = cur.stack; 8 9 // code in the 1st trail 10 // await e; 11 TRLS[0].evt = EVT_e; 12 TRLS[0].lbl = Awake_e; 13 TRLS[0].seq = cur.seqno; 14 break; 15 16 case And.sub.2: 17 // await A; 18 TRLS[1].evt = EVT_A; 19 TRLS[1].lbl = Awake.A.1; 20 TRLS[1].seq = cur.seqno; 21 break; 22 23 <...> // other labels 24 } 25 } </pre>
[a]	[b]	[c]

Fig. 16. [a] Static allocation of trails: the comments identify the trail indexes inferred by the compiler; [b] Entry-point labels: each trail segment has an associated numeric identifier generated by the compiler. [c] Dispatch function: uses a switch to associate each segment identifier with the corresponding code to execute.

ous execution. For the boot reaction, the scheduler starts from the *dispatch* pass, given that the Main label is the only one that can be active at the stack level 0 (ln. [b]: 13).

The code in [c] dispatches a trail segment according to the current label to execute. For the first reaction, it executes the Main label in TRAIL-0. When the Main label reaches the par/and (ln. [a]:4), it stacks TRAIL-1 (ln. [c]:4–7) and proceeds to the code in TRAIL-0 (ln. [c]:9–14) to respect lexical execution order. The code sets the running TRAIL-0 to await EVT_e on label Awake_e, and then halts with a break. Then, dispatch takes TRAIL-1 and executes its registered label And.sub.2 (ln. 16–21), which sets TRAIL-1 to await EVT_A and also halts.

Regarding abortion and finalization, when a par/or terminates, the scheduler makes a *broadcast* pass for the FINALIZE event, but limited to the range of trails covered by the terminating par/or. Trails that do not match the FINALIZE are set to INACTIVE, as they have to be aborted. Given that trails in parallel are allocated in subsequent slots in the static vector TRLS (ln. [b]:12–15), this pass only aborts the desirable trails. The subsequent *dispatch* pass executes the finalization code properly. Escaping a loop that contains parallel compositions also triggers the same abortion process.

3.5. Single-Threaded Dispatching

The implementation of CÉU dispatches active trails sequentially in a single thread, taking no advantage of multi-core CPUs. This decision comes not only from the fact that CÉU targets constrained embedded systems with a single CPU, but also because CÉU imposes deterministic execution for intra-reaction statements.

Note that, as discussed in Section 2.3, the temporal analysis of CÉU infers precisely trails that are concurrent yet do not share resources. Hence, these non-conflicting trails could potentially execute with real parallelism in multiple cores. However, our experiments with multi-core execution are actually slower than single-core execution in the same system. Considering that we use CÉU primarily in control-dominated applica-

tions, this result is not surprising and also appears in related work [Yuan et al. 2011; Haribi 2012]. One reason is due to continuous fork-and-rejoin overhead from small reactions. Another reason is due to excessive locality of data from stackless trails sharing contiguous static memory.

If we consider data-dominated applications, multi-core implementations can offer considerable speed-up gains. However, data-intensive computations do not typically require a disciplined step-wise execution and can actually execute in isolated asynchronous calls. Esterel provides a `task` primitive for this purpose [Berry 2000], while CÉU provides an equivalent `async/thread` primitive. Asynchronous execution is out of the scope of this paper.

Single-threaded dispatching may not be suitable for hard real-time activities. In a previous work [Sant’Anna et al. 2013b], we measure how synchronous lengthy computations in *C* (e.g., hashing and compression) can block the scheduler and affect higher-priority activities such as a radio driver. In such cases, the system requires careful testing to avoid undersized hardware deployment. For instance, we currently do not perform any worst-case reaction times analysis [Boldt et al. 2008; Li et al. 2005].

3.6. Interaction with the Environment

As a reactive language, the execution of programs in CÉU is guided entirely by the occurrence of external input events. The binding for a specific platform (environment) calls hook functions in the API of the CÉU runtime whenever an external event occurs. However, to preserve the sequential/discrete semantics of time in CÉU, these calls must never interleave or parallelize execution.

Figure 17 shows our binding for *TinyOS* [Hill et al. 2000], which maps system callbacks to input events in CÉU. The file `ceu_app.h` (ln. 3) contains all definitions for the compiled CÉU program, which are further queried through `#ifdef`’s. The file `ceu_app.c` (ln. 4) contains the runtime of CÉU with the scheduler and dispatcher pointing to the labels defined in the program. The callback `Boot.booted` (ln. 6–11) is called by *TinyOS* on startup, so we initialize CÉU inside it (ln. 7). If the CÉU program uses timers, we also start a periodic timer (ln. 8–10) that triggers callback `Timer.fired` (ln. 13–17) every 10 milliseconds to advance the wall-clock time of CÉU (ln. 15)⁵. The remaining lines map pre-defined *TinyOS* events that can be used in CÉU programs, such as the light sensor (ln. 19–23) and the radio transceiver (ln. 25–36). The scheduler of *TinyOS* is already synchronous by default and always execute event handlers atomically, hence, the API calls to CÉU are properly serialized.

3.7. The Terra Virtual Machine

Terra is a system for programming wireless sensor network applications which uses CÉU as its scripting language [Branco et al. 2015]. Figure 18 shows the three basic elements of Terra: CÉU as the scripting language, a set of customized pre-built components, and the embedded virtual-machine engine which can disseminate and install bytecode images dynamically. This approach aims to combine the flexibility of remotely uploading code with the expressiveness and safety guarantees of CÉU.

The main difference between the standard *C* back end and the Terra *VM* is the *code generation phase*, which here outputs assembly instructions for the *VM*, instead of statements in *C*. To reduce the memory footprint of applications, the *VM* includes special instructions for complex and recurrent operations from the runtime of CÉU, such as for handling events and trails.

⁵We also offer a mechanism to start the underlying timer on demand to avoid the “battery unfriendly” 10ms polling.

```

1 implementation
2 {
3     #include "ceu_app.h"
4     #include "ceu_app.c"
5
6     event void Boot.booted () {
7         ceu_init();
8 #ifdef CEU_WCLOCKS
9         call Timer.startPeriodic(10);
10    #endif
11    }
12
13 #ifdef CEU_WCLOCKS
14     event void Timer.fired () {
15         ceu_wclock(10000);
16     }
17 #endif
18
19 #ifdef EVT_PHOTO_READDONE
20     event void Photo.readDone (int val) {
21         ceu_go(EVT_PHOTO_READDONE, &val);
22     }
23 #endif
24
25 #ifdef EVT_RADIO_SENDDONE
26     event void RadioSend.sendDone (message_t* msg) {
27         ceu_go(EVT_RADIO_SENDDONE, &msg);
28     }
29 #endif
30
31 #ifdef EVT_RADIO_RECEIVE
32     event message_t* RadioReceive.receive (message_t* msg) {
33         ceu_go(EVT_RADIO_RECEIVE, &msg);
34         return msg;
35     }
36 #endif
37
38     <...> // other events
39 }

```

Fig. 17. The *TinyOS* binding for CÉU. This platform-dependent template includes the *C* files generated from the original application in CÉU (*ceu_app.h* and *ceu_app.c*) for the *code generation phase*.




Fig. 18. Terra programming system basic elements.

In Terra, CÉU scripts cannot execute arbitrary *C* code, instead, they rely on pre-built components that can be customized for different application domains. Considering the domain of sensor networks, Terra already provides components organized in four areas: radio communication, group management, data aggregation, and local operations (e.g., access to sensors and actuators). When creating an instance of the *VM*, the programmer can choose whether or not to include each component, setting differ-

<pre> // Output events output void REQUEST_TEMPERATURE; output int REQUEST_SEND; // sends int value // Input events input int TEMPERATURE_DONE; // recvs int value input void SEND_DONE; // System calls function int getRadioID (void); </pre>	<pre> 1 // Output events 2 void VM.out(int evt_id, void* args) { 3 switch (id){ 4 case O_REQUEST_TEMPERATURE: 5 call TINYOS_TEMP.read(); 6 <...>; // O_REQUEST_SEND 7 } 8 } 9 10 // Input events 11 event TINYOS_TEMP.done (int val) { 12 VM.enqueue(I_TEMPERATURE_DONE, &val); 13 } 14 <...> // TINYOS_SEND.done 15 16 // System calls 17 void VM.function(int id, void* params) { 18 switch (id) { 19 case F_GET_RADIO_ID: 20 VM.push(TINYOS_NODE_ID); 21 } 22 } </pre>
[a]	[b]

Fig. 19. [a] C  U interface with customized VM. [b] The routine VM.out redirects all output events to the corresponding OS calls (ln. 1–8). Each *TinyOS* event callback calls VM.enqueue for the corresponding input event (ln. 10–14). System calls use VM.push for immediate return values (ln. 16–22).

ent abstraction boundaries for scripts. The generated VM has to be preloaded into the embedded devices before they are physically distributed.

The communication between scripts in C  U and the components in the VM is mostly through events: scripts emit requests through output events and await answers through input events. Terra also provides system calls for initialization and configuration of components (e.g., *getters* and *setters*). Figure 19.a shows a C  U interface with the available functionality for a customized VM (with temperature and radio components). Figure 19.b shows the associated bindings for output events (ln. 1–8), input events (ln. 10–14), and system calls (ln. 16–22). Note that all applications for the customized VM must comply with the same interface. In contrast, the template-based C back end (illustrated in Figure 17) allows applications to choose all possible combinations of functionalities from the underlying platform at compile time.

4. RELATED WORK

C  U has a strong influence from Esterel but differ in the most fundamental aspect of the notion of time. In C  U, instead of clock ticks, atomic external event occurrences that define time units. The event-driven approach of C  U is widespread [Ousterhout 1996] and popular in many software communities, such as web frameworks (e.g., *jQuery* [Chaffer 2009] and *Node.js* [Tilkov and Vinoski 2010]), GUI toolkits (e.g., *Tcl/Tk* [Ousterhout 1991] and *Java Swing* [Eckstein et al. 1998]), and Games [Nystrom 2014]. We consider that this approach is more familiar to programmers and simplifies the reasoning about concurrency. As far as we know, C  U is the first to encrust this event-triggered notion of time in the core of the language, which is a prerequisite for the temporal analysis that enables safe shared-memory concurrency.

Another unique aspect of C  U is the distinction between external and internal events. Internal events support stack-based micro reactions within external reactions, providing more fine-grained control for intra-reaction execution.

Like C  U, many synchronous languages rely on deterministic scheduling to preserve intra-reaction determinism (*Reactive C* [Boussinot 1991], *Protothreads* [Dunkels et al. 2006], *SOL* [Karpinski and Cahill 2007], *SC* [Von Hanxleden 2009], and *PRET-C* [An-

dalam et al. 2010]). CÉU goes one step further and performs a temporal analysis to detect trails that, when reordered, change the observable behavior of the program, i.e., trails that actually rely on deterministic scheduling. CÉU also guarantees deterministic behavior for timer compositions by adjusting inaccuracies in the system clock.

Regarding resource management, Esterel supports a finalization mechanism to unconditionally execute a series of statements on abortion. In addition, CÉU also tracks pointers representing resources that cross *C* boundaries and enforce the programmer to provide associated finalizers.

Esterel has different compilation back ends that synthesizes to software and also to hardware circuits [Dayaratne et al. 2005; Edwards 2003]. Among the software-based approaches, *SAXO-RT* [Closse et al. 2002] is the closest to our implementation with respect to trail allocation and scheduling: the compiler slices programs into “control points” (analogous to our “entry points”) and rearranges them into a directed acyclic graph respecting the constructive semantics of Esterel. Then, it flattens the graph into sequential code in *C* suitable for static scheduling.

A number of virtual machines have been proposed for embedded systems. *Darjeeling* [Brouwers et al. 2008] and *TakaTuka* [Aslam et al. 2010] are complete *Java VMs* targeting constrained embedded systems with support for multithreading and garbage collection. Java has antagonistic design choices in comparison to CÉU: it does not impose static bounds on memory usage and execution time, and provides preemptive multithreading which requires synchronization primitives for accessing shared memory. Plummer et al. [Plummer et al. 2006] propose a Esterel-based *VM* with similar design choices to our work. To reduce code size, the *VM* has a specialized instruction set to deal with events and concurrency constructs that are particular to Esterel. However, the proposed *VM* is only a proof of concept, with no support for arithmetic operations, external system calls, or remote reprogramming.

5. CONCLUSION

We present the design and implementation of CÉU, a synchronous reactive language inspired by Esterel with simpler semantics and more fine-grained control for intra-reaction execution.

CÉU is a concurrency-safe language, employing static analysis to ensure that the high degree of concurrency in embedded systems does not pose safety threats to applications. As a summary, the following safety properties hold for all programs that successfully compile in CÉU: time and memory-bounded reactions to the environment (except for system calls), no race conditions in shared memory, reliable abortion for activities handling resources, and automatic synchronization for timers. These properties are usually desirable in embedded applications and are guaranteed as preconditions in CÉU by design.

CÉU is a resource-efficient language suitable for constrained embedded systems. The reference implementation compiles to portable event-driven code in *C*, with no special requirements for OS threads or per-trail data stacks. The *VM* implementation uses the same front end and imposes no extra restrictions, being equally suitable for constrained systems.

CÉU is a practical language with expressive control constructs, such as lexically scoped parallel compositions, convenient first-class timers, and a unique stack-based mechanism for internal signalling. Programs interoperate seamlessly with *C*, and can take advantage of existing libraries, lowering the entry barrier for adoption. CÉU has

an open source implementation and bindings for *TinyOS*, *Arduino*, and the *SDL* graphical library.⁶

For the past three years, we have been teaching CÉU for undergraduate and graduate students in research projects and two hands-on courses on *distributed systems* and *reactive programming*. Our experience shows that students take advantage of the sequential-imperative style of CÉU and can implement non-trivial concurrent applications in a few weeks. More recently, a company specialized in embedded systems (not related to our research group) released a product based on CÉU.

REFERENCES

- Sidharta Andalam and others. 2010. Predictable multithreading of embedded applications using PRET-C. In *Proceeding of MEMOCODE'10*. IEEE, 159–168.
- Faisal Aslam and others. 2010. Optimized java binary and virtual machine for tiny motes. In *Distributed Computing in Sensor Systems*. Springer, 15–30.
- Albert Benveniste and others. 2003. The synchronous languages twelve years later. In *Proceedings of the IEEE*, Vol. 91. 64–83.
- Alexander Bernauer and Kay Römer. 2013. A Comprehensive Compiler-Assisted Thread Abstraction for Resource-Constrained Systems. In *Proceedings of IPSN'13*. Philadelphia, USA.
- Gérard Berry. 1993. Preemption in Concurrent Systems.. In *FSTTCS (LNCS)*, Vol. 761. Springer, 72–93.
- Gérard Berry. 1996. The Constructive Semantics of Pure Esterel. (1996).
- Gérard Berry. 2000. *The Esterel-V5 Language Primer*. CMA and Inria, Sophia-Antipolis, France. Version 5.10, Release 2.0.
- Gérard Berry, S Ramesh, and RK Shyamasundar. 1993. Communicating reactive processes. In *Proceedings of the 20th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*. ACM, 85–98.
- Gérard Berry and Ellen Sentovich. 2001. Multiclock esterel. In *Correct Hardware Design and Verification Methods*. Springer, 110–125.
- Marian Boldt, Claus Traulsen, and Reinhard von Hanxleden. 2008. Worst case reaction time analysis of concurrent reactive programs. *Electronic Notes in Theoretical Computer Science* 203, 4 (2008), 65–79.
- Frédéric Boussinot. 1991. Reactive C: An extension of C to program reactive systems. *Software: Practice and Experience* 21, 4 (1991), 401–428.
- Frédéric Boussinot. 1998. SugarCubes implementation of causality. (1998).
- F. Boussinot and R. de Simone. 1991. The Esterel language. *Proc. IEEE* 79, 9 (Sep 1991), 1293–1304.
- Adriano Branco, Francisco Sant'anna, Roberto Ierusalimsky, Noemi Rodriguez, and Silvana Rossetto. 2015. Terra: Flexibility and Safety in Wireless Sensor Networks. *ACM Trans. Sen. Netw.* 11, 4, Article 59 (Sept. 2015), 27 pages. DOI: <http://dx.doi.org/10.1145/2811267>
- Niels Brouwers, Peter Corke, and Koen Langendoen. 2008. Darjeeling, a Java compatible virtual machine for microcontrollers. In *Proceedings of the ACM/IFIP/USENIX Middleware'08 Conference Companion*. ACM, 18–23.
- Jonathan Chaffer. 2009. *Learning JQuery 1.3: Better Interaction and Web Development with Simple JavaScript Techniques*. Packt Publishing Ltd.
- Etienne Closse and others. 2002. Saxo-RT: Interpreting esterel semantic on a sequential execution structure. *Electronic Notes in Theoretical Computer Science* 65, 5 (2002), 80–94.
- Sajeewa Dayaratne and others. 2005. Direct Execution of Esterel Using Reactive Microprocessors. In *Proceedings of SLAP'05*.
- Dunkels and others. 2006. Protothreads: simplifying event-driven programming of memory-constrained embedded systems. In *Proceedings of SenSys'06*. ACM, 29–42.
- Robert Eckstein, Marc Loy, and Dave Wood. 1998. *Java swing*. O'Reilly & Associates, Inc.
- Stephen A. Edwards. 1999. Compiling Esterel into sequential code. In *7th International Workshop on Hardware/Software Codesign*. ACM, 147–151.
- Stephen A. Edwards. 2003. Tutorial: Compiling concurrent languages for sequential processors. *ACM Transactions on Design Automation of Electronic Systems* 8, 2 (2003), 141–187.
- Stephen A. Edwards. 2005. Using and Compiling Esterel. MEMOCODE'05 Tutorial. (July 2005).

⁶Website of CÉU: <http://www.ceu-lang.org/>

- David Gay and others. 2003. The nesC Language: A Holistic Approach to Networked Embedded Systems. In *PLDI'03*. 1–11.
- Wahbi Haribi. 2012. Compiling Esterel for Multi-Core Execution. *Synchrone Sprachen* (2012), 45.
- Hill and others. 2000. System architecture directions for networked sensors. *SIGPLAN Notices* 35 (November 2000), 93–104. Issue 11.
- Marcin Karpinski and Vinny Cahill. 2007. High-Level Application Development is Realistic for Wireless Sensor Networks. In *Proceedings of SECON'07*. 610–619.
- Oliver Kasten and Kay Römer. 2005. Beyond Event Handlers: Programming Wireless Sensors with Attributed State Machines. In *Proceedings of IPSN '05*. 45–52.
- Xin Li, Jan Lukoschus, Marian Boldt, Michael Harder, and Reinhard Von Hanxleden. 2005. An Esterel processor with full preemption support and its worst case reaction time analysis. In *Proceedings of the 2005 international conference on Compilers, architectures and synthesis for embedded systems*. ACM, 225–236.
- Ana Lúcia De Moura and Roberto Ierusalimsky. 2009. Revisiting coroutines. *ACM TOPLAS* 31, 2 (Feb. 2009), 6:1–6:31.
- Robert Nystrom. 2014. *Game Programming Patterns*. Genever Benning.
- John K. Ousterhout. 1991. An X11 Toolkit Based on the Tcl Language.. In *USENIX Winter*. 105–116.
- John K. Ousterhout. 1996. Why Threads Are A Bad Idea (for most purposes). (January 1996).
- Becky Plummer, Mukul Khajanchi, and Stephen A. Edwards. 2006. An Esterel virtual machine for embedded systems. In *International Workshop on Synchronous Languages, Applications, and Programming (SLAP'06)*. Citeseer, Vienna, Austria.
- Dumitru Potop-Butucaru and others. 2005. The Synchronous Hypothesis and Synchronous Languages. In *Embedded Systems Handbook*, R. Zurawski (Ed.).
- Partha S Roop, Zoran Salic, and MW Dayaratne. 2004. Towards direct execution of Esterel programs on reactive processors. In *Proceedings of the 4th ACM international conference on Embedded software*. ACM, 240–248.
- Francisco Sant'Anna. 2013. *Safe System-level Concurrency on Resource-Constrained Nodes with Céu*. Ph.D. Dissertation. PUC-Rio.
- Francisco Sant'Anna and others. 2012. *Céu: Embedded, Safe, and Reactive Programming*. Technical Report 12/12. PUC-Rio.
- Francisco Sant'Anna and others. 2013a. Advanced Control Reactivity for Embedded Systems. Workshop on Reactivity, Events and Modularity (REM'13). (2013).
- Francisco Sant'Anna and others. 2013b. Safe System-level Concurrency on Resource-Constrained Nodes. In *Proceedings of SenSys'13*. ACM.
- Klaus Schneider and Michael Wenz. 2001. A new method for compiling schizophrenic synchronous programs. In *Proceedings of the 2001 international conference on Compilers, architecture, and synthesis for embedded systems*. ACM, 49–58.
- Ellen M Sentovich. 1997. Quick conservative causality analysis. In *System Synthesis, 1997. Proceedings., Tenth International Symposium on*. IEEE, 2–8.
- Thomas R Shiple, Gerard Berry, and Hemé Touati. 1996. Constructive analysis of cyclic circuits. In *European Design and Test Conference, 1996. ED&TC 96. Proceedings*. IEEE, 328–333.
- Olivier Tardieu and Robert De Simone. 2004. Curing schizophrenia by program rewriting in Esterel. In *Formal Methods and Models for Co-Design, 2004. MEMOCODE'04. Proceedings. Second ACM and IEEE International Conference on*. IEEE, 39–48.
- Stefan Tilkov and Steve Vinoski. 2010. Node.js: Using JavaScript to build high-performance network programs. *IEEE Internet Computing* 6 (2010), 80–83.
- Reinhard Von Hanxleden. 2009. SyncCharts in C: a proposal for light-weight, deterministic concurrency. In *Proceedings of the seventh ACM international conference on Embedded software*. ACM, 225–234.
- Simon Yuan, Li Hsien Yoong, and Partha S Roop. 2011. Compiling Esterel for multi-core execution. In *Digital System Design (DSD), 2011 14th Euromicro Conference on*. IEEE, 727–735.
- Jeong-Han Yun, Chul-Joo Kim, Seonggun Kim, Kwang-Moo Choe, and Taisook Han. 2013. Detection of harmful schizophrenic statements in esterel. *ACM Transactions on Embedded Computing Systems (TECS)* 12, 3 (2013), 80.